



Systems




Desktop PC




Server



Macbook Pro w/ Retina



iPhone 6/6S




Memory in a Modern System

Overview

- **Review of Memory Technologies**
- **Overview of Memory Hierarchy**
- **Cache Design Principles**
- **Learning Objectives**
 - Why is that some memories slow ?
 - What is memory hierarchy?
 - Why do we need memory hierarchy?
 - What is a Cache?

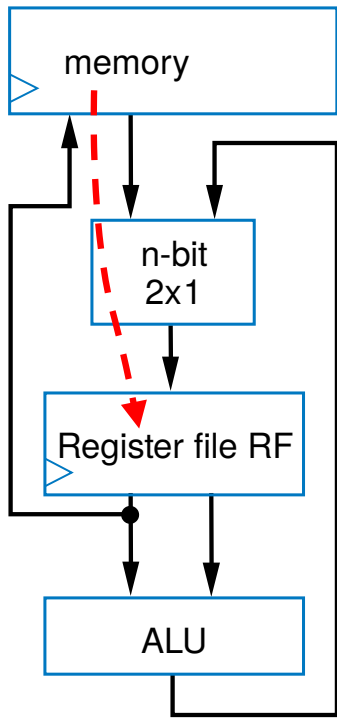
Processor-DRAM Memory Gap (latency)



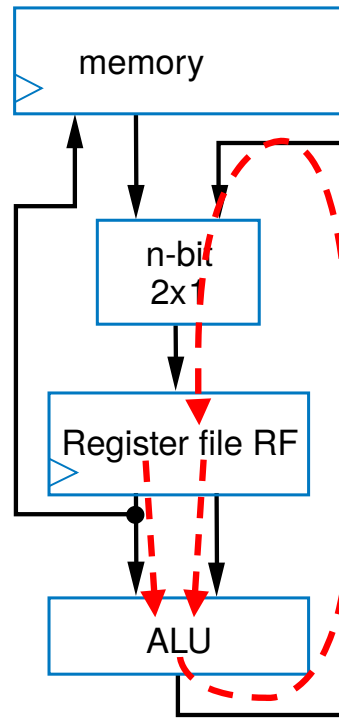
Source: David Patterson, UCB

Review: Datapath Operations

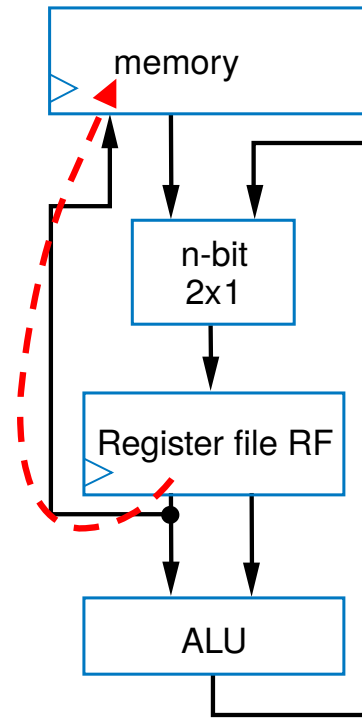
- Load operation: Load data from data memory to RF
- ALU operation: Transforms data by passing one or two RF register values through ALU, performing operation (ADD, SUB, AND, OR, etc.), and writing back into RF.
- Store operation: Stores RF register value back into data memory
- Each operation can be done in one clock cycle




Load operation



ALU operation

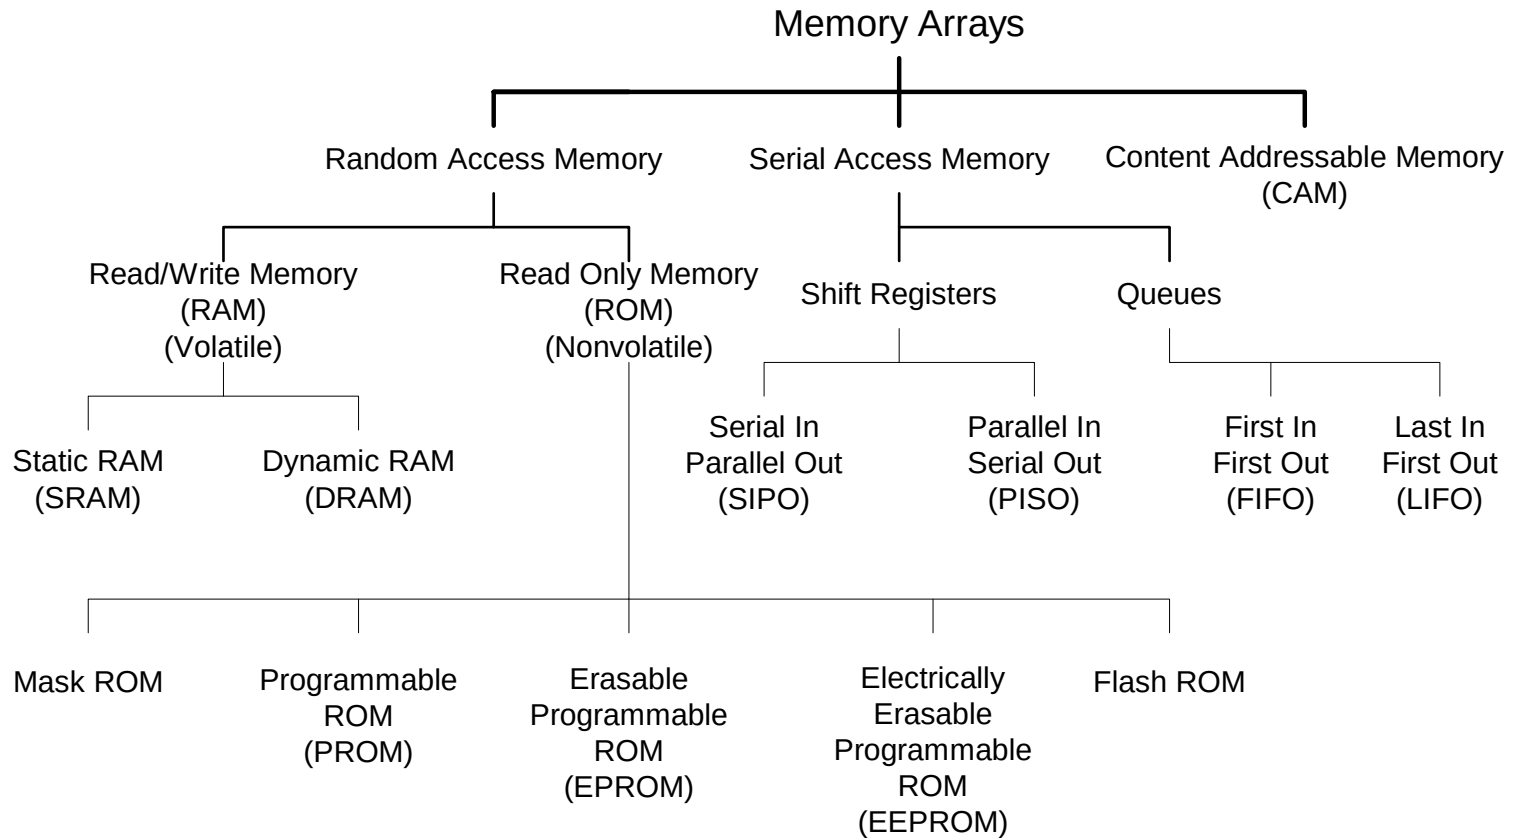


Store operation



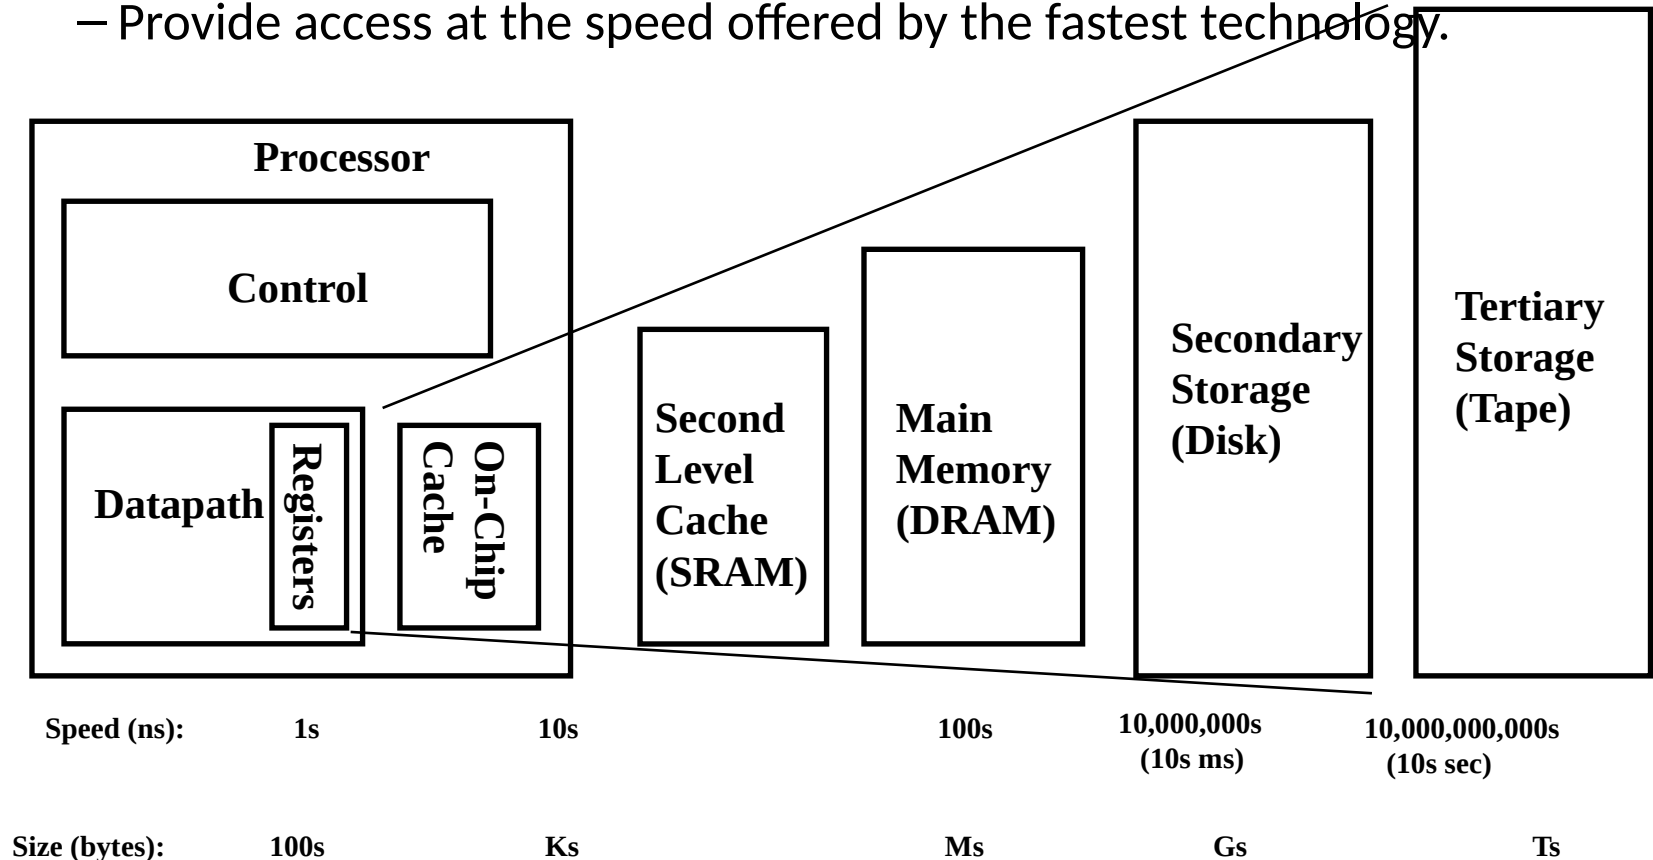
What are the Common Memory Technologies?

Memory




Memory Hierarchy of a Modern Computer System

- By taking advantage of the principle of locality:
 - Present the user with as much memory as is available in the cheapest technology.
 - Provide access at the speed offered by the fastest technology.



Memory

Address	Data
000000	00111110
000001	01101011
000010	01011101
000011	01100011
000100	00111110
000101	00000000
000110	11111111
000111	01010101
001000	10101010
001001	00100001
001010	11011010



D3
D2
D1

8x4 RAM

Address

Data


000	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>
001	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>
010	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>
011	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>
100	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>
101	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>
110	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>
111	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>

A2

A1

A0

8x4 RAM




Static RAM Organization

Chip Select Line (active lo)

Write Enable Line (active lo)

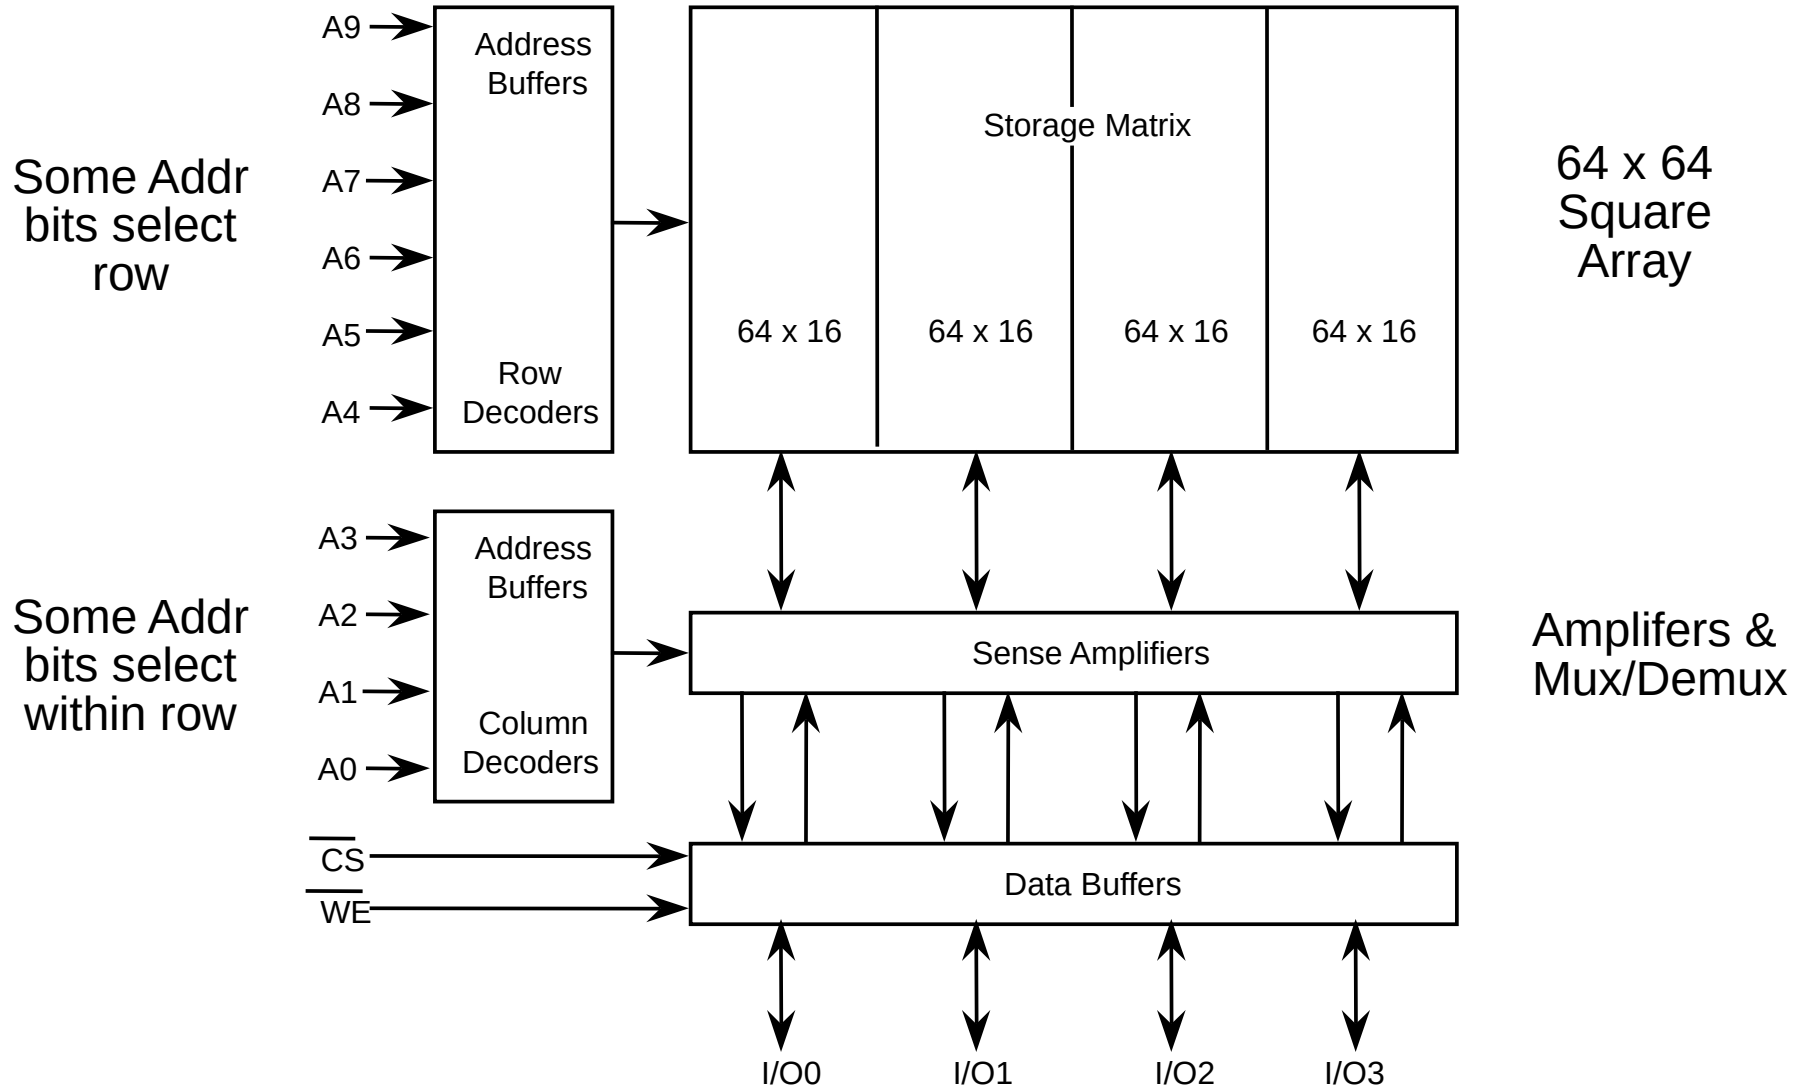
10 Address Lines

4 Bidirectional Data Lines

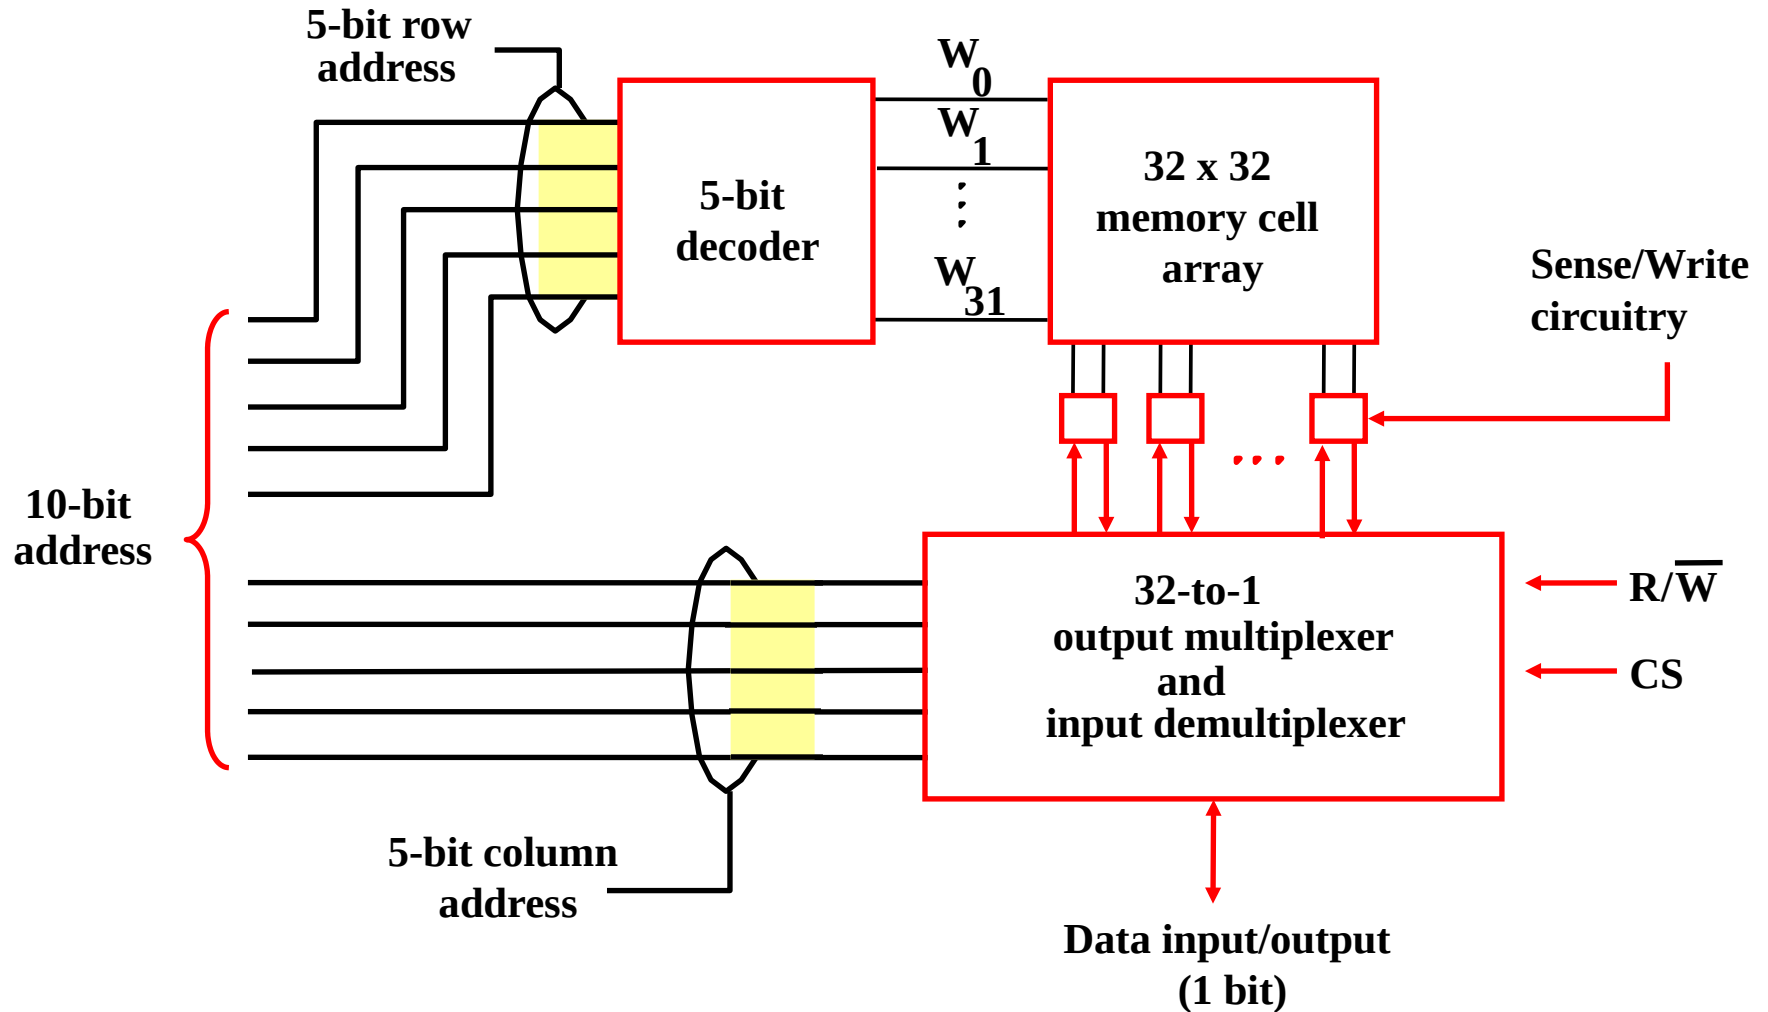


RAM Organization

Long thin layouts are not the best organization for a RAM



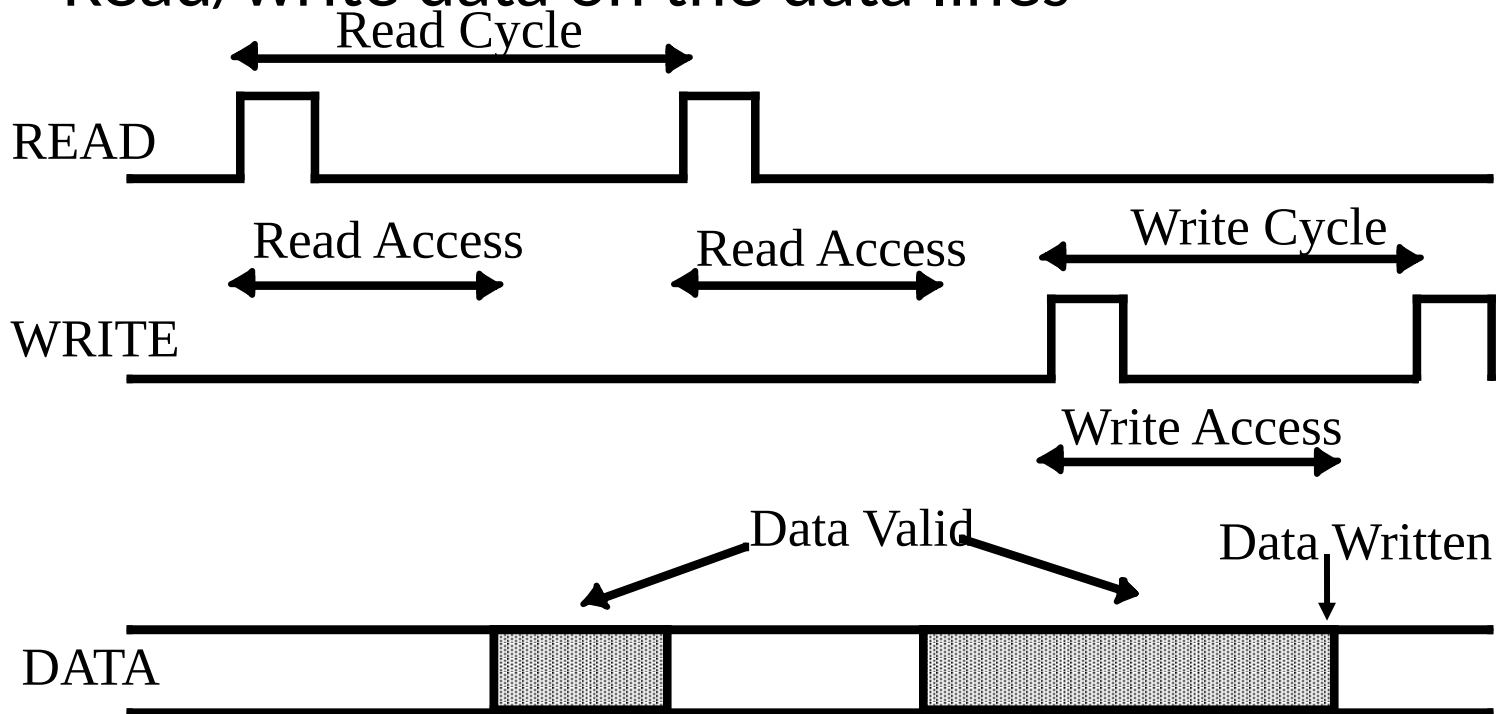
1Kx 1 bit RAM Organization



Organization of a $1K \times 1$ memory chip.

Memory Access Timing: the Big Picture

- Timing:
 - Send address on the address lines, wait for the word line to become stable
 - Read/write data on the data lines



Content of a memory

- Each **word in memory** is assigned an **identification number**, called an **address**, starting **from 0 up to 2^k-1** , where **k** is the **number of address lines**.
- The number of words in a memory with one of the letters $K=2^{10}$, $M=2^{20}$, or $G=2^{30}$.

$$64K = 2^{16} \quad 2M = 2^{21}$$

$$4G = 2^{32}$$

Memory address		Memory content
Binary	decimal	
0000000000	0	1011010101011101
0000000001	1	1010101110001001
0000000010	2	0000110101000110
	⋮	⋮
1111111101	1021	1001110100010100
1111111110	1022	0000110100011110
1111111111	1023	1101111000100101

1024x16 Memory Module

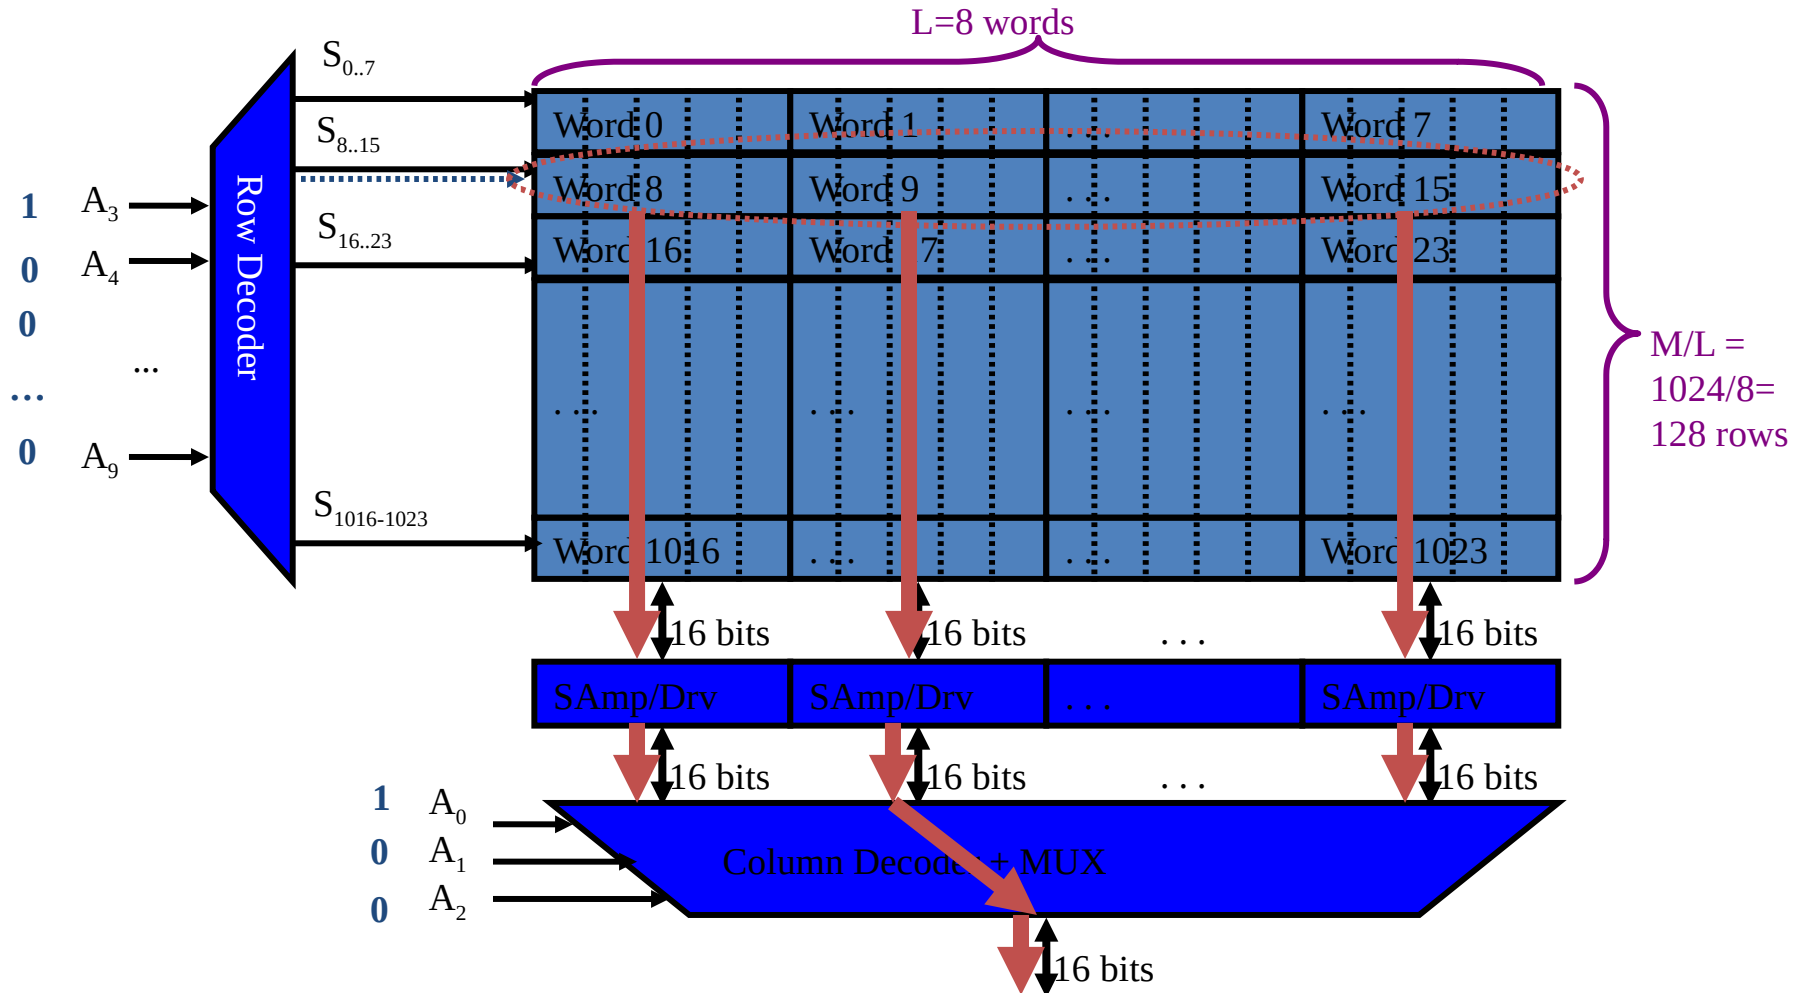
Word-Addressable Memory

- Each 32-bit data word has a unique address

Word Address	Data	
⋮	⋮	⋮
00000003	4 0 F 3 0 7 8 8	Word 3
00000002	0 1 E E 2 8 4 2	Word 2
00000001	F 2 F 1 A C 0 7	Word 1
00000000	A B C D E F 7 8	Word 0

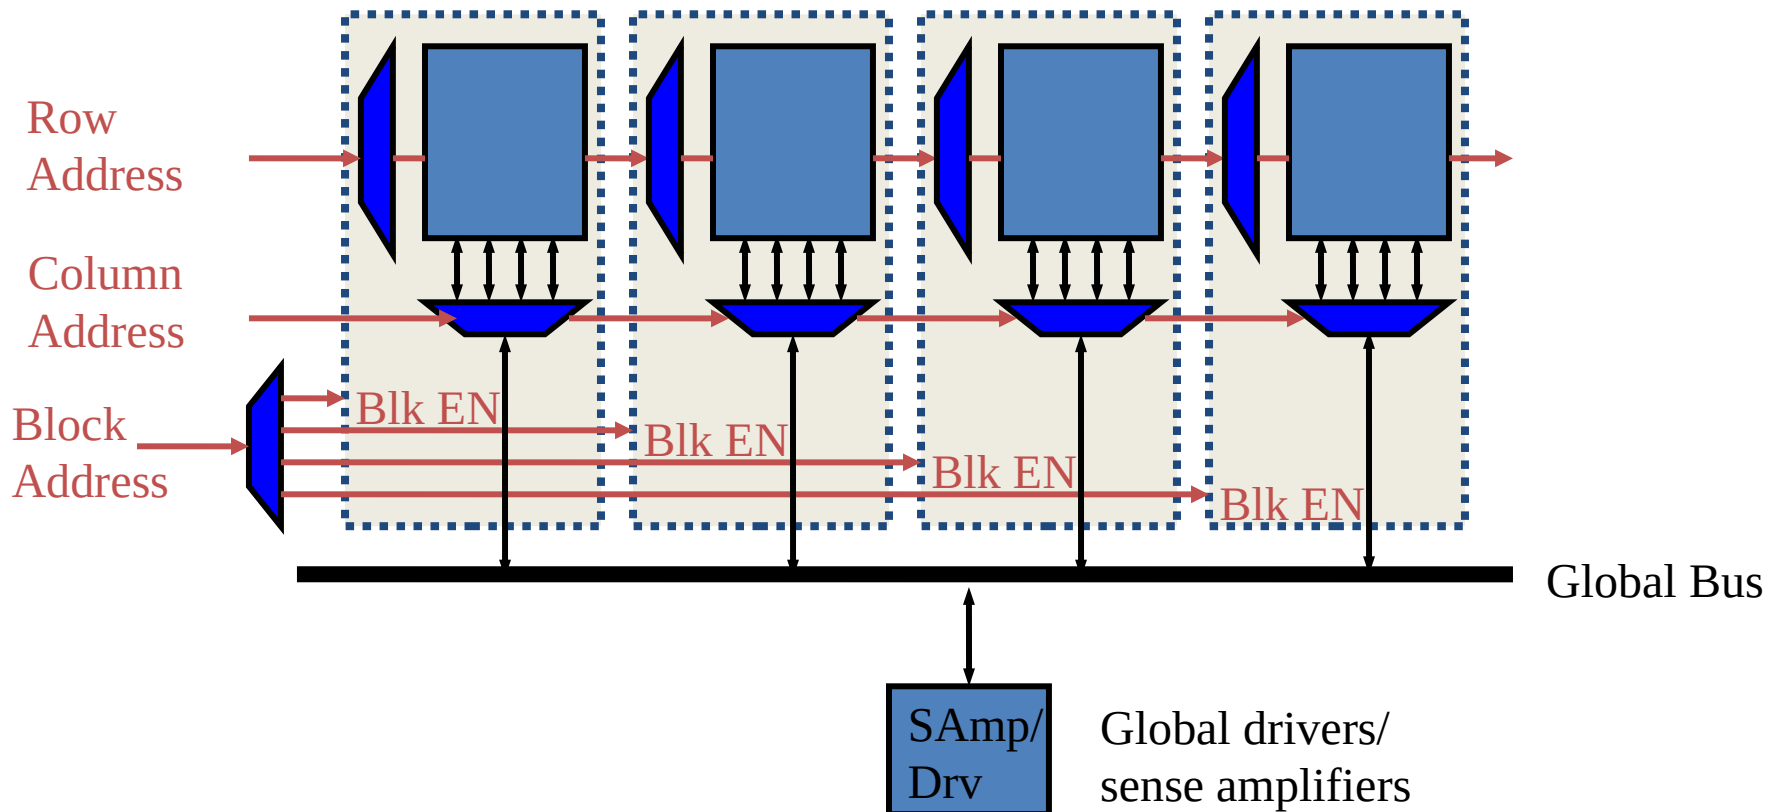
Memory Cell Array Access Example

- word=16-bit wide(N), row=8 words(L), address=10 bits (k)
- Accessing word 9= 0000001001_2




Memory Structures


- Taking the idea one step further
 - Shorter wires within each block
 - Enable only one block addr decoder \Rightarrow power savings



Larger Memories Using Multiple Chips




Question ? Design a 2Mx32 given: 512kx8




Organization of a $2M \times 32$ memory module using $512K \times 8$ static memory chips (16 chips).

19-bit internal chip address





Organization of a $2M \times 32$ memory module using $512K \times 8$ static memory chips (16 chips).



Organization of a $2M \times 32$ memory module using $512K \times 8$ static memory chips (16 chips).

Random Access Memory (RAM)


- RAM – Readable and writable memory
 - “Random access memory”
 - Strange name—Created several decades ago to contrast with sequentially-accessed storage like tape drives
 - Logically same as register file—Memory with address inputs, data inputs/outputs, and control
 - RAM usually one port; RF usually two or more
 - RAM vs. RF
 - RAM typically larger than *about* 512 or 1024 words
 - RAM typically stores bits using a bit storage approach that is more efficient than a flip-flop
 - RAM typically implemented on a chip in a square rather than rectangular shape—keeps longest wires (hence delay) short



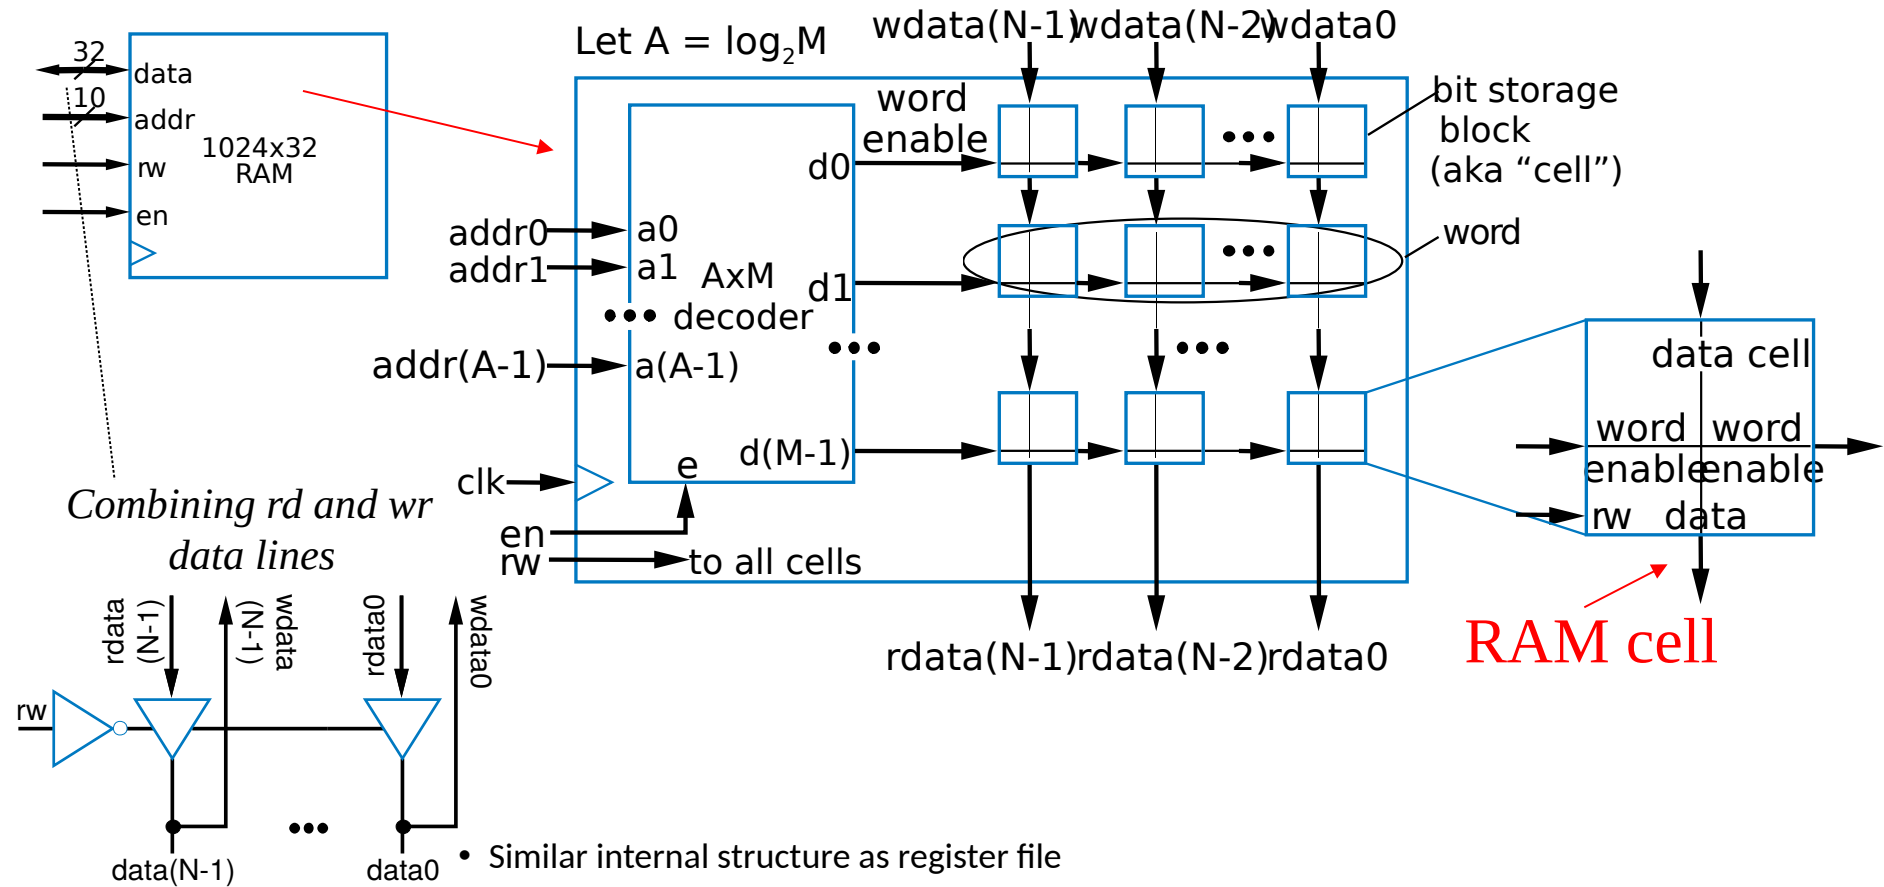
RAM block symbol

Implementing Registers in CMOS

- Uses transmission gate
 - When “WR” asserted, “write” operation will take place
 - Stack D latch structures to get n-bit register

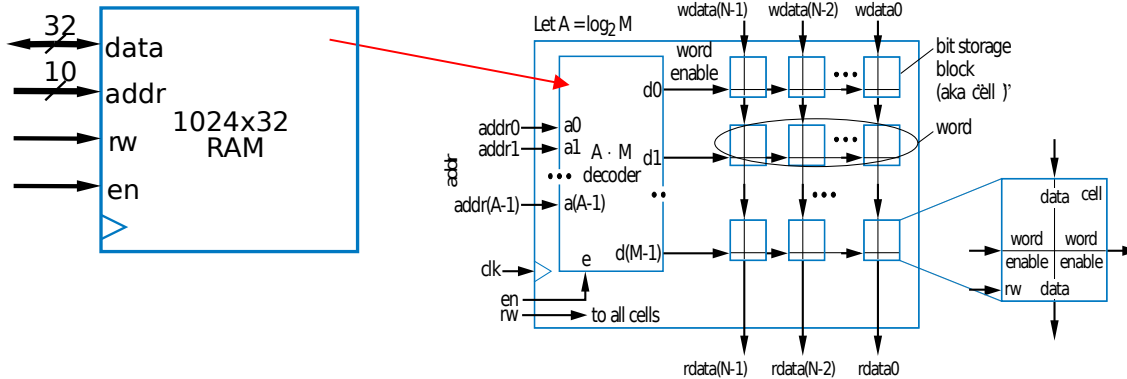


RAM Internal Structure




- Similar internal structure as register file
 - Decoder enables appropriate word based on address inputs
 - rw controls whether cell is written or read
 - rd and wr data lines typically combined
 - Let's see what's inside each RAM cell

Static RAM (SRAM)





- “Static” RAM cell
 - 6 transistors (recall inverter is 2 transistors)
 - Writing this cell
 - *word enable* input comes from decoder
 - When 0, value *d* loops around inverters
 - That loop is where a bit stays stored
 - When 1, the *data* bit value enters the loop
 - *data* is the bit to be stored in this cell
 - *data'* enters on other side
 - Example shows a “1” being written into cell

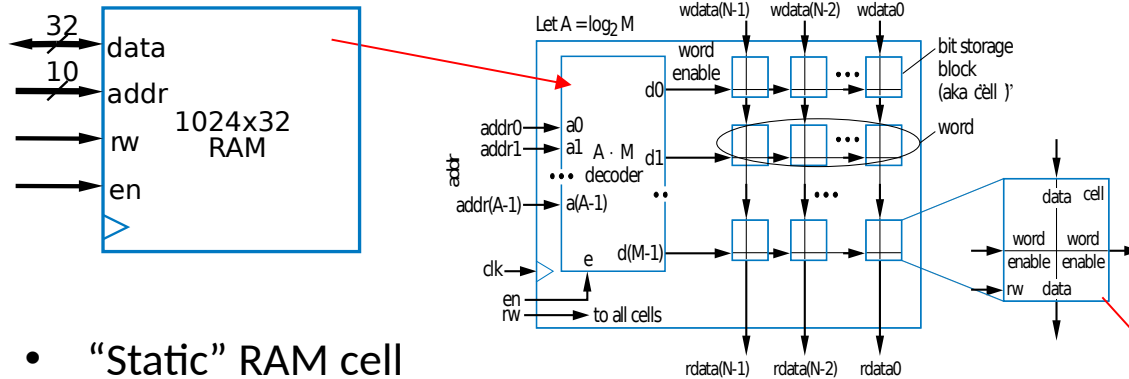
SRAM cell



SRAM cell

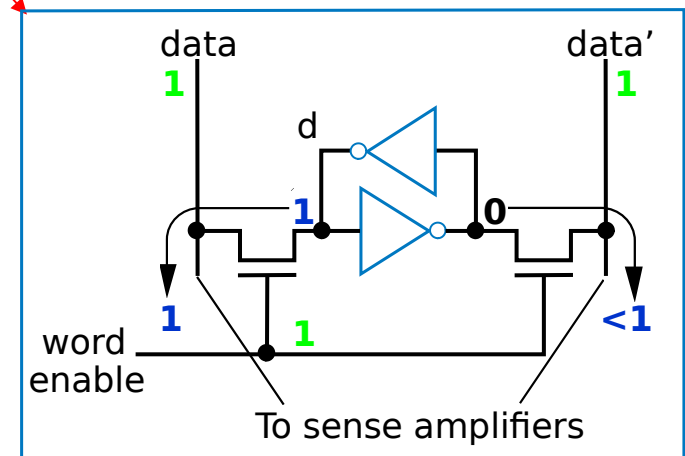


Static RAM (SRAM)





- “Static” RAM cell
 - Reading this cell
 - Somewhat trickier
 - When rw set to read, the RAM logic sets both *data* and *data'* to 1
 - The stored bit *d* will pull either the left line or the right bit down slightly below 1
 - “Sense amplifiers” detect which side is slightly pulled down
 - The electrical description of SRAM is really beyond our scope – just general idea here, mainly to contrast with DRAM...

SRAM cell

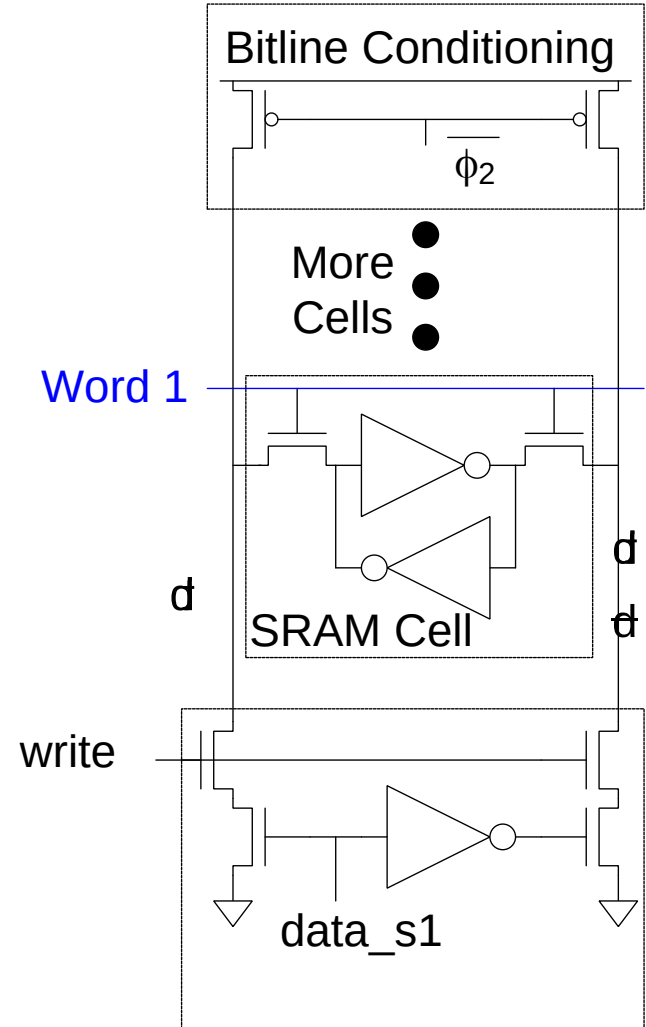


SRAM Column Example

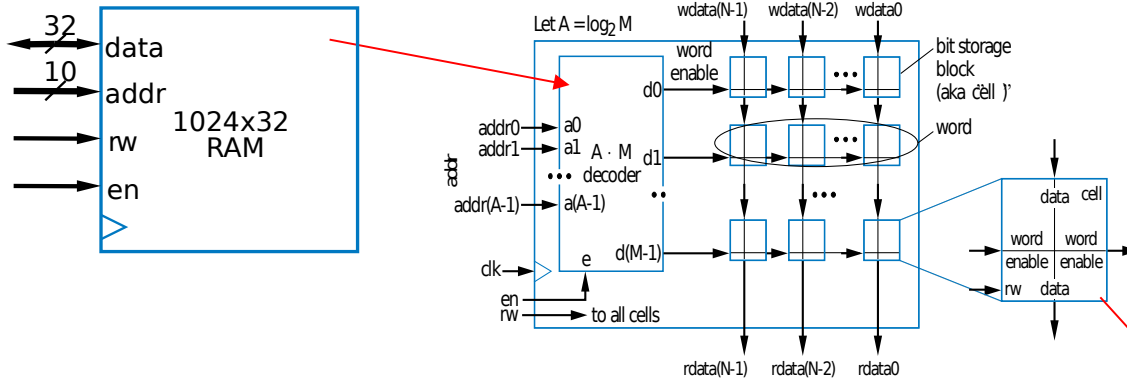
Read



Write

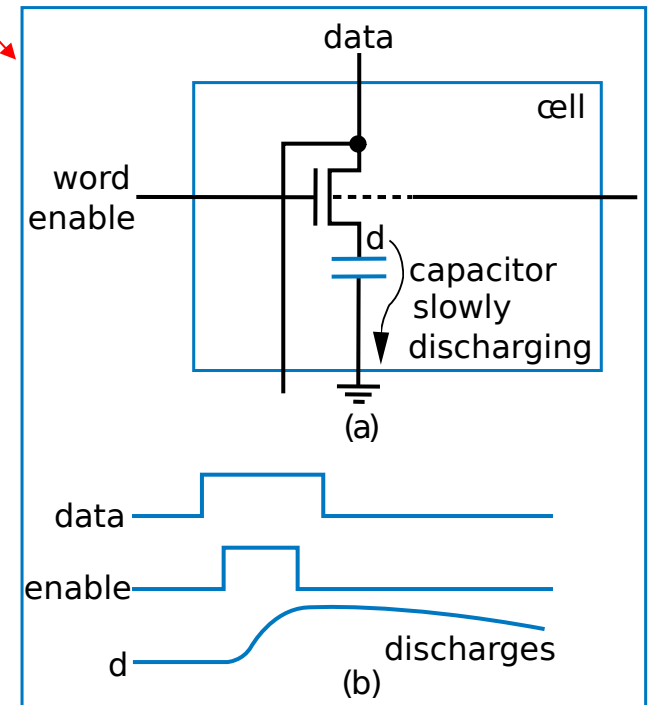


Dynamic RAM (DRAM)



- “Dynamic” RAM cell
 - 1 transistor (rather than 6)
 - Relies on *large* capacitor to store bit
 - Write: Transistor conducts, data voltage level gets stored on top plate of capacitor
 - Read: Just look at value of d
 - Problem: Capacitor discharges over time
 - Must “refresh” regularly, by reading d and then writing it right back

DRAM cell




- “Dynamic” RAM cell
 - 1 transistor (rather than 6)
 - Relies on capacitor to store bit





(a)

(b)

Dynamic RAM 1-Transistor Cell: Layout




Trench Cell



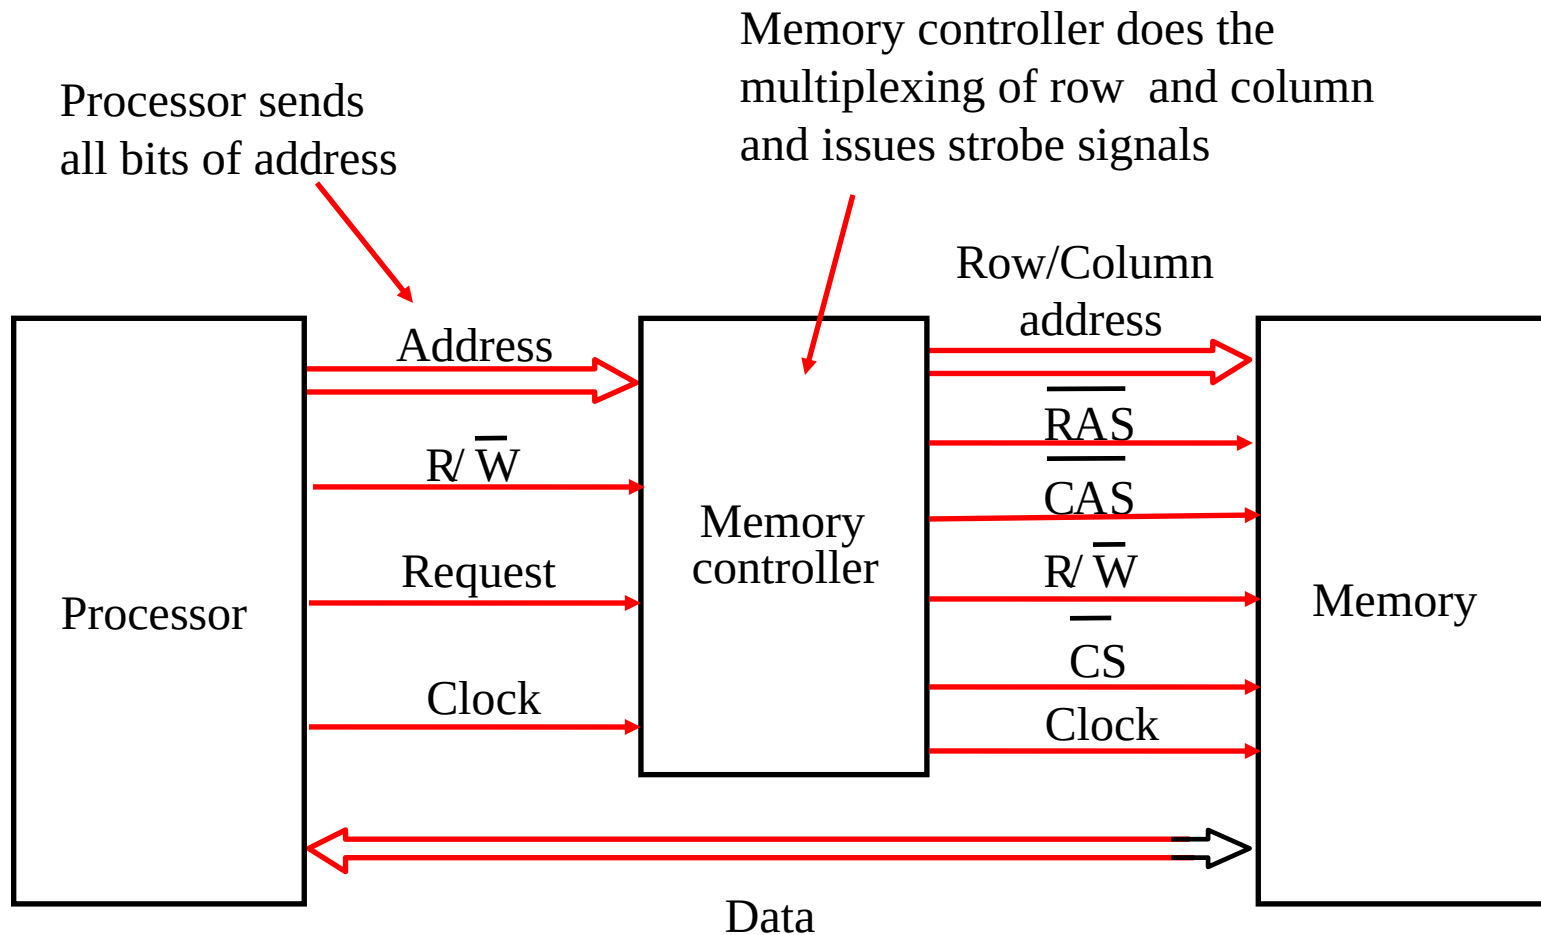
Stacked-capacitor Cell

Comparing Memory

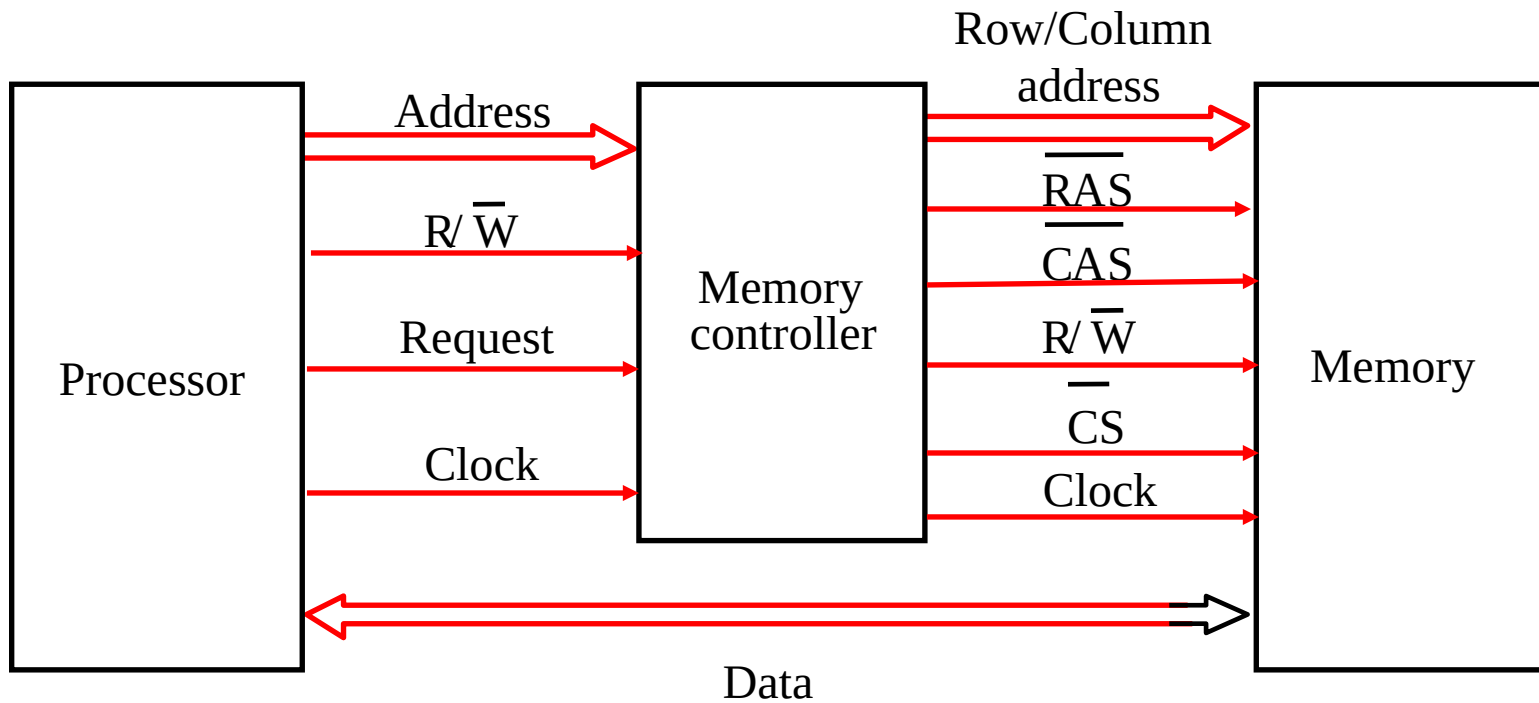
- Register file
 - Fastest
 - But biggest size
- SRAM
 - Fast
 - More compact than register file
- DRAM
 - Slowest (capacitor)
 - And refreshing takes time
 - But very compact (lower cost)
- Use register file for small items, SRAM for large items, and DRAM for huge items
 - Note: DRAM's big capacitor requires a special chip design process, so DRAM is often a separate chip



Size comparison for same number of bits (not to scale)



Use of a memory controller.



Memory controller provides the refresh control if not done on the chip


Refreshing typically once every 64 ms. At a cost of .2ms

Less than .4% overhead

Use of a memory controller.

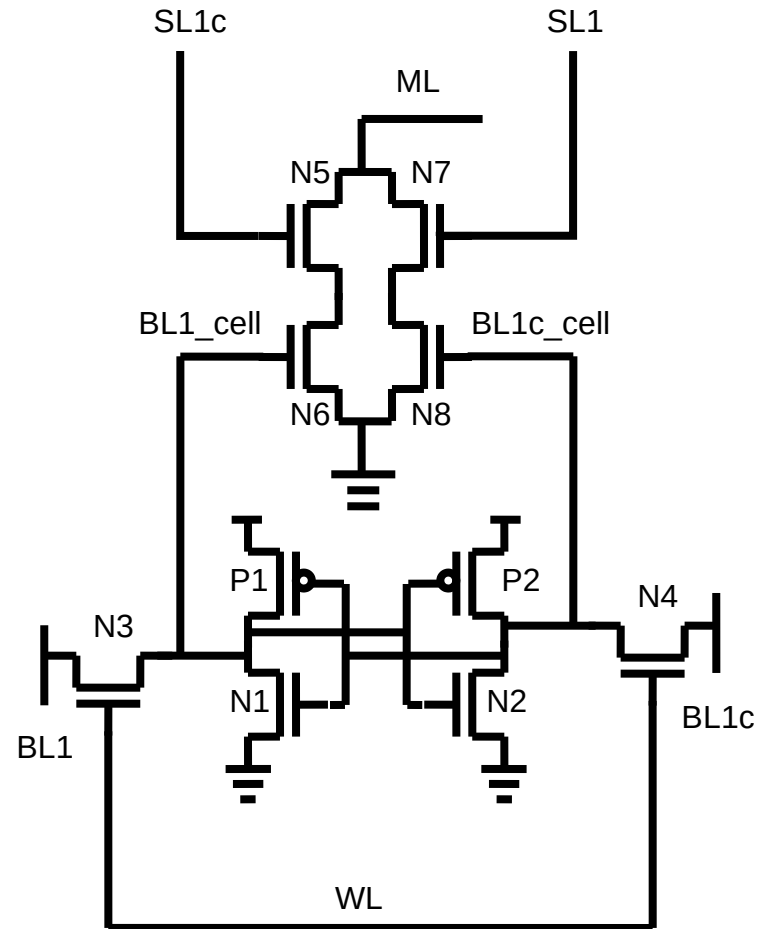
CAM

- CAM vs. RAM




CAM

- Binary CAM Cell
 - ML pre-charged to V_{DD}
 - Match: ML remains at V_{DD}
 - Mismatch: ML discharges



Read-Only Memory

- A block diagram of a ROM is shown below. It consists of k address inputs and n data outputs.
- The number of words in a ROM is determined from the fact that k address input lines are needed to specify 2^k words.





ROM Block diagram $2^k \times n$ Module

Read-Only Memory Cells



Bit line (BL) is resistively clamped to the ground, so its default value is 0

Diode disadvantage – no electrical isolation between bit and word lines


BL is resistively clamped to VDD, so its default value is 1



Diode ROM




MOS ROM 1



MOS ROM 2


Nonvolatile Memory

ROM
PROM
EPROM




Read-only memory organization, with the fixed contents shown on the right.

MOS NOR ROM




MOS NAND ROM



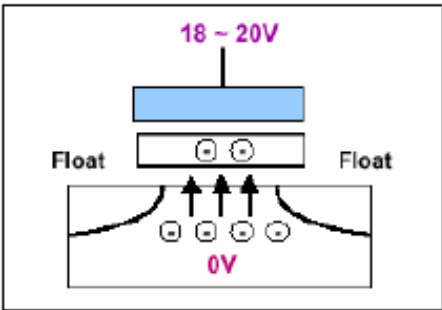
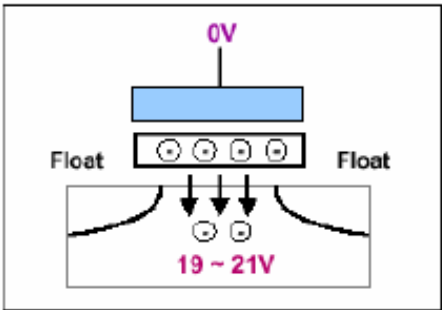
All word lines high by default with exception of selected row

Flash Memory




Flash memory

- Flash memory
 - Non volatile read only memory (ROM)
 - Erase Electrically or UV (EPROM)
 - Uses F-N tunneling for program & erase
 - Reads like DRAM (~ns)
 - Writes like DISK (~ms).

Program	Erase
	
<ul style="list-style-type: none">● Use F-N Tunneling● Channel Inversion	<ul style="list-style-type: none">● Use F-N Tunneling● Channel Accumulation


Reading Memory State



Change in Threshold Voltage due to **Screening Effect** of Floating Gate

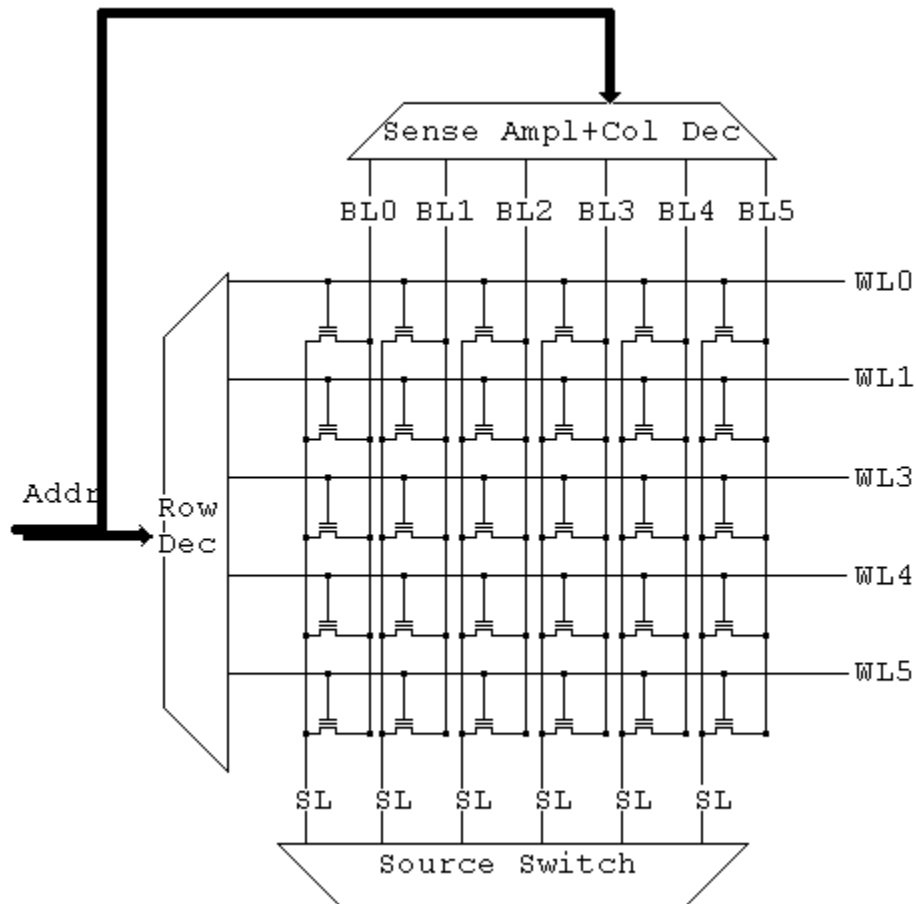
Read mode: Apply intermediate voltage, check whether current is flowing or not

Writing Memory State



Control gate voltage determines whether electrons are injected to, or push/pulled out of floating gate.

NOR Array




Reading:


Assert a single word line.
The source lines are asserted and the read of the bitline gives the contents of the cell.

Multi-Levels

- By using reference cells set at given levels and comparing them to the value from the bitline, we can determine the value stored.




Review of memory technologies




Memory type	SRAM	DRAM	Flash
Speed	Very fast	Slow	Slow
Density	Low	High	Very high
Power	High	Low	Very low
Refresh	No	Yes	No
Mechanism	Bi-stable latch	Capacitor	FN tunneling

Memory Hierarchy


Actual Memory Systems



+



Type1: Fast Small



Can we achieve?: Fast, Large
Type2: Slow, Large


(Cache Principle)

Locality

- **Principle of Locality:** Programs tend to use data and instructions with addresses near or equal to those they have used recently


- **Temporal locality:**

- Recently referenced items are likely to be referenced again in the near future



- **Spatial locality:**

- Items with nearby addresses tend to be referenced close together in time



Locality Example

- Data references
 - Reference array elements in succession
 - Reference variable sum each iteration.
- Instruction references
 - Reference instructions in sequence.
 - Cycle through loop repeatedly.

```
sum = 0;  
for (i = 0; i < n; i++)  
    sum += a[i];  
return sum;
```

Spatial locality

Temporal locality

Spatial locality

Temporal locality

Qualitative Estimates of Locality

- **Claim:** Being able to look at code and get a qualitative sense of its locality is a key skill for a professional programmer.
- **Question:** Does this function have good locality with respect to array *a*?

```
int sum_array_rows(int a[M][N])
{
    int i, j, sum = 0;

    for (i = 0; i < M; i++)
        for (j = 0; j < N; j++)
            sum += a[i][j];
    return sum;
}
```


Locality Example

- **Question:** Does this function have good locality with respect to array a?

```
int sum_array_cols(int a[M][N])
{
    int i, j, sum = 0;

    for (j = 0; j < N; j++)
        for (i = 0; i < M; i++)
            sum += a[i][j];
    return sum;
}
```

Locality reference example




Good locality example:

The X axis corresponds to time (in cycles) and the Y axis corresponds to memory locations. The highlighted sections show examples of temporal and spatial locality.

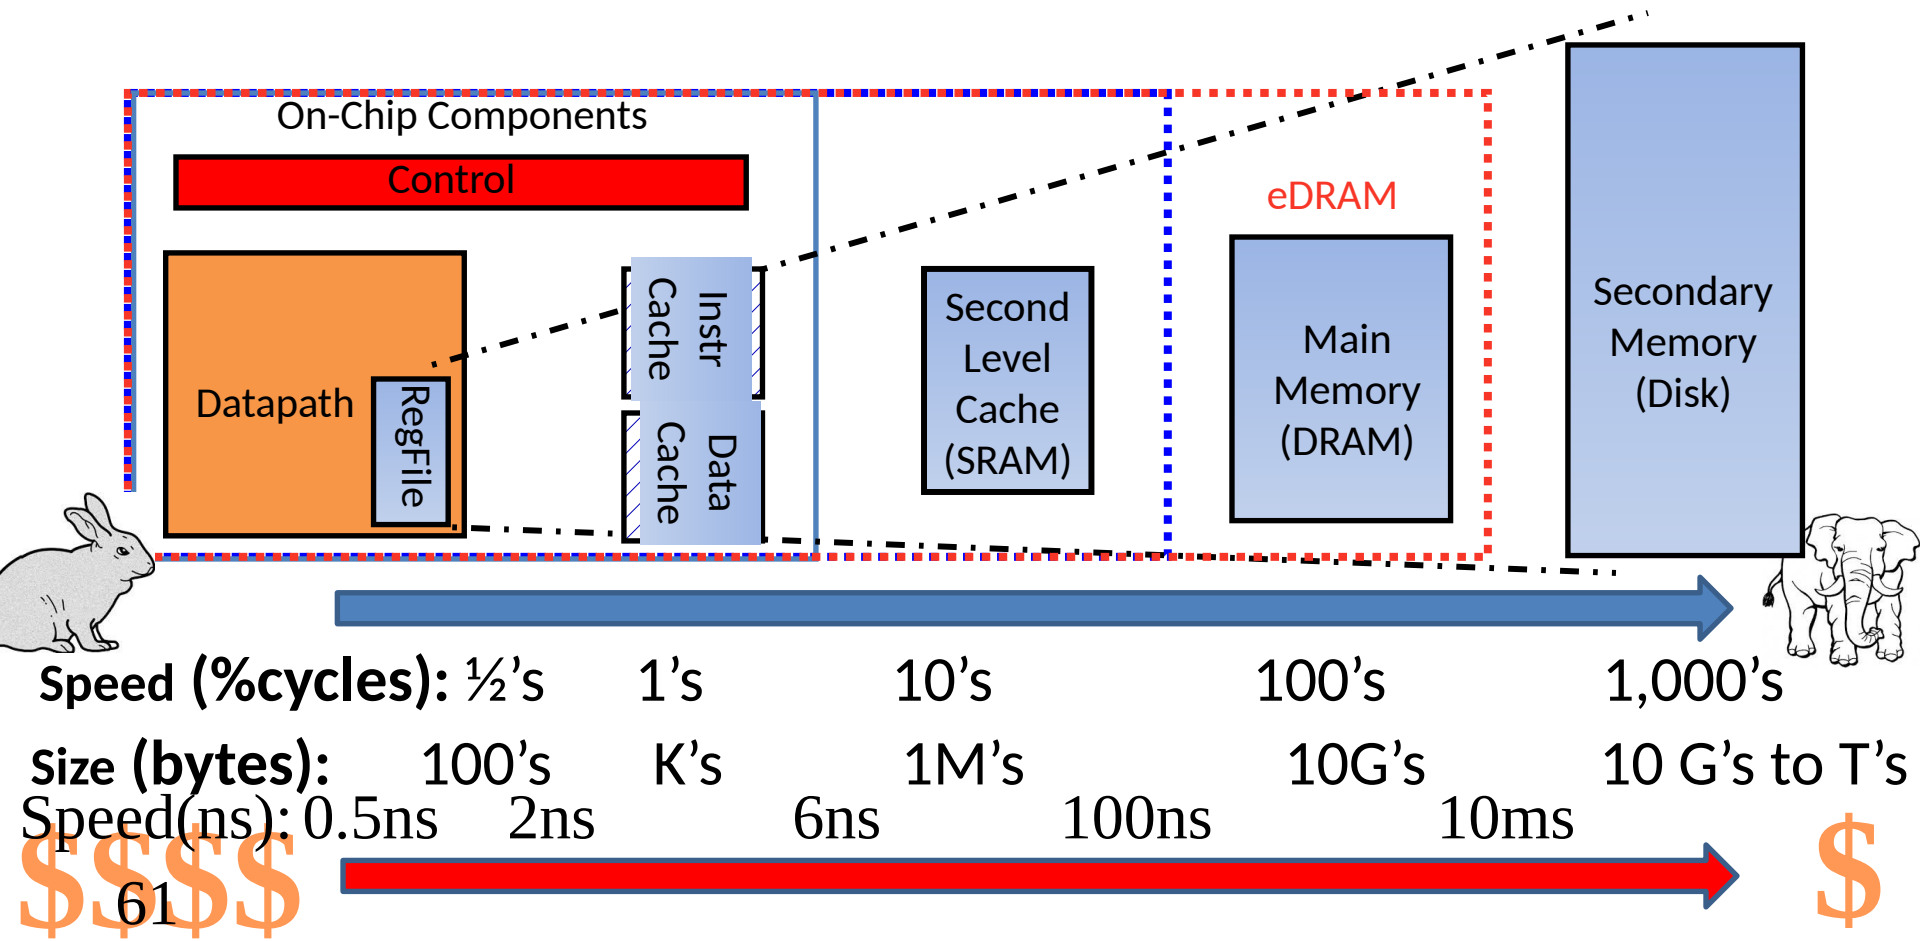
Poor locality example:

We can see that the program accesses a large memory footprint without a specific order, giving the pattern poor spatial locality. We also see little temporal re-use, which also makes this code cache unfriendly.



The Memory Hierarchy

- By taking advantage of the **principle of locality**
 - Can present the user with as much memory as is available in the cheapest technology
 - at the speed offered by the fastest technology



Cache Design

Maurice Wilkes, “Slave Memories and Dynamic Storage Allocation,” IEEE Trans. On Electronic Computers, 1965

Awards: Turing Award,
IEEE John von Neumann Medal, Faraday Medal,
Eckert–Mauchly Award, Mountbatten Medal



Example Memory Hierarchy

↑
Smaller,
faster,
and
costlier
(per byte)
storage
devices

↓
Larger,
slower,
and
cheaper
(per byte)
storage
devices

L0:

Regs

CPU registers hold words
retrieved from the L1 cache.

L1:

L1 cache
(SRAM)

L1 cache holds cache lines
retrieved from the L2 cache.

L2:

L2 cache
(SRAM)

L2 cache holds cache lines
retrieved from L3 cache

L3:

L3 cache
(SRAM)

L3 cache holds cache lines
retrieved from main memory.

L4:

Main memory
(DRAM)

Main memory holds disk
blocks retrieved from
local disks.

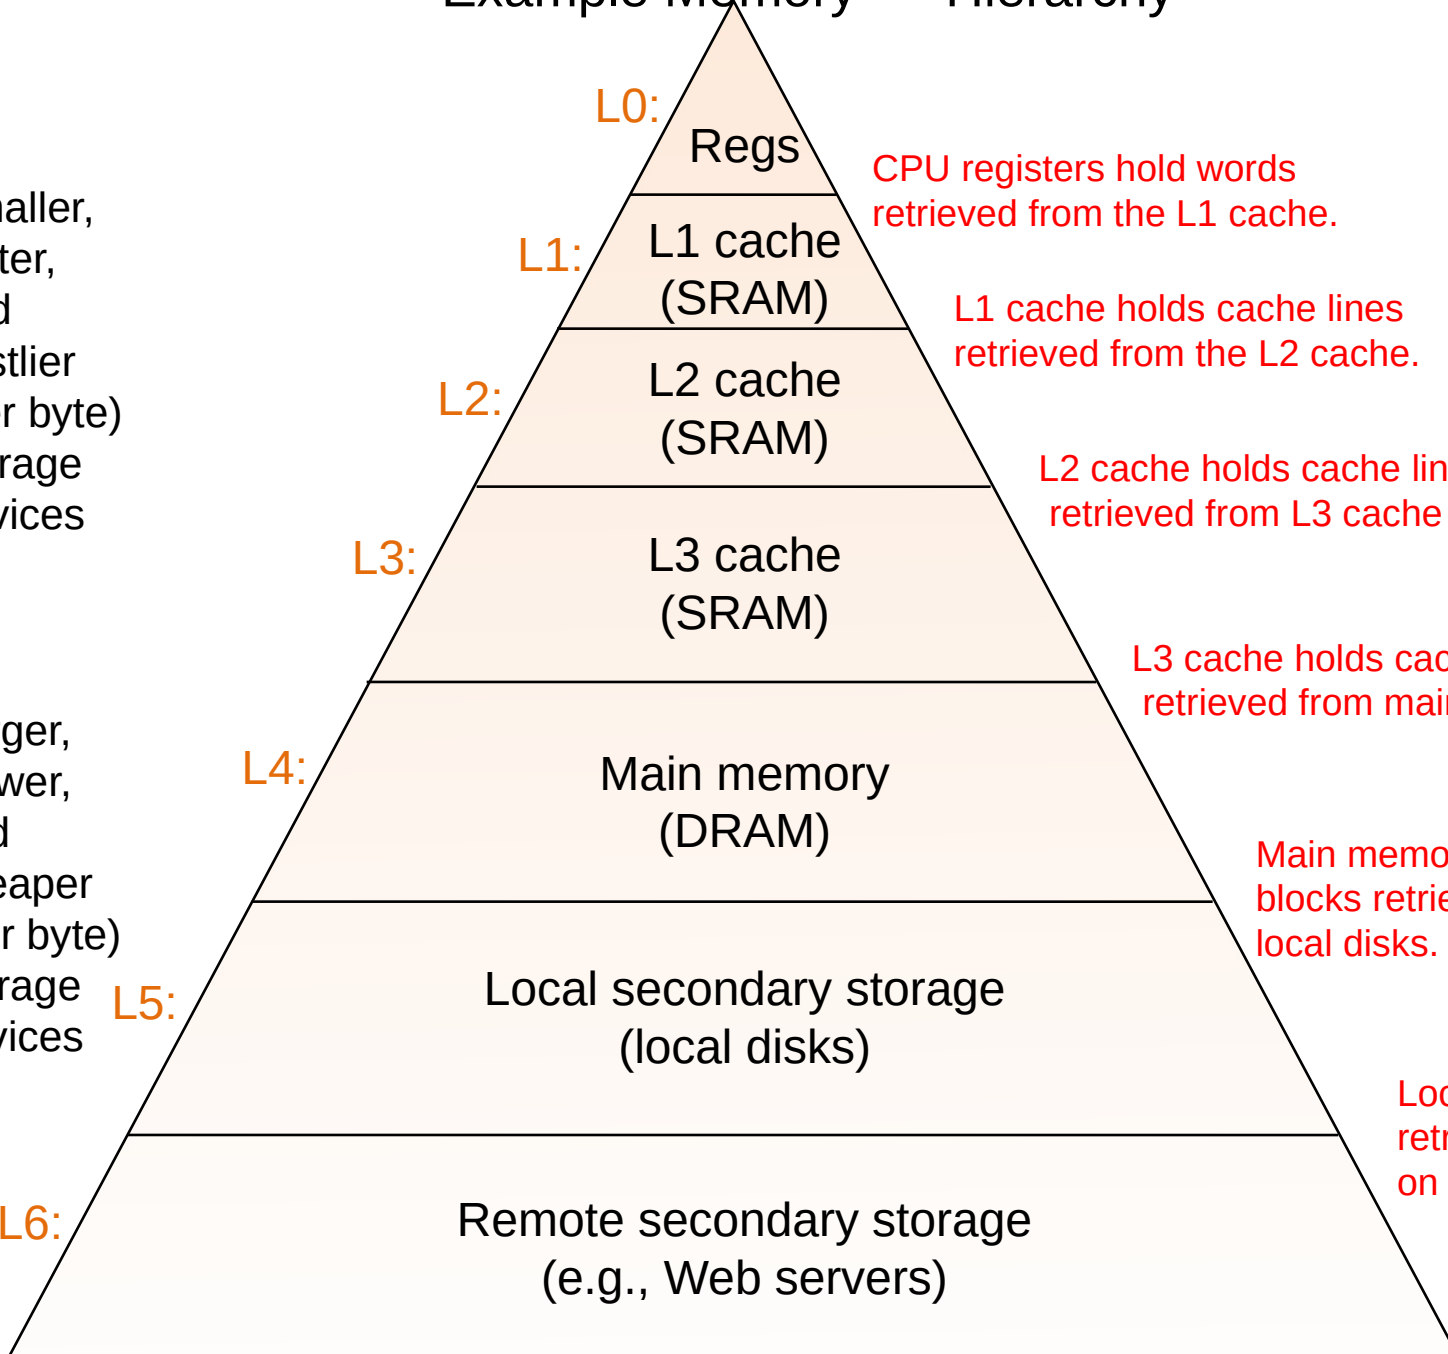
L5:

Local secondary storage
(local disks)

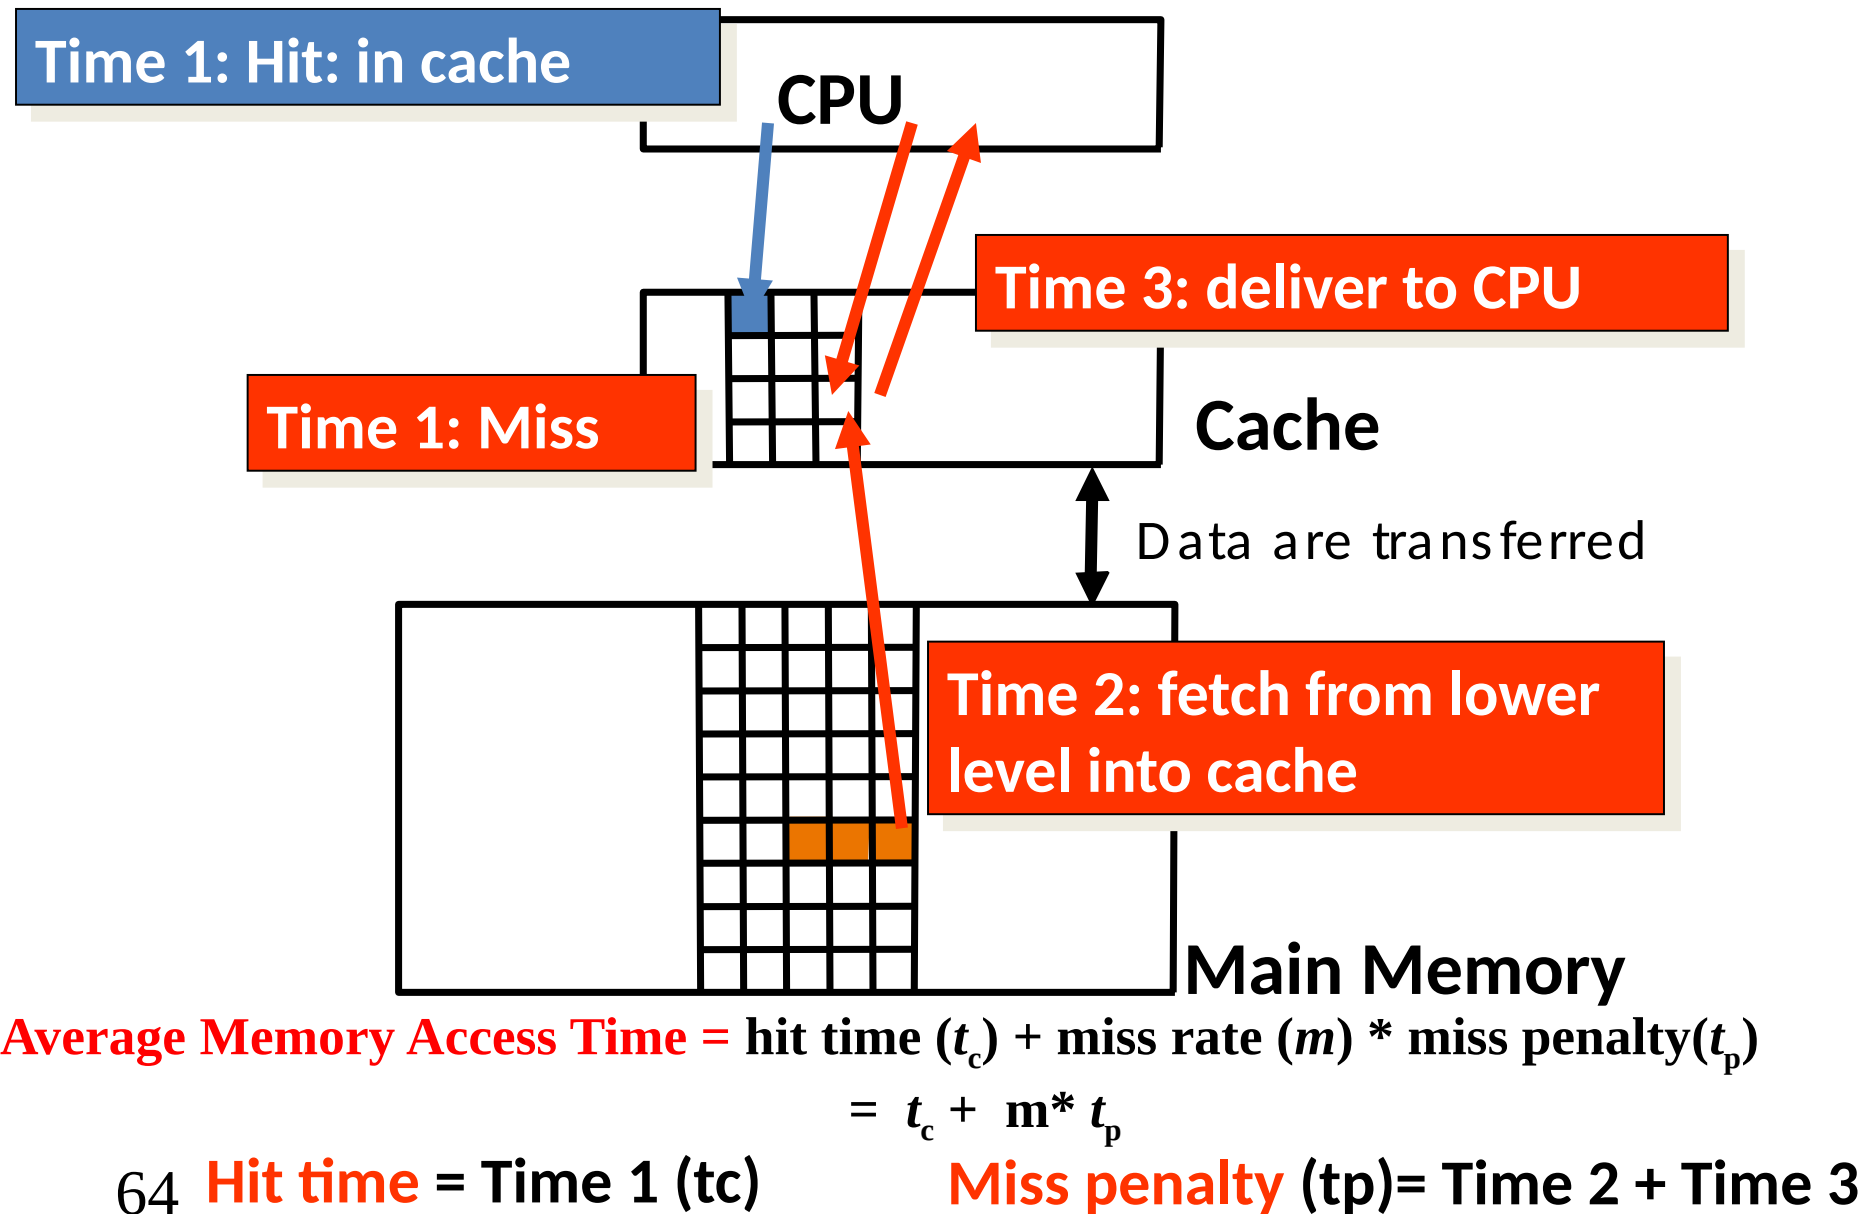
Local disks hold files
retrieved from disks
on remote servers

L6:

Remote secondary storage
(e.g., Web servers)



Cache Example



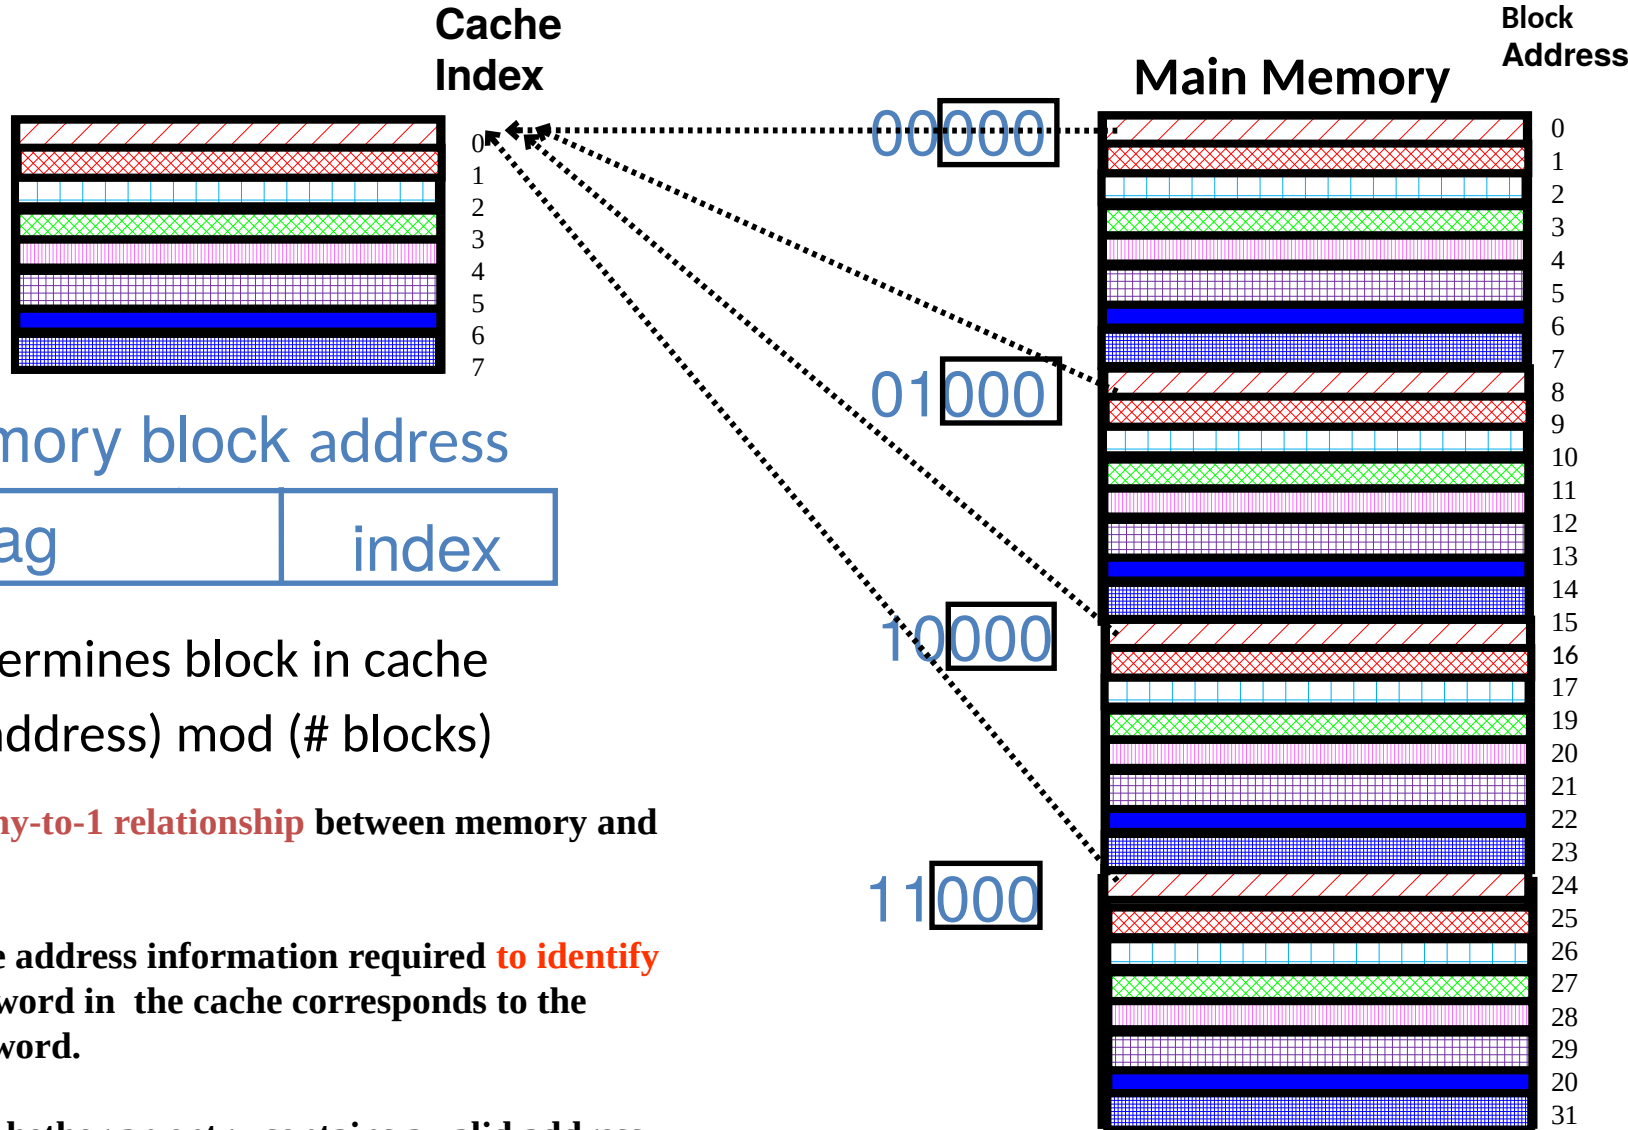
Example: Impact on Performance

- Suppose a processor executes at
 - Clock Rate = 1000 MHz (1 ns per cycle)
 - CPI = 1.1
 - 50% arith/logic, 30% ld/st, 20% control
- Suppose that 10% of memory operations get 50 cycle miss penalty
- $\text{CPI} = \text{ideal CPI} + \text{average stalls per instruction} = 1.1(\text{cyc}) + (0.30 \text{ (datamops/ins)} \times 0.10 \text{ (miss/datamop)} \times 50 \text{ (cycle/miss)})$
 $= 1.1 \text{ cycle} + 0.3 \times 0.1 \times 50 \text{ cycle} = 1.1 \text{ cycle} + 1.5 \text{ cycle} = 2.6$
- 58 % of the time the processor is stalled waiting for memory!
- a 1% instruction miss rate would add an additional 0.5 cycles to the CPI!

Let's think about those numbers

- Huge difference between a hit and a miss
 - Could be 100x, if just L1 and main memory
- Would you believe 99% hits is twice as good as 97%?
 - Consider:
cache hit time of 1 cycle
miss penalty of 100 cycles
 - Average access time:
97% hits: $1 \text{ cycle} + 0.03 * 100 \text{ cycles} = 4 \text{ cycles}$
99% hits: $1 \text{ cycle} + 0.01 * 100 \text{ cycles} = 2 \text{ cycles}$
- This is why “miss rate” is used instead of “hit rate”

Simplest Cache: Direct Mapped



- index determines block in cache
- $\text{index} = (\text{address}) \bmod (\# \text{ blocks})$

There is a **Many-to-1 relationship** between memory and cache

tags

- contain the address information required **to identify** whether a word in the cache corresponds to the requested word.

valid bit

- indicates whether an entry contains a valid address

67

**Cache Table
Entry**

V	tag	Data
---	-----	------

Direct Mapped Cache: Temporal Example

lw \$1,10 110 (\$0)

lw \$2,11 010 (\$0)

lw \$3,10 110 (\$0)

Miss: valid

Miss: valid

Hit!

lw \$1,22(\$0)

lw \$2,26(\$0)

lw \$3,22(\$0)

Index	Valid	Tag	Data
000	N		
001	N		
010	Y	11	Memory[11010]
011	N		
100	N		
101	N		
110	Y	10	Memory[10110]
111	N		

Direct Mapped Cache: Worst case, always miss!

lw \$1,10 110 (\$0)

lw \$2,11 110 (\$0)

lw \$3,00 110 (\$0)

Miss: valid

Miss: tag

Miss: tag

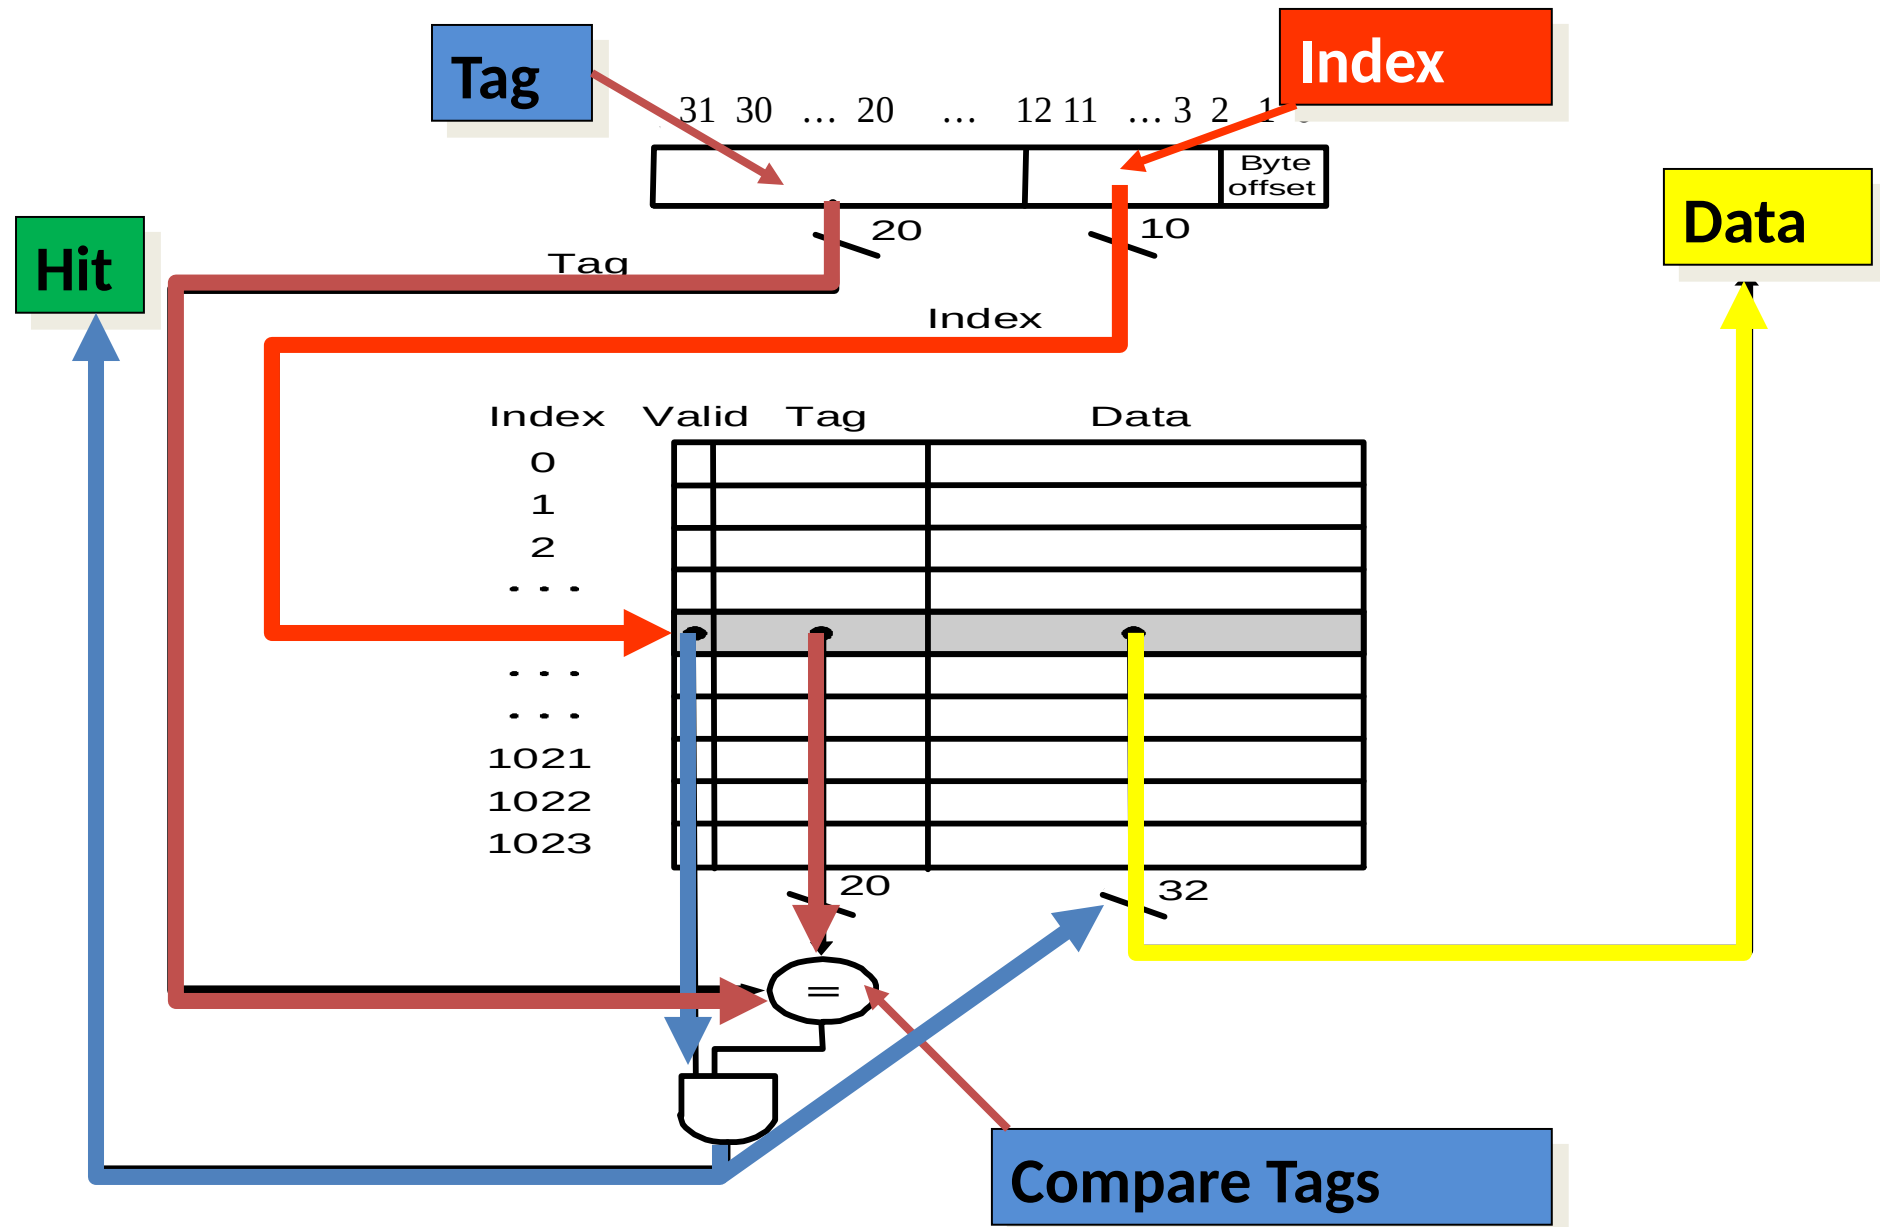
lw \$1,22(\$0)

lw \$2,30(\$0)

lw \$3,6(\$0)

Index	Valid	Tag	Data
000	N		
001	N		
010	N		
011	N		
100	N		
101	N		
110	Y	00	Memory[00110]
111	N		

Direct Mapped Cache: MIPS Example



4KB Direct Mapped Cache

Example

4 Kbytes = Nominal Cache Capacity

$1K = 2^{10} = 1024$ Blocks

Each block = one word

(4 bytes)

Can cache up to


2^{32} bytes = 4 GB

of memory

Mapping function:

Cache Block frame number =
(Block address) MOD (1024)

i.e . Index field or 10 low bits of
block address




Tag

Index

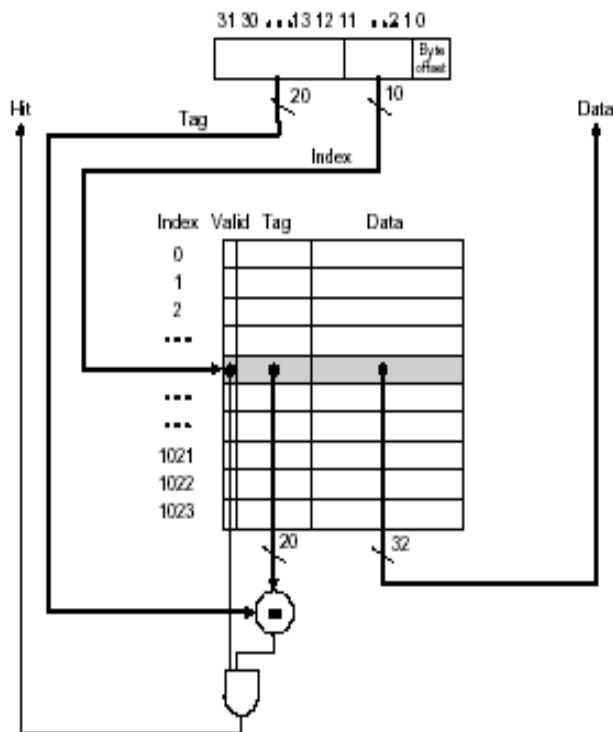
Offset

Mapping

Hit Access Time = SRAM Delay + Hit/Miss Logic Delay



Direct Mapped Cache [contd...]



- What is the size of cache ?

4K

- If I read
0000 0000 0000 0000 0000 0000 1000 0001
- What is the index number checked ?
- If the number was found, what are the inputs to comparator ?

Nominal Capacity

64KB Direct Mapped Cache Example

$4K = 2^{12} = 4096$ blocks

Each block = four words = 16 bytes

Can cache up to 2^{32} bytes = 4 GB of memory

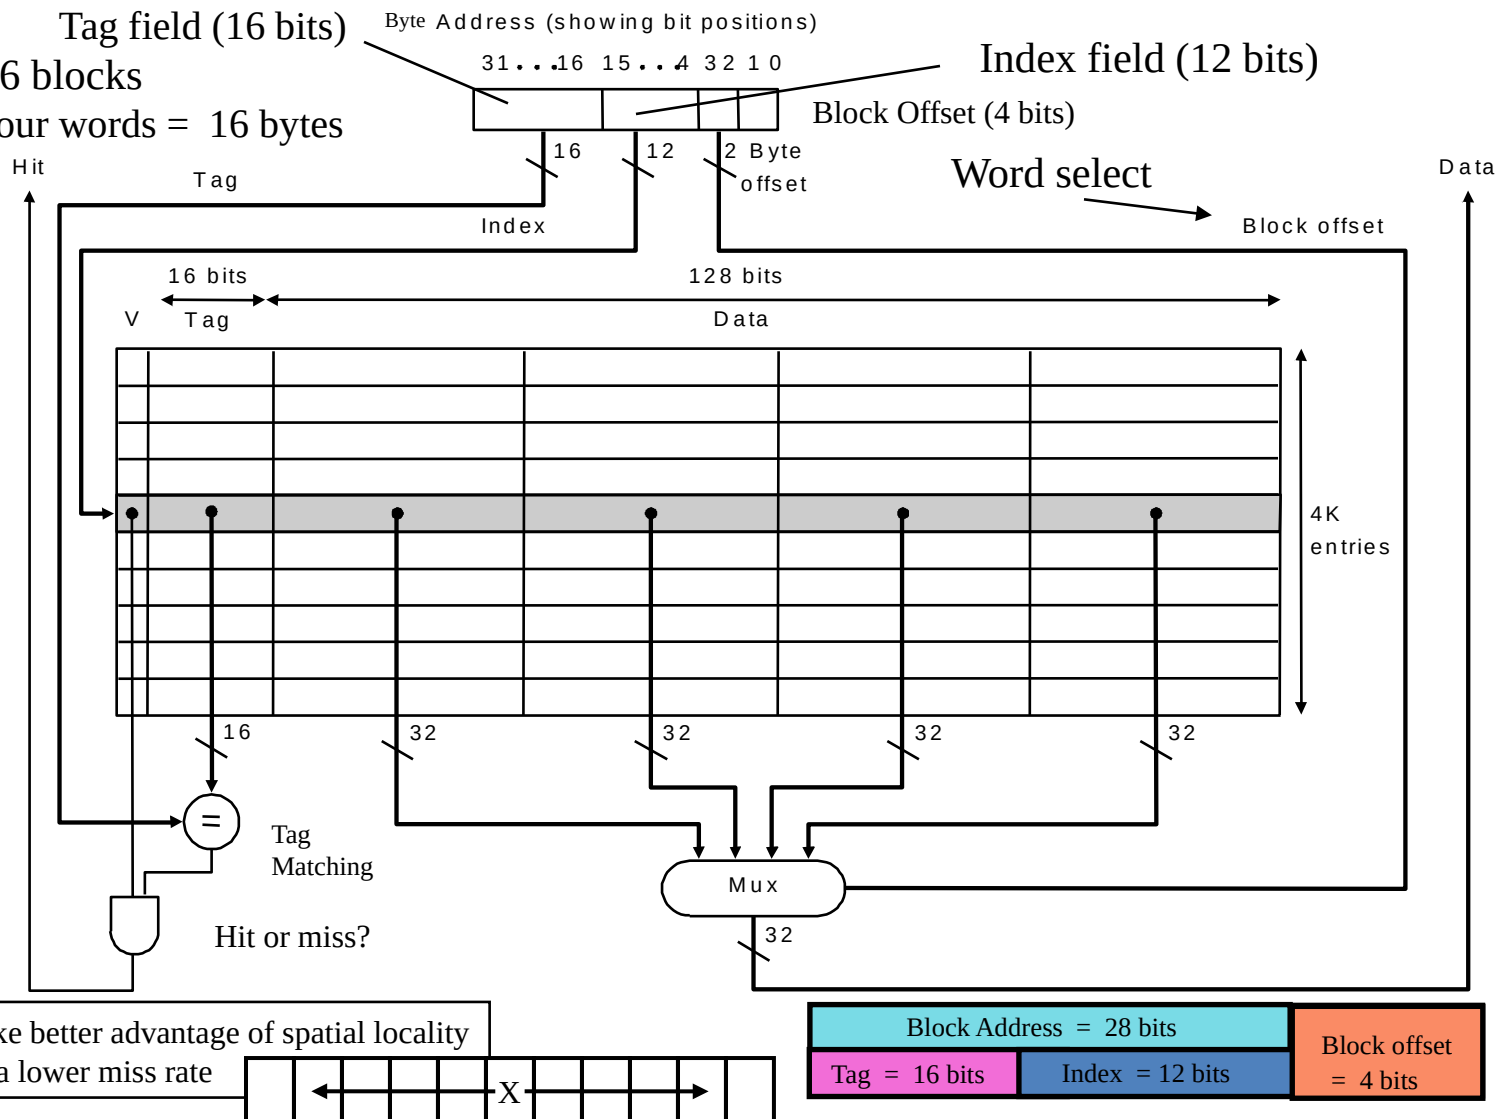
SRAM

Typical cache Block or line size: 32-64 bytes

Larger cache blocks take better advantage of spatial locality and thus may result in a lower miss rate

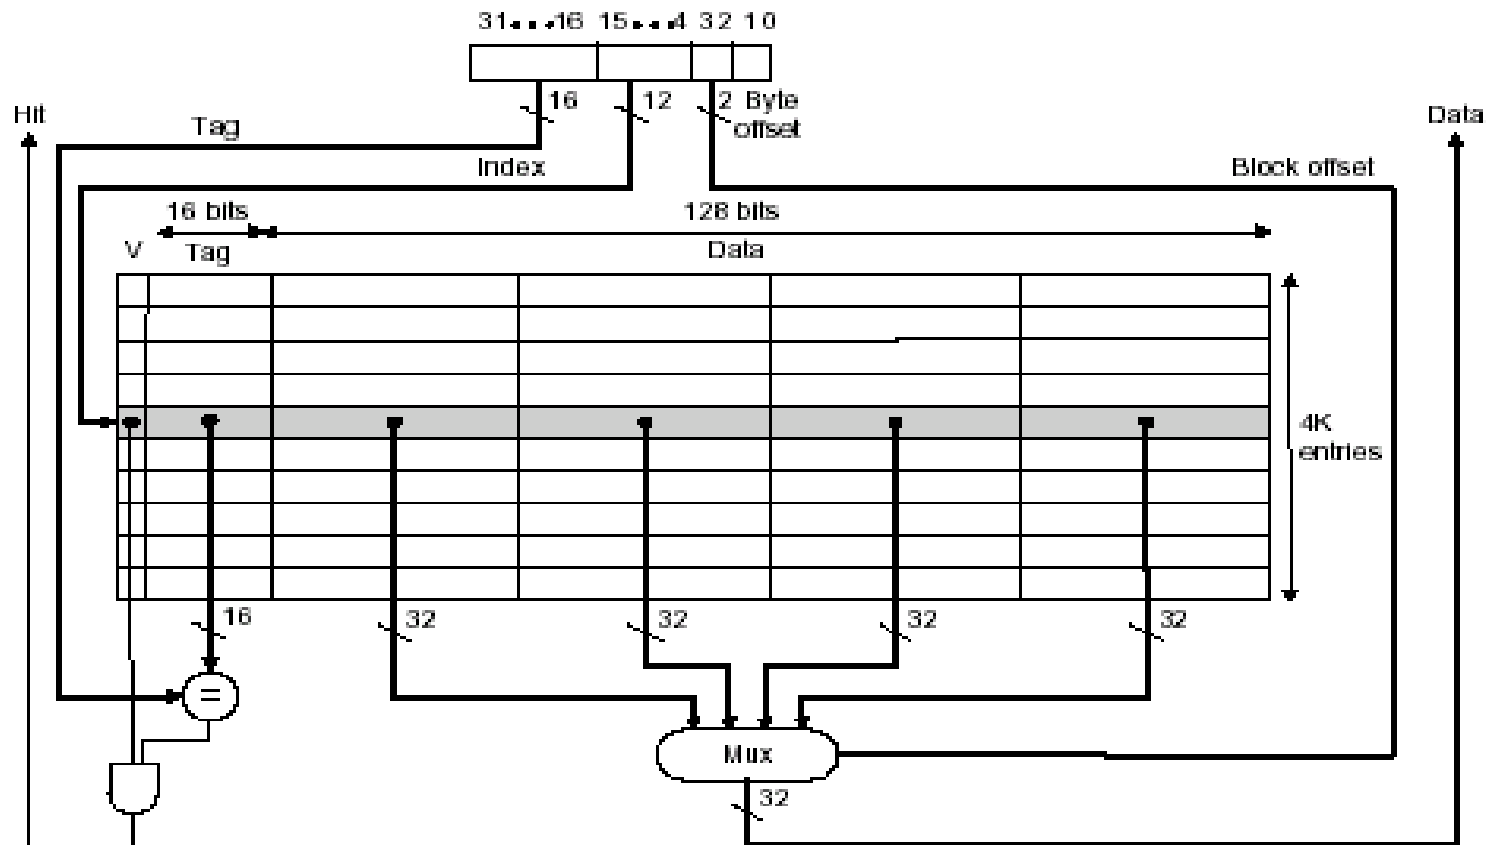
Mapping Function: Cache Block frame number = (Block address) MOD (4096)
i.e. index field or 12 low bit of block address

Hit Access Time = SRAM Delay + Hit/Miss Logic Delay



Direct Mapped Cache [contd...]

Taking advantage of spatial locality, we read 4 bytes at a time



Caching Terminology

- **block/line** : the unit of information that can be exchanged between two adjacent levels in the memory hierarchy
- **cache hit**: when CPU finds a requested data item in the cache
- **cache miss**: when CPU does not find a requested data item in the cache
- **miss rate (m)** : fraction of cache accesses that result in a miss
- **hit rate (h)** : $(1 - \text{miss rate})$ fraction of memory access found in the faster level (cache)
- **hit time (t_c)** : time required to access the level 1 (cache)
- **miss penalty (t_p)** : time required to replace a data unit in the faster level (cache) + time involved to deliver the requested data to the processor from the cache. (hit time \ll miss penalty)
- **average memory access time (amat)** = hit time (t_c) + miss rate (m) * miss penalty (t_p)
$$\text{amat} = t_c + (1-h) * t_p$$

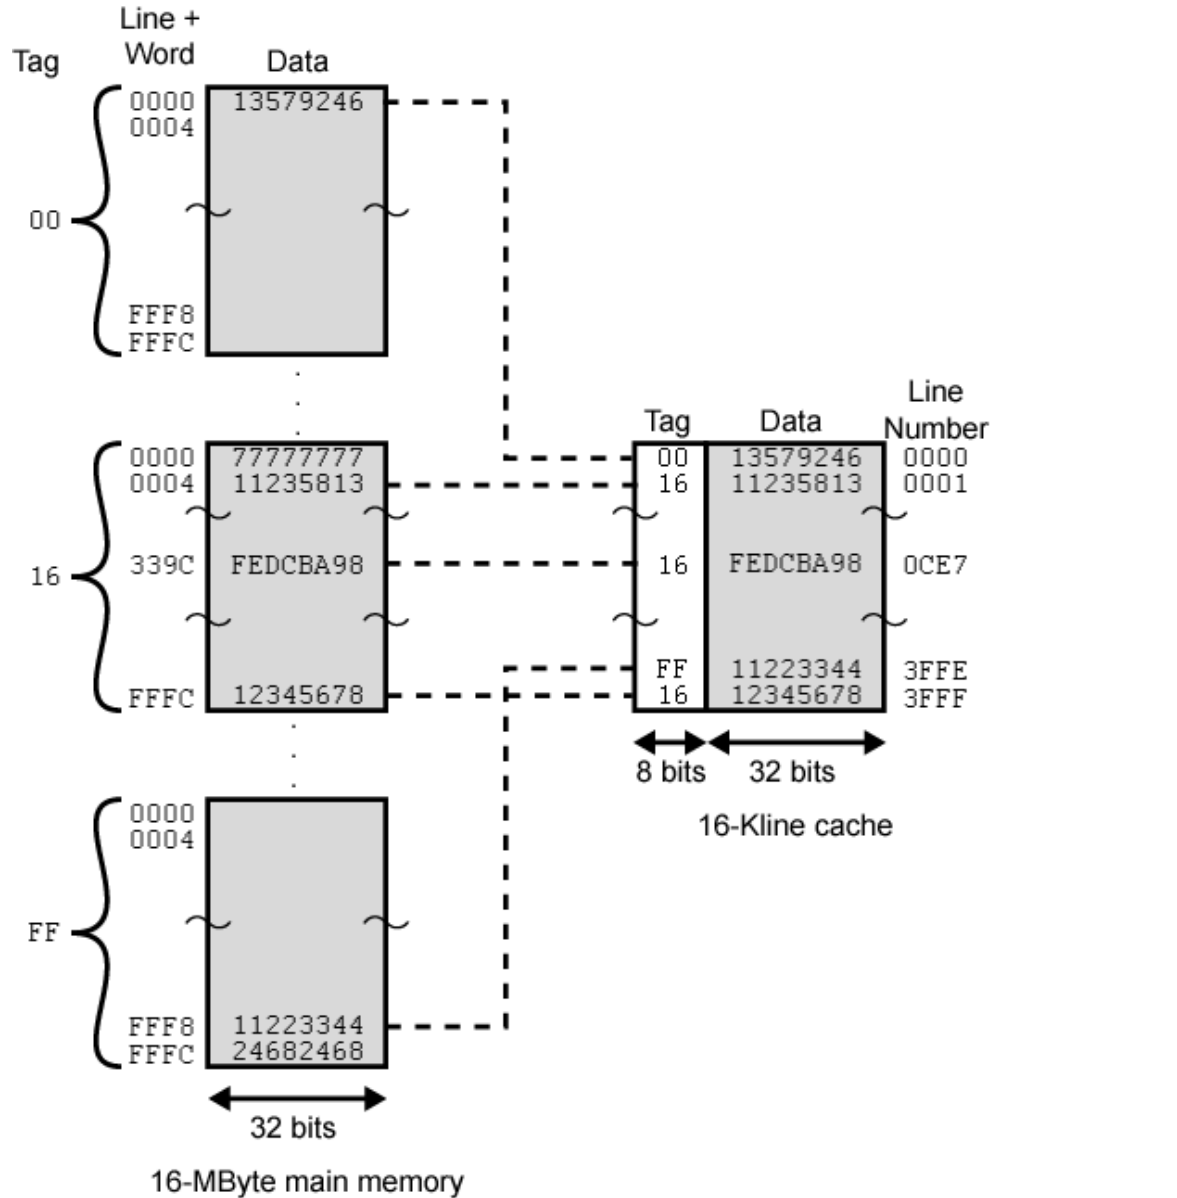
Cache Misses

- **Compulsory** (cold start or process migration, first reference): first access to a block
 - “Cold” fact of life: not a whole lot you can do about it
 - Note: If you are going to run “billions” of instruction, Compulsory Misses are insignificant
- **Capacity**:
 - Cache cannot contain all blocks access by the program
 - Solution: increase cache size
- **Conflict** (collision):
 - Multiple memory locations mapped to the same cache location
 - Solution 1: increase cache size
 - Solution 2: increase associativity
- **Coherence** (Invalidation): other process (e.g., I/O) updates memory

Direct Mapped Cache [contd...]

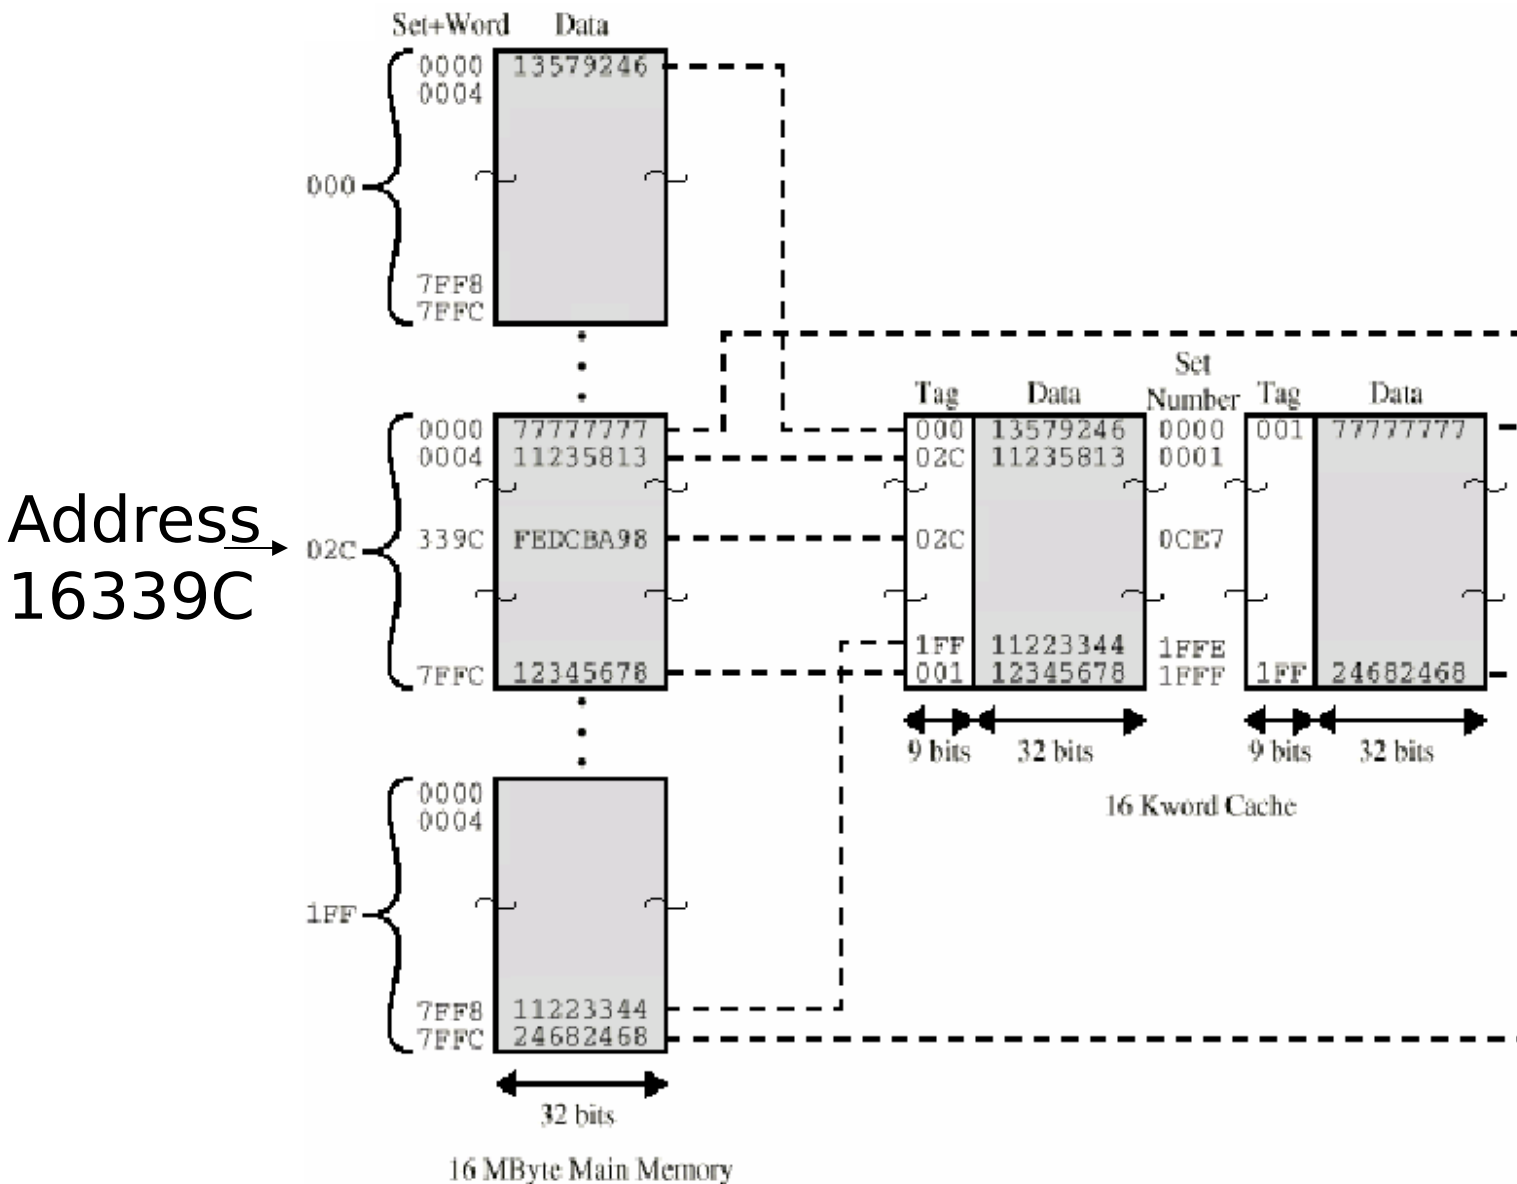
- Advantage
 - Simple
 - Fast
- Disadvantage
 - Mapping is fixed !!!

Direct Mapping Example



	Tag	Line	Word
Main memory address =	8	14	2

Two Way Set Associative Example




4K Four-Way Set Associative Cache: MIPS Implementation Example

Nominal Capacity

1024 block frames
Each block = one word
4-way set associative
 $1024 / 4 = 2^8 = 256$ sets

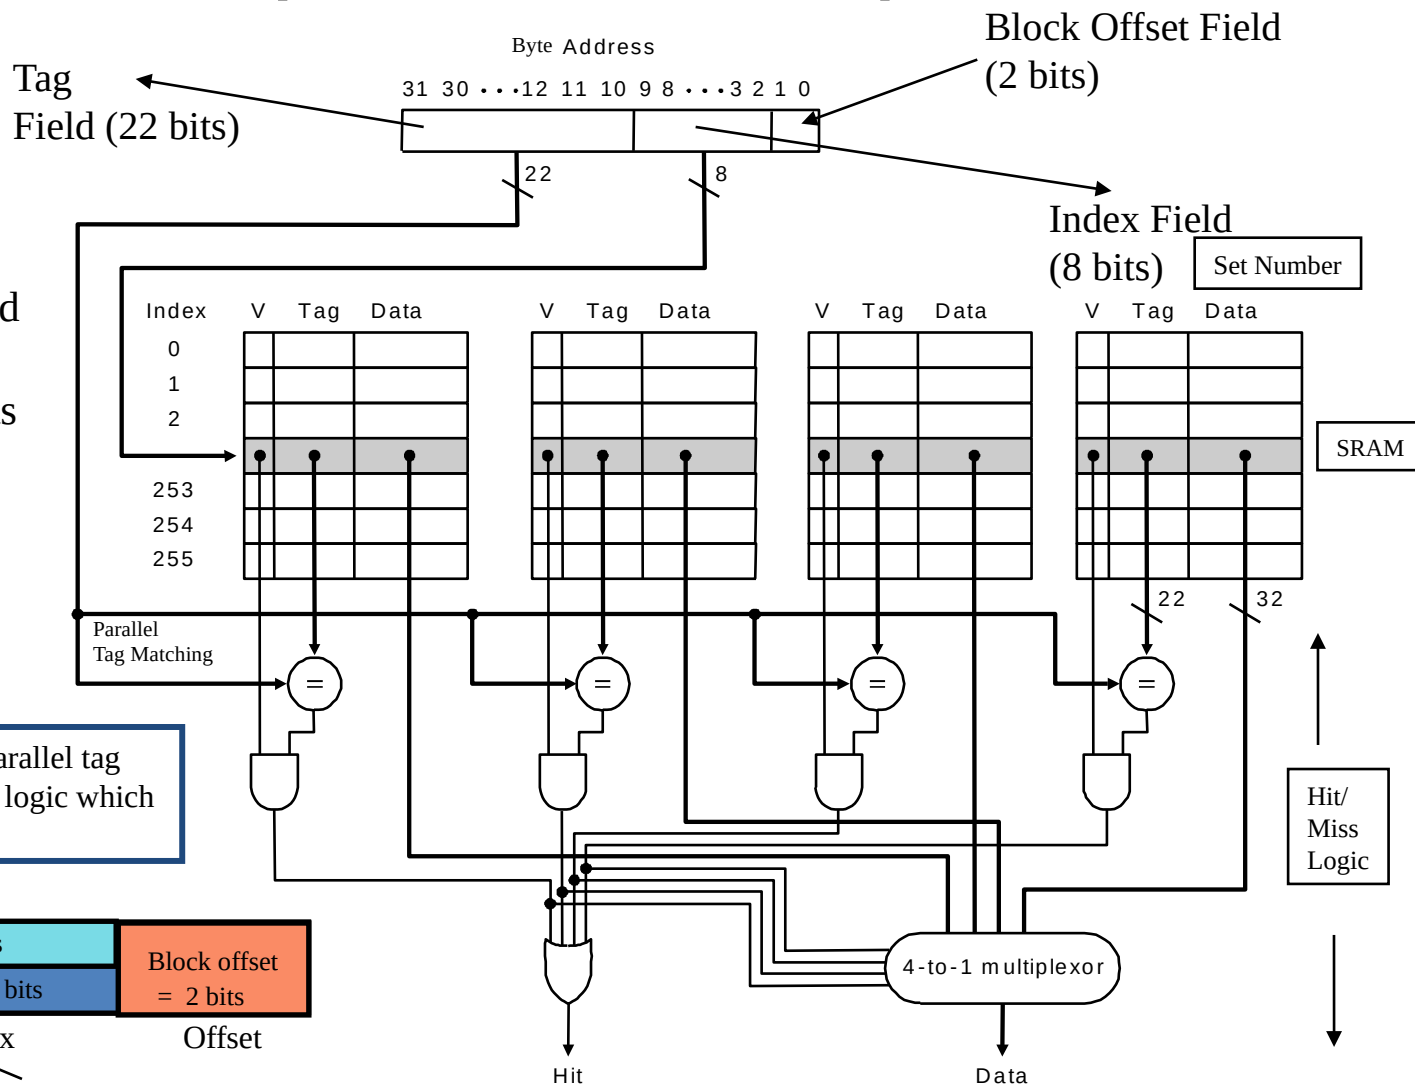
Can cache up to
 2^{32} bytes = 4 GB
of memory

Set associative cache requires parallel tag matching and more complex hit logic which may increase hit time



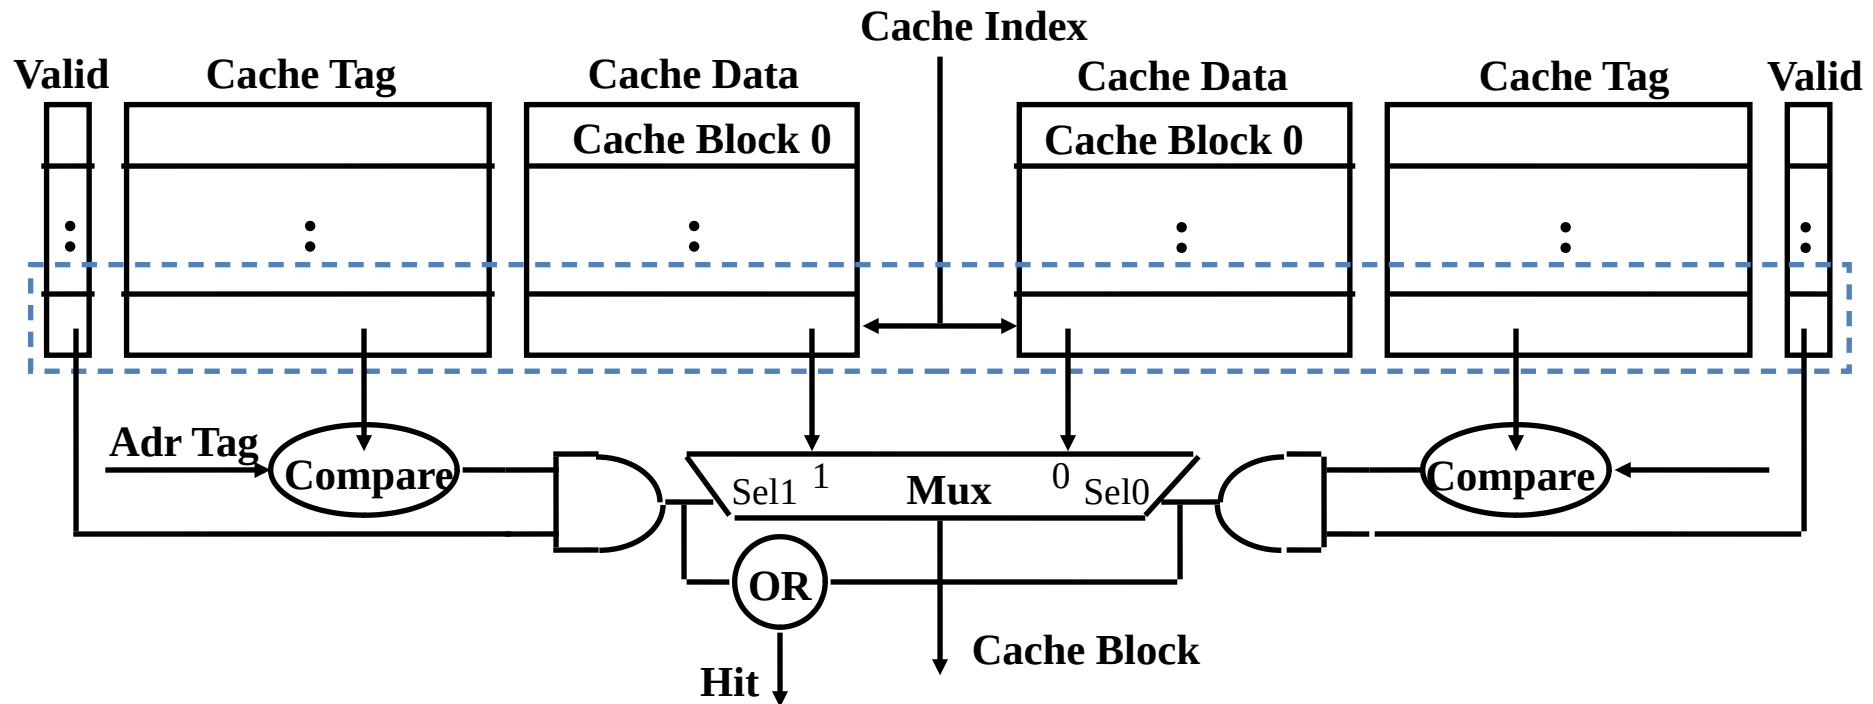
Mapping Function: Cache Set Number = index = (Block address) MOD (256)

Hit Access Time = SRAM Delay + Hit/Miss Logic Delay



Set Associative Cache

- **N-way set associative**: N entries for each Cache Index
 - N direct mapped caches operates in parallel
- Example: Two-way set associative cache
 - Cache Index selects a “set” from the cache
 - The two tags in the set are compared to the input in parallel
 - Data is selected based on the tag result



Problem 1

- (a) A 64KB, direct mapped cache has 16 byte blocks. If addresses are 32 bits, how many bits are used the tag, index, and offset in this cache?
- (b) How would the address be divided if the cache were 4-way set associative instead?
- (c) How many bits is the index for a fully associative cache.

Problem 2

An ISA has 44 bit addresses with each addressable item being a byte. Design a four-way set associative cache with each of the four lines in a set containing 64 bytes. Assume that you have 256 sets in the cache.

Calculating Number of Cache Bits Needed



Address Fields



Cache Block Frame (or just cache block)

- How many total bits are needed for a direct-mapped cache with 64 KBytes of data and one word blocks, assuming a 32-bit address?
 - 64 Kbytes = 16 K words = 2^{14} words = 2^{14} blocks
 - Block size = 4 bytes \Rightarrow offset size = $\log_2(4) = 2$ bits,
 - #sets = #blocks = 2^{14} \Rightarrow index size = 14 bits
 - Tag size = address size - index size - offset size = $32 - 14 - 2 = 16$ bits
 - Bits/block = data bits + tag bits + valid bit = $32 + 16 + 1 = 49$
 - Bits in cache = #blocks x bits/block = $2^{14} \times 49 = 98$ Kbytes
- How many total bits would be needed for a 4-way set associative cache to store the same amount of data?
 - Block size and #blocks does not change.
 - #sets = #blocks/4 = $(2^{14})/4 = 2^{12}$ \Rightarrow index size = 12 bits
 - Tag size = address size - index size - offset = $32 - 12 - 2 = 18$ bits
 - Bits/block = data bits + tag bits + valid bit = $32 + 18 + 1 = 51$
 - Bits in cache = #blocks x bits/block = $2^{14} \times 51 = 102$ Kbytes
- Increase associativity \Rightarrow increase bits in cache

i.e nominal cache Capacity = 64 KB

Number of cache block frames

Actual number of bits in a cache block frame

Word = 4 bytes

More bits in tag

1 k = 1024 = 2^{10}

Calculating Cache Bits Needed



Address Fields



Cache Block Frame (or just cache block)

Nominal size

- How many total bits are needed for a direct-mapped cache with 64 KBytes of data and 8 word (32 byte) blocks, assuming a 32-bit address (it can cache 2^{32} bytes in memory)?

– 64 Kbytes = 2^{14} words = $(2^{14})/8 = 2^{11}$ blocks

Number of cache block frames

– block size = 32 bytes

=> offset size = block offset + byte offset = $\log_2(32) = 5$ bits,

– #sets = #blocks = 2^{11} => index size = 11 bits

– tag size = address size - index size - offset size = $32 - 11 - 5 = 16$ bits

– bits/block = data bits + tag bits + valid bit = $8 \times 32 + 16 + 1 = 273$ bits

– bits in cache = #blocks x bits/block = $2^{11} \times 273 = 68.25$ Kbytes

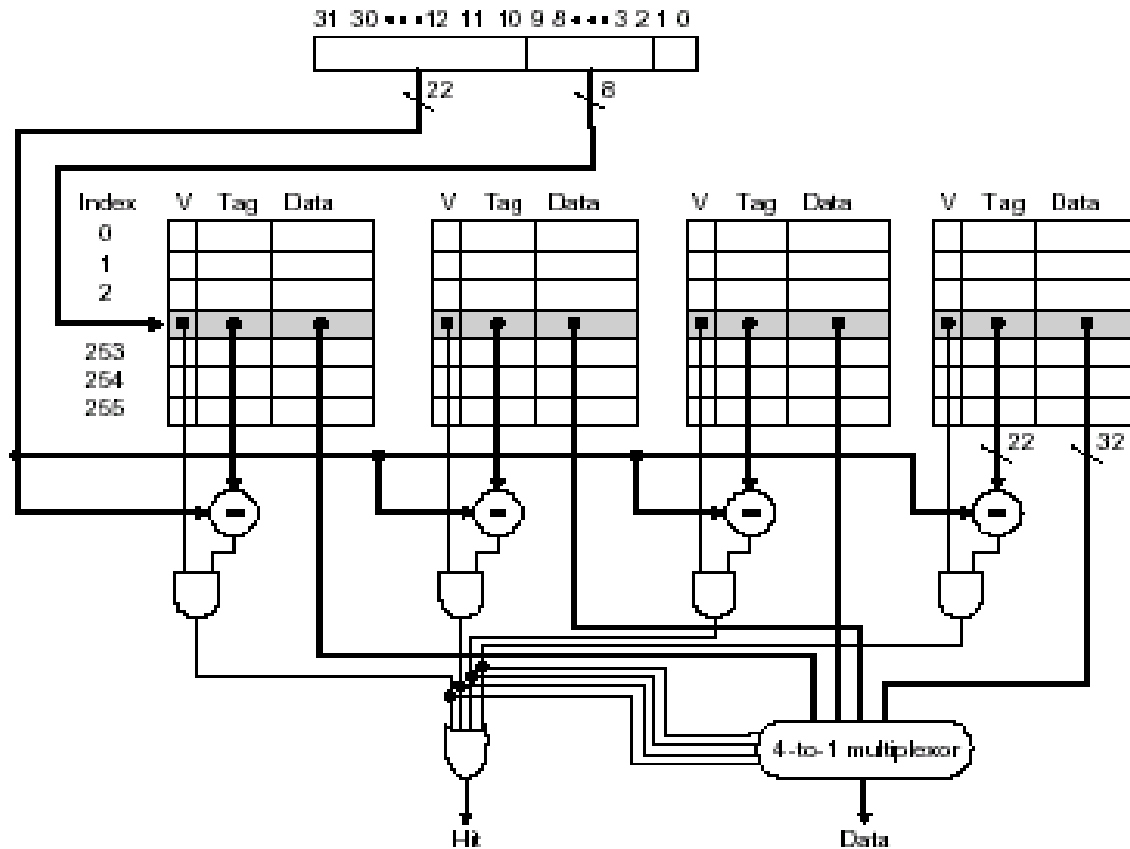
Actual number of bits in a cache block frame

- In Fewer cache block frames thus fewer tags/valid bits e.

Word = 4 bytes

1 k = $1024 = 2^{10}$

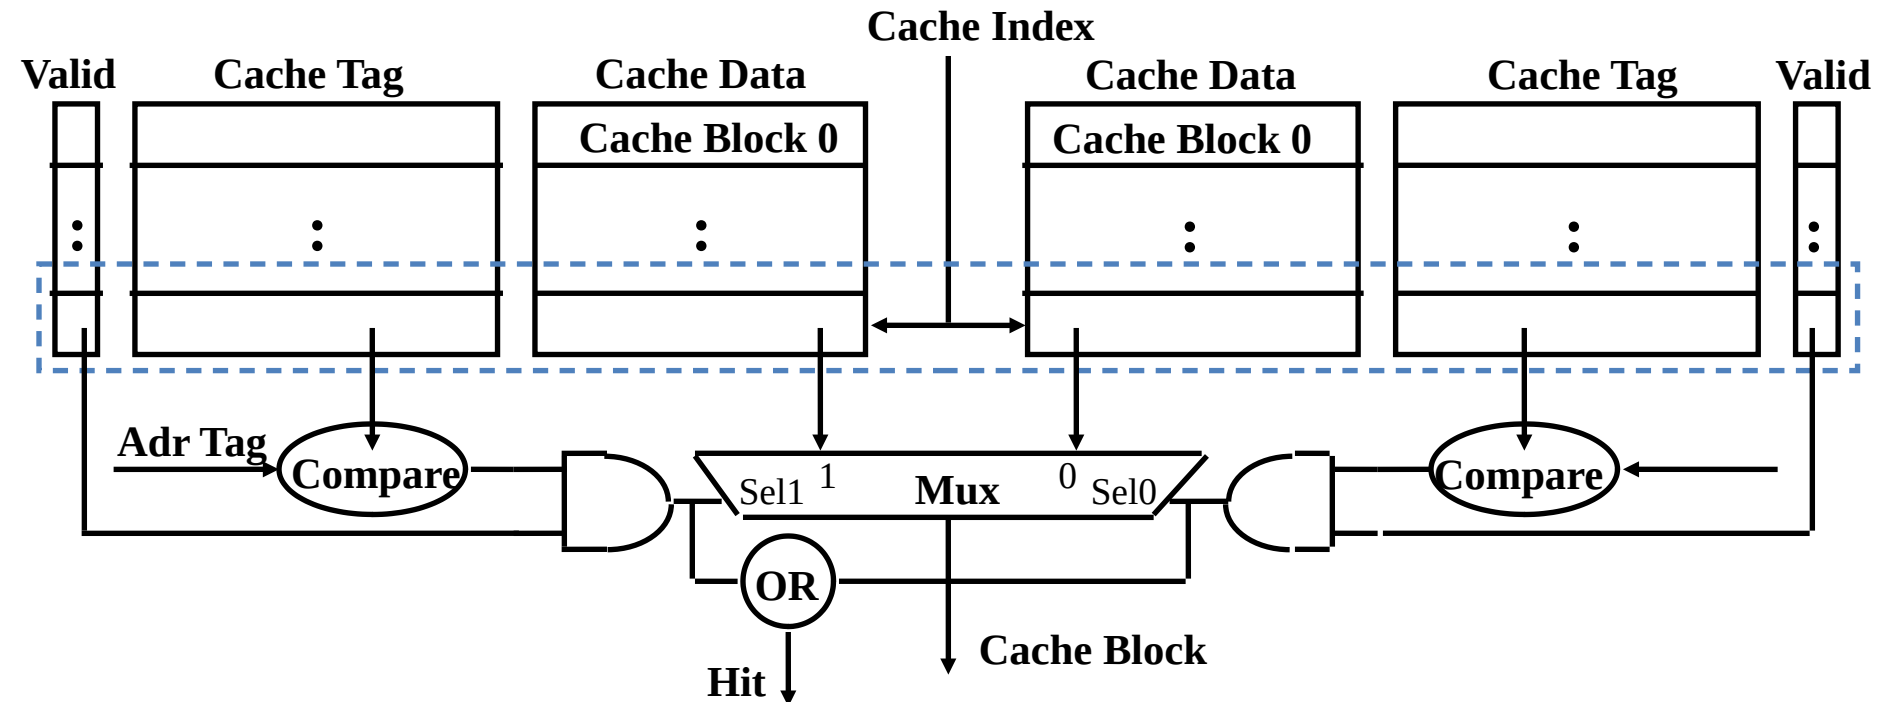
Example: 4-way set associative Cache



What is the cache size in this case ?

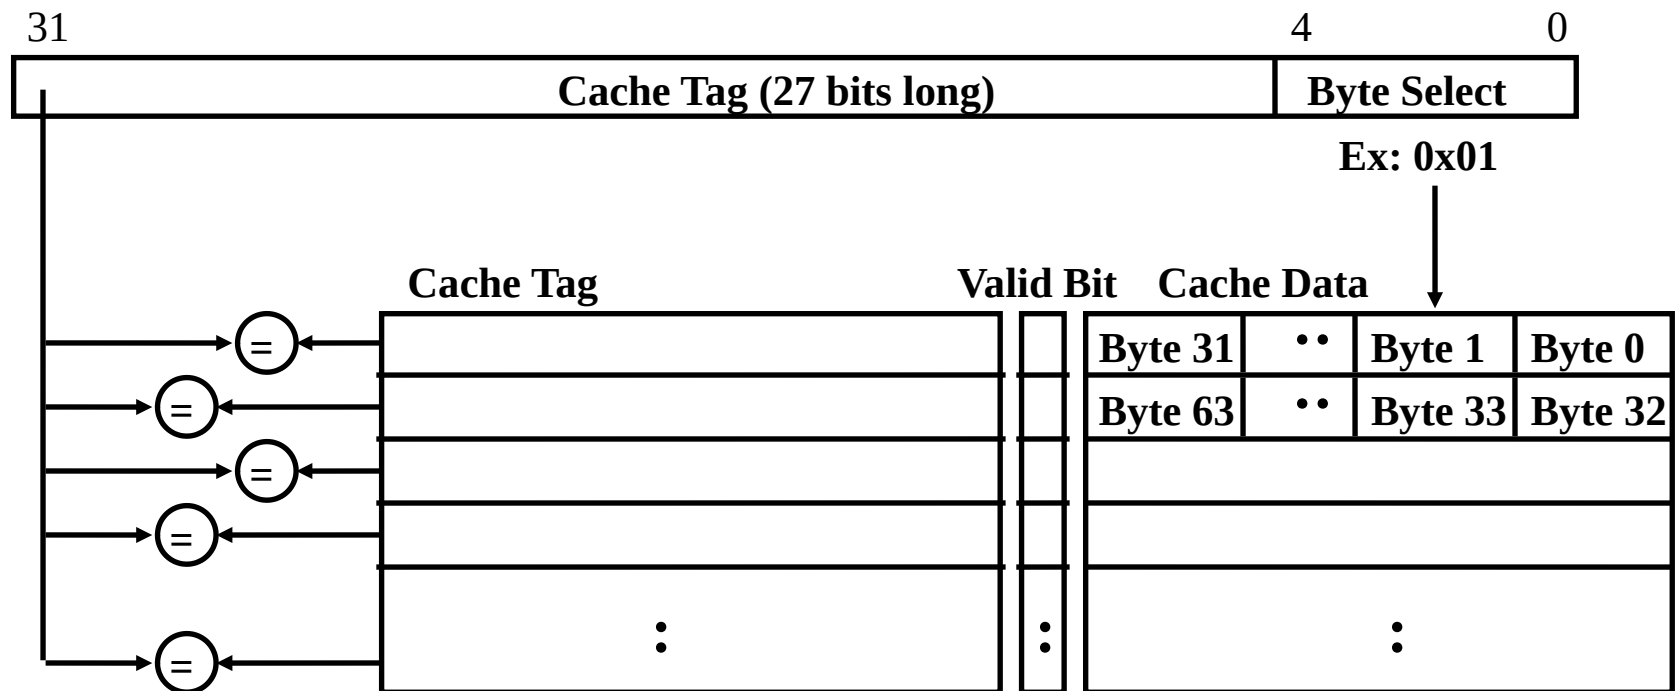
Disadvantages of Set Associative Cache

- N-way Set Associative Cache versus Direct Mapped Cache:
 - N comparators vs. 1
 - Extra MUX delay for the data
 - Data comes **AFTER** Hit/Miss decision and set selection
- In a direct mapped cache, Cache Block is available **BEFORE** Hit/Miss:

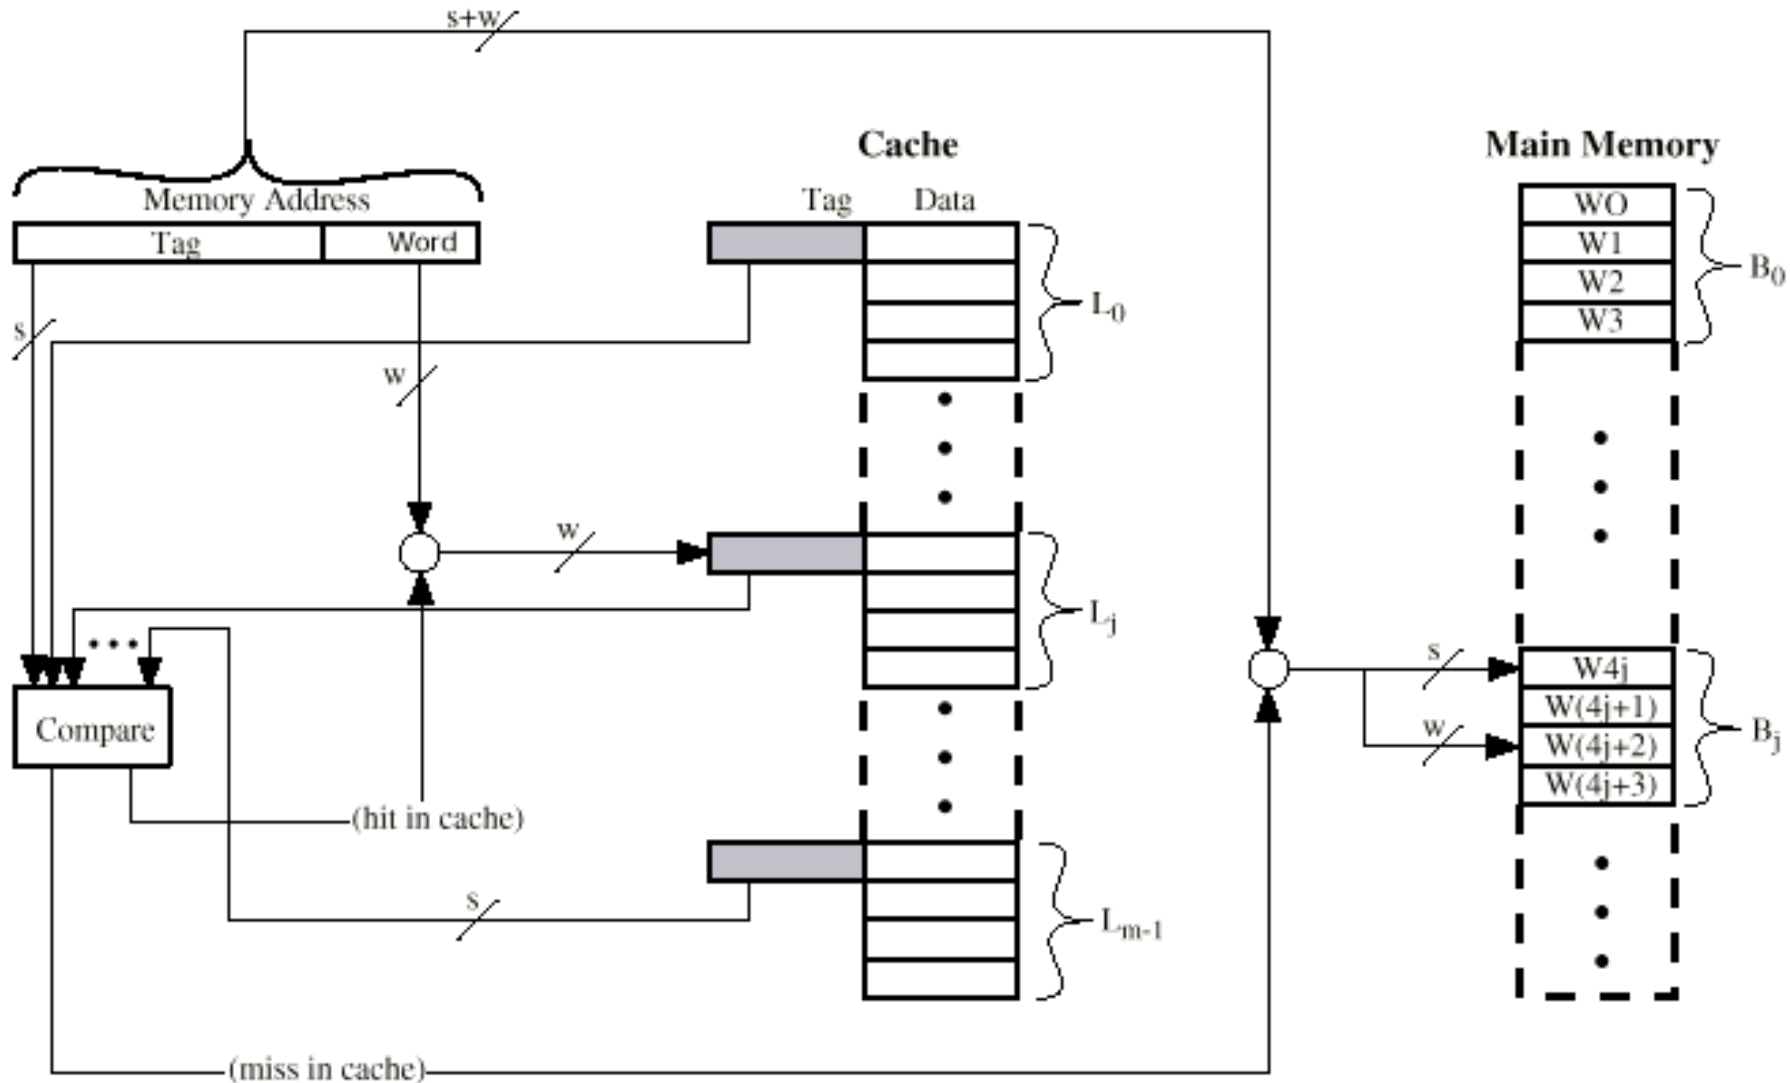


Fully Associative Cache

- Fully Associative Cache
 - Forget about the Cache Index
 - Compare the Cache Tags of all cache entries in parallel
 - Example: Block Size = 32 B blocks, we need N 27-bit comparators
- By definition: Conflict Miss = 0 for a fully associative cache



Fully Associative Cache Organization



Cache Organization:



Cache Block Frame

Set Associative Cache

Why set associative?

One-way set associative
(direct mapped)

1-way set associative:
(direct mapped)
1 block frame per set

Block	Tag	Data
0		
1		
2		
3		
4		
5		
6		
7		

Set associative cache reduces cache misses by reducing conflicts between blocks that would have been mapped to the same cache block frame in the case of direct mapped cache

Two-way set associative

Set	Tag	Data	Tag	Data
0				
1				
2				
3				

2-way set associative:
2 blocks frames per set

Four-way set associative

Set	Tag	Data	Tag	Data	Tag	Data	Tag	Data
0								
1								

4-way set associative:
4 blocks frames per set

8-way set associative:
8 blocks frames per set
In this case it becomes fully associative
since total number of block frames = 8

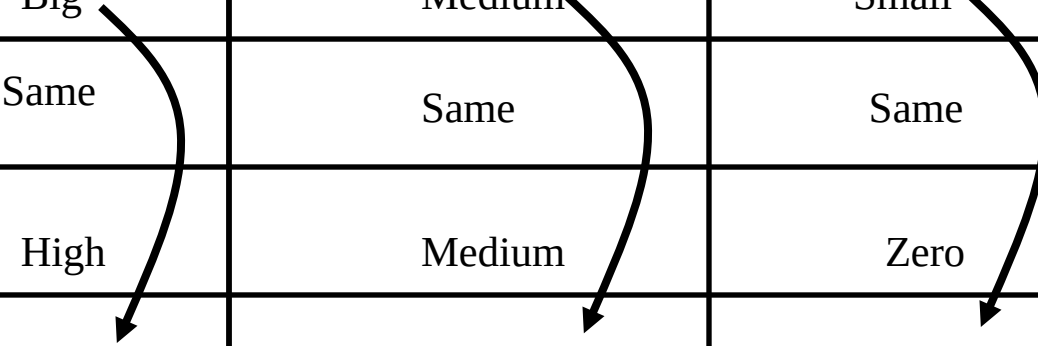
Eight-way set associative (fully associative)

Tag	Data	Tag	Data	Tag	Data	Tag	Data	Tag	Data	Tag	Data	Tag	Data	Tag	Data

A cache with a total of 8 cache block frames shown

Design Options at Constant Cost

	Direct Mapped	N-way Set Associative	Fully Associative
Cache Size	Big	Medium	Small
Compulsory Miss	Same	Same	Same
Conflict Miss	High	Medium	Zero
Capacity Miss	Low	Medium	High
Coherence Miss	Same	Same	Same



Four Questions for Cache Design

- Q1: Where can a block be placed in the upper level?
(Block placement)
- Q2: How is a block found if it is in the upper level?
(Block identification)
- Q3: Which block should be replaced on a miss?
(Block replacement)
- Q4: What happens on a write?
(Write strategy)

Where can a block be placed in the upper level?

- Direct Mapped
- Set Associative
- Fully Associative

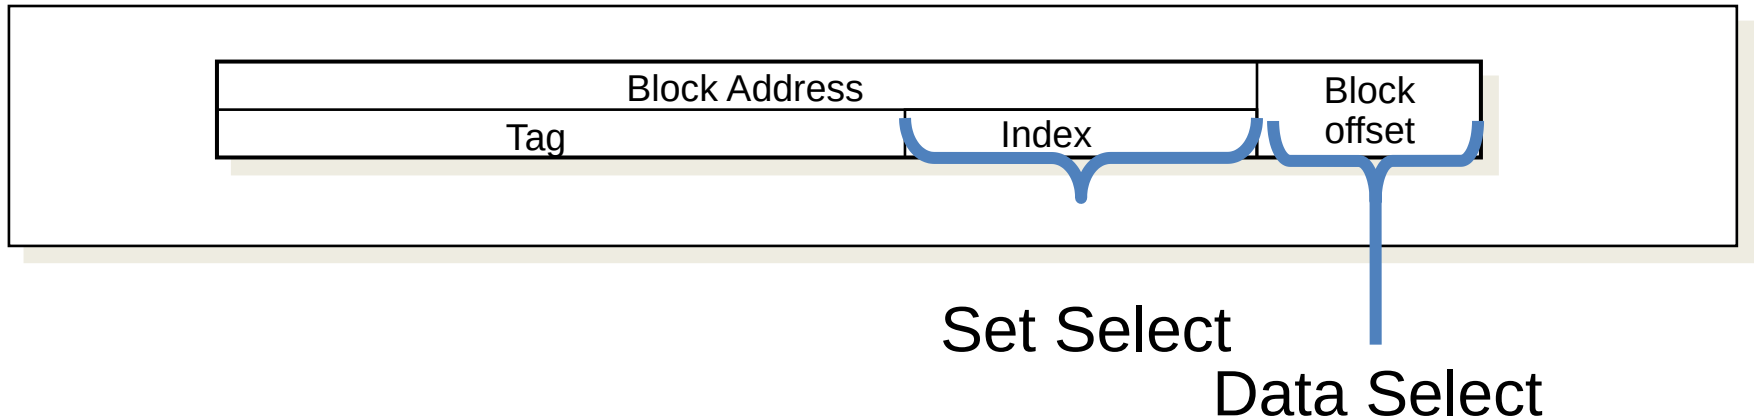
(direct mapped)		
Block	Tag	Data
0		
1		
2		
3		
4		
5		
6		
7		

Set	Tag	Data	Tag	Data
0				
1				
2				
3				

Four-way set associative								
Set	Tag	Data	Tag	Data	Tag	Data	Tag	Data
0								
1								

[illegible]

How is a block found if it is in the upper level?



- Direct indexing (using index and block offset), tag compares, or combination
- Increasing associativity shrinks index, expands tag


Which block should be replaced on a miss?

- Easy for Direct Mapped
- Set Associative or Fully Associative:
 - Random
 - LRU (Least Recently Used)

Associativity:		2-way		4-way		8-way	
Size	LRU	Random	LRU	Random	LRU	Random	
16 KB	5.2%	5.7%	4.7%	5.3%	4.4%	5.0%	
64 KB	1.9%	2.0%	1.5%	1.7%	1.4%	1.5%	
256 KB	1.15%	1.17%	1.13%	1.13%	1.12%	1.12%	

What happens on a write?

- Write through—The information is written to both the block in the cache and to the block in the lower-level memory.
- Write back—The information is written only to the block in the cache. The modified cache block is written to main memory only when it is replaced.
 - is block clean or dirty?
- WT always combined with write buffers so that don't wait for lower level memory



Intel Cache Evolution

Processor on which feature first appears

Problem	Solution	Processor on which feature first appears
External memory slower than the system bus.	Add external cache using faster memory technology.	386
Increased processor speed results in external bus becoming a bottleneck for cache access.	Move external cache on-chip, operating at the same speed as the processor.	486
Internal cache is rather small, due to limited space on chip	Add external L2 cache using faster technology than main memory	486
Contention occurs when both the Instruction Prefetcher and the Execution Unit simultaneously require access to the cache. In that case, the Prefetcher is stalled while the Execution Unit's data access takes place.	Create separate data and instruction caches.	Pentium
Increased processor speed results in external bus becoming a bottleneck for L2 cache access.	Create separate back-side bus that runs at higher speed than the main (front-side) external bus. The BSB is dedicated to the L2 cache.	Pentium Pro
	Move L2 cache on to the processor chip.	Pentium II
Some applications deal with massive databases and must have rapid access to large amounts of data. The on-chip caches are too small.	Add external L3 cache.	Pentium III
	Move L3 cache on-chip.	Pentium 4

Main memory Evolution

Year introduced	Chip size	\$ per MB	Total access time to a new row/column	Column access time to existing row
1980	64 Kbit	\$1500	250 ns	150 ns
1983	256 Kbit	\$500	185 ns	100 ns
1985	1 Mbit	\$200	135 ns	40 ns
1989	4 Mbit	\$50	110 ns	40 ns
1992	16 Mbit	\$15	90 ns	30 ns
1996	64 Mbit	\$10	60 ns	12 ns
1998	128 Mbit	\$4	60 ns	10 ns
2000	256 Mbit	\$1	55 ns	7 ns
2002	512 Mbit	\$0.25	50 ns	5 ns
2004	1024 Mbit	\$0.10	45 ns	3 ns

Comparison of Cache Sizes

Processor	Type	Year of Introduction	L1 cache ^a	L2 cache	L3 cache
IBM 360/85	Mainframe	1968	16 to 32 KB	—	—
PDP-11/70	Minicomputer	1975	1 KB	—	—
VAX 11/780	Minicomputer	1978	16 KB	—	—
IBM 3033	Mainframe	1978	64 KB	—	—
IBM 3090	Mainframe	1985	128 to 256 KB	—	—
Intel 80486	PC	1989	8 KB	—	—
Pentium	PC	1993	8 KB/8 KB	256 to 512 KB	—
PowerPC 601	PC	1993	32 KB	—	—
PowerPC 620	PC	1996	32 KB/32 KB	—	—
PowerPC G4	PC/server	1999	32 KB/32 KB	256 KB to 1 MB	2 MB
IBM S/390 G4	Mainframe	1997	32 KB	256 KB	2 MB
IBM S/390 G6	Mainframe	1999	256 KB	8 MB	—
Pentium 4	PC/server	2000	8 KB/8 KB	256 KB	—
IBM SP	High-end server/ supercomputer	2000	64 KB/32 KB	8 MB	—
CRAY MTAb	Supercomputer	2000	8 KB	2 MB	—
Itanium	PC/server	2001	16 KB/16 KB	96 KB	4 MB
SGI Origin 2001	High-end server	2001	32 KB/32 KB	4 MB	—
Itanium 2	PC/server	2002	32 KB	256 KB	6 MB
IBM POWER5	High-end server	2003	64 KB	1.9 MB	36 MB
CRAY XD-1	Supercomputer	2004	64 KB/64 KB	1MB	—

Memory Hierarchy in Power 4/5

Cache	Capacity	Associativity	Line size	Write policy	Repl. alg	Comments
L1 I-cache	64 KB	Direct/2-way	128 B		LRU	Sector cache (4 sectors)
L1 D-cache	32 KB	2-way/4-way	128 B	Write-through	LRU	
L2 Unified	1.5 MB/2 MB	8-way/10-way	128 B	Write-back	Pseudo-LRU	
L3	32 MB/36MB	8-way/12-way	512 B	Write-back	?	Sector cache (4 sectors)

Latency	P4 (1.7 GHz)	P5 (1.9 GHz)
L1 (I and D)	1	1
L2	12	13
L3	123	87
Main Memory	351	220

Virtual Memory

Provides *illusion* of very large memory

- sum of the memory of many jobs greater than physical memory
- address space of each job larger than physical memory

Allows available (fast and expensive) physical memory to be very well utilized

Simplifies memory management

Exploits memory hierarchy to keep average access time low.

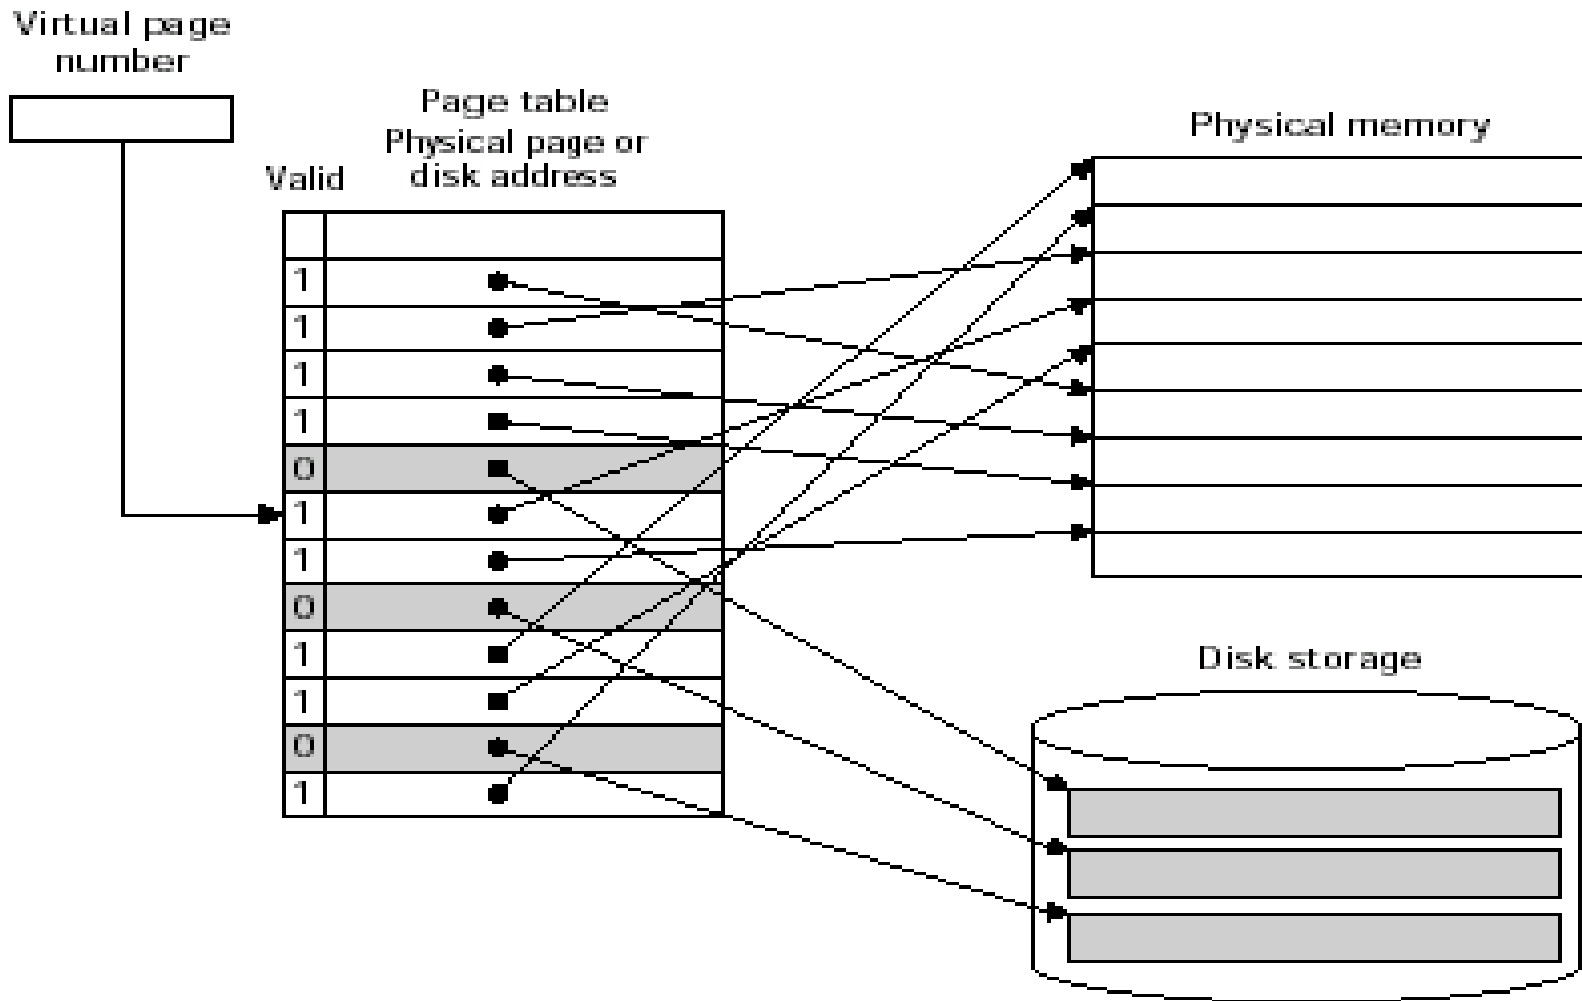
Involves at least two storage levels: *main* and *secondary*

Virtual Address -- address used by the programmer

Virtual Address Space -- collection of such addresses

Memory Address -- address of word in physical memory
also known as “physical address” or “real address”

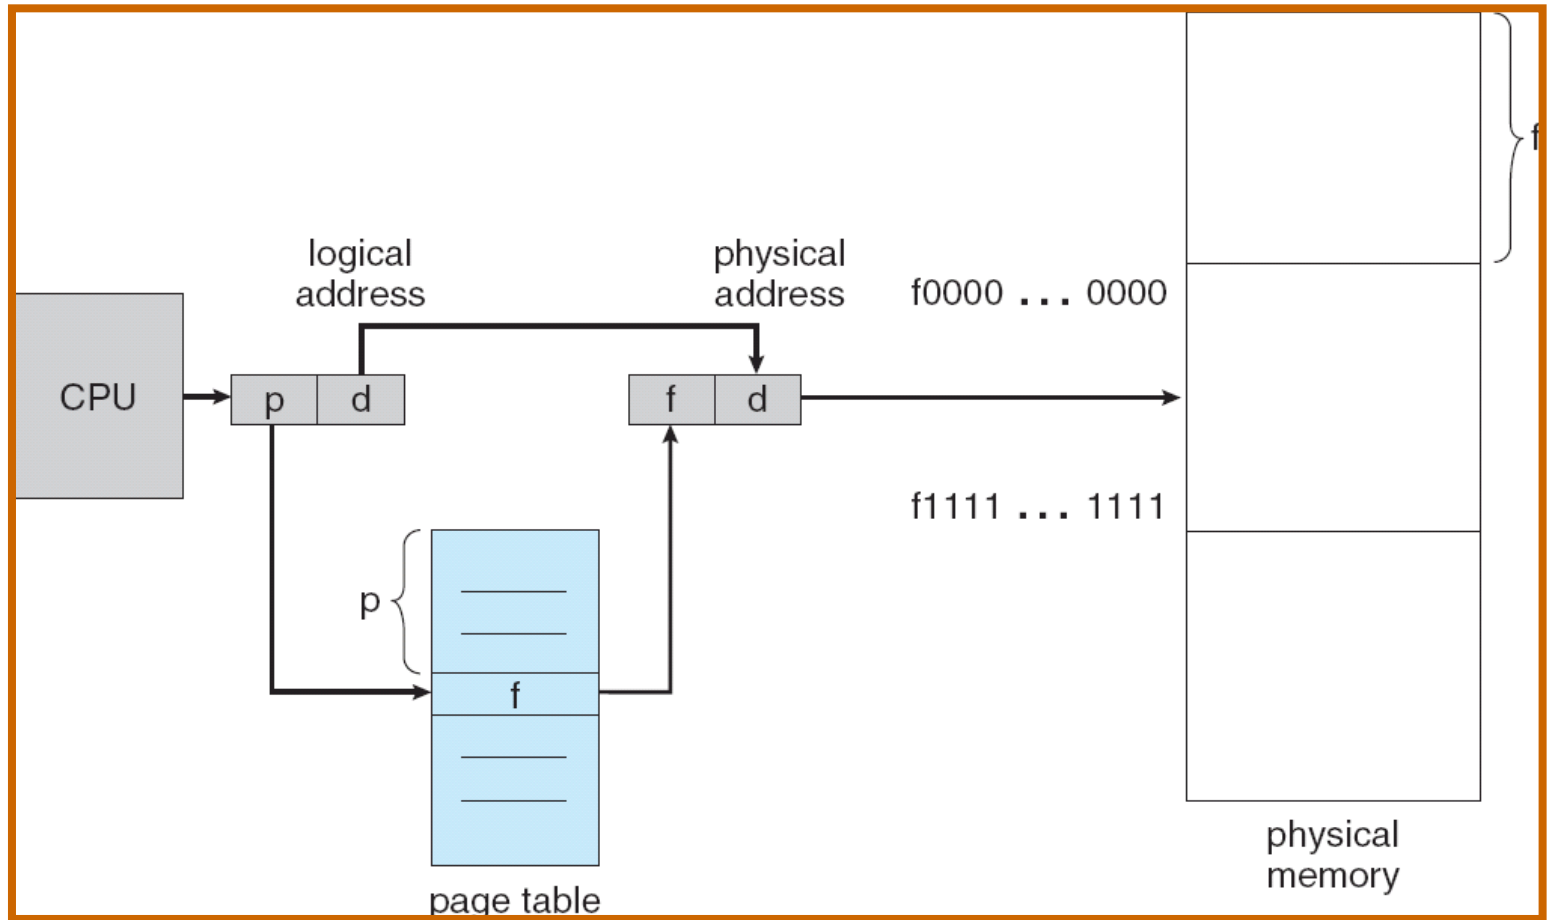
Page Table address Mapping



Implementing paging

- Keep track of all free physical page frames
- To run a program of size n pages, need to find n free frames and load program
 - Pages **do not** need to be contiguous!
- Set up mapping from virtual to physical in **page table**
- At memory reference time, translate virtual address to physical address using **page table**

Paging Hardware



Paging Examples

Assume a page size of 1K and a 15-bit logical address space.

How many pages are in the system?

- Answer: $2^5 = 32$.
- Assuming a 15-bit address space with 8 logical pages. How large are the pages?

- Answer: $2^5 = 32$.
- Assuming a 15-bit address space with 8 logical pages. How large are the pages?
- Answer: $2^{12} = 4K$. It takes 3 bits to reference 8 logical pages ($2^3 = 8$). This leaves 12 bits for the page size thus pages are 2^{12} .

Consider logical address 1025 and the following

page table for some process P0. Assume a 15-bit address space with a page size of 1K. What is the physical address to which logical address 1025 will be mapped?

8
0
2

Consider logical address 1025 and the following

page table for some process P0. Assume a 15-bit address space with a page size of 1K. What is the physical address to which logical address 1025 maps?

8
0
2

Step 1. Convert to binary:

00001000000000
1

Consider logical address 1025 and the following

page table for some process P0. Assume a 15-bit address space with a page size of 1K. What is the physical address to which logical address 1025 maps?

8
0
2

Step2. Determine the logical page number:

Since there are 5-bits allocated to the logical page, the address is broken up as follows:

00001

0000000001

Logical page number offset within
page

Consider logical address 1025 and the following

page table for some process P0. What is the physical address?

00001 →	8
	0
	2

Step 3. Use logical page number as an index into the page table.

00001 0000000001

Consider logical address 1025 and the following

page table for some process P0. What is the physical address?

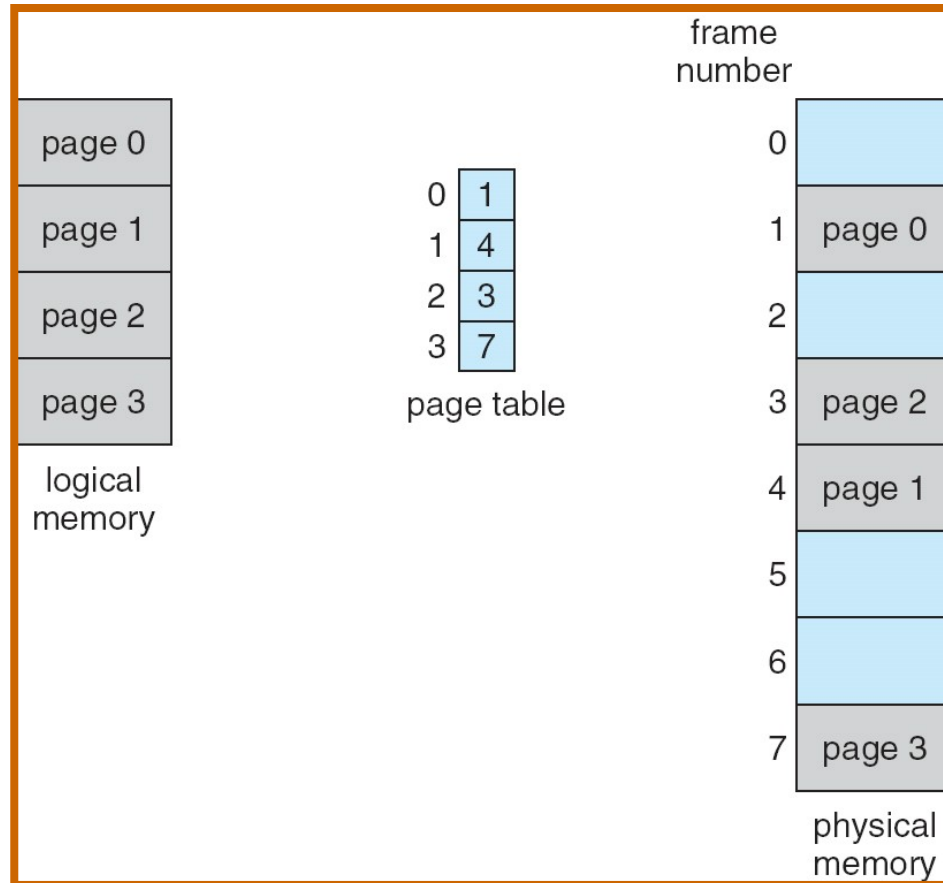
00001 →	8
	0
	2

Take the physical page number from the page table and concatenate the offset.

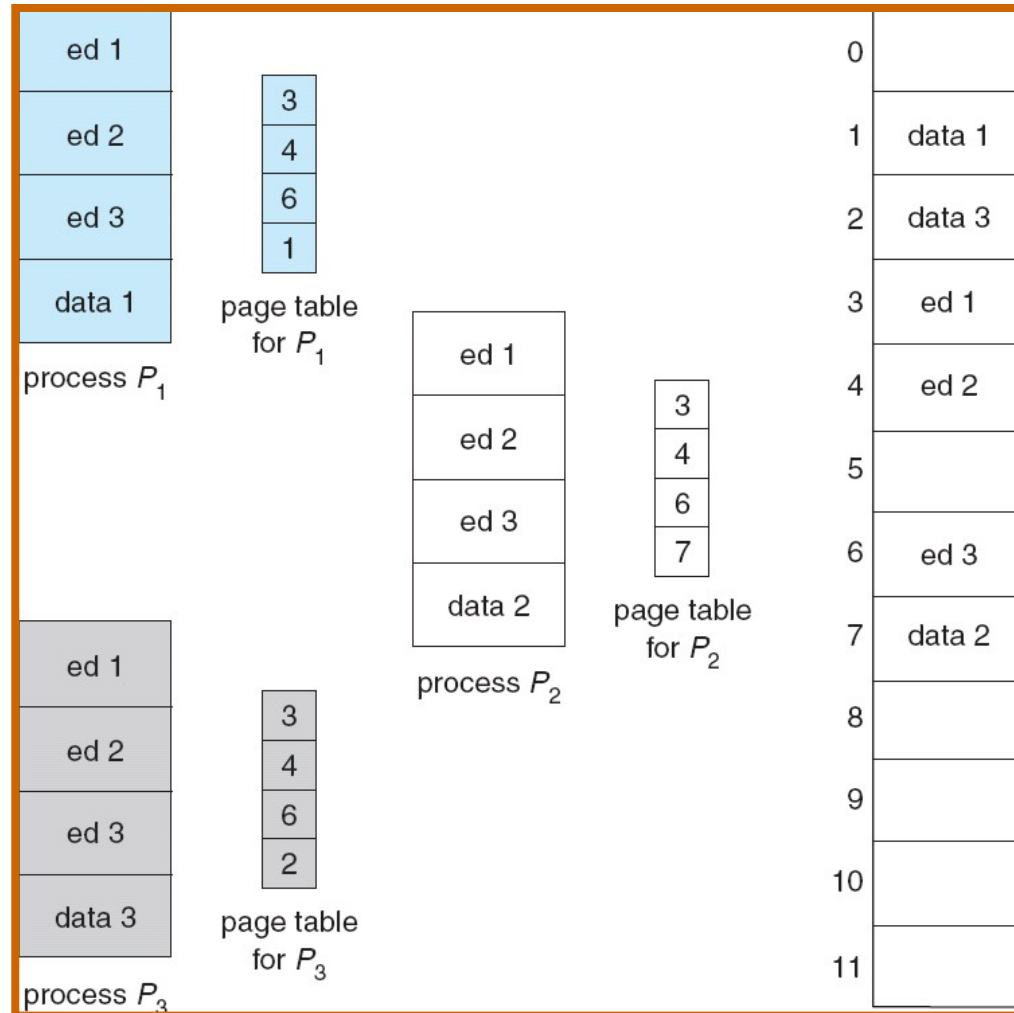
The physical address is byte 1.

0000000000000001

Paging Model of Logical and Physical Memory



Shared Pages Example



Valid (v) or Invalid (i) Bit in a Page Table

00000	page 0
	page 1
	page 2
	page 3
	page 4
10,468	page 5
12,287	


frame number		valid-invalid bit
0	2	v
1	3	v
2	4	v
3	7	v
4	8	v
5	9	v
6	0	i
7	0	i

page table

0	
1	
2	page 0
3	page 1
4	page 2
5	
6	
7	page 3
8	page 4
9	page 5
	⋮
	page n

Hierarchical Page Tables

- Break up the logical address space into multiple page tables
- A simple technique is a two-level page table



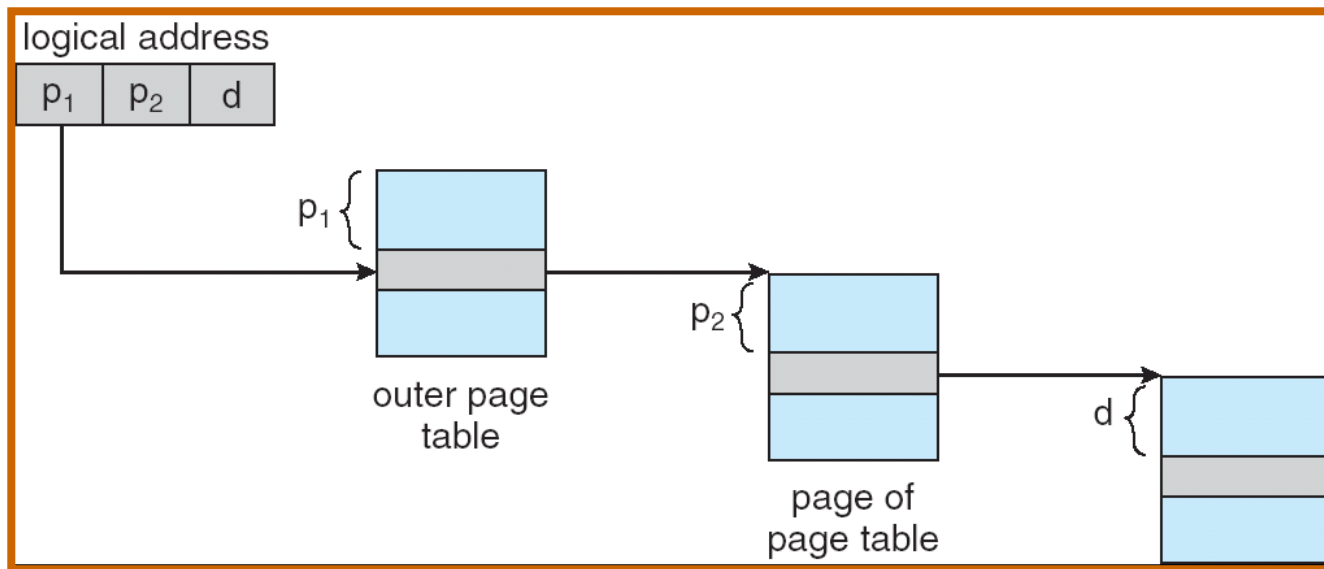
Two-Level Paging Example

- A logical address (on 32-bit machine with 1K page size) is divided into:
 - a page number consisting of 22 bits
 - a page offset consisting of 10 bits
- Since the page table is paged, the page number is further divided into:
 - a 12-bit page number
 - a 10-bit page offset
- Thus, a logical address is as follows:

page number		page offset
p_1	p_2	d

where p_1 is an index into the outer page table, and p_2 is the displacement within the page of the outer page table

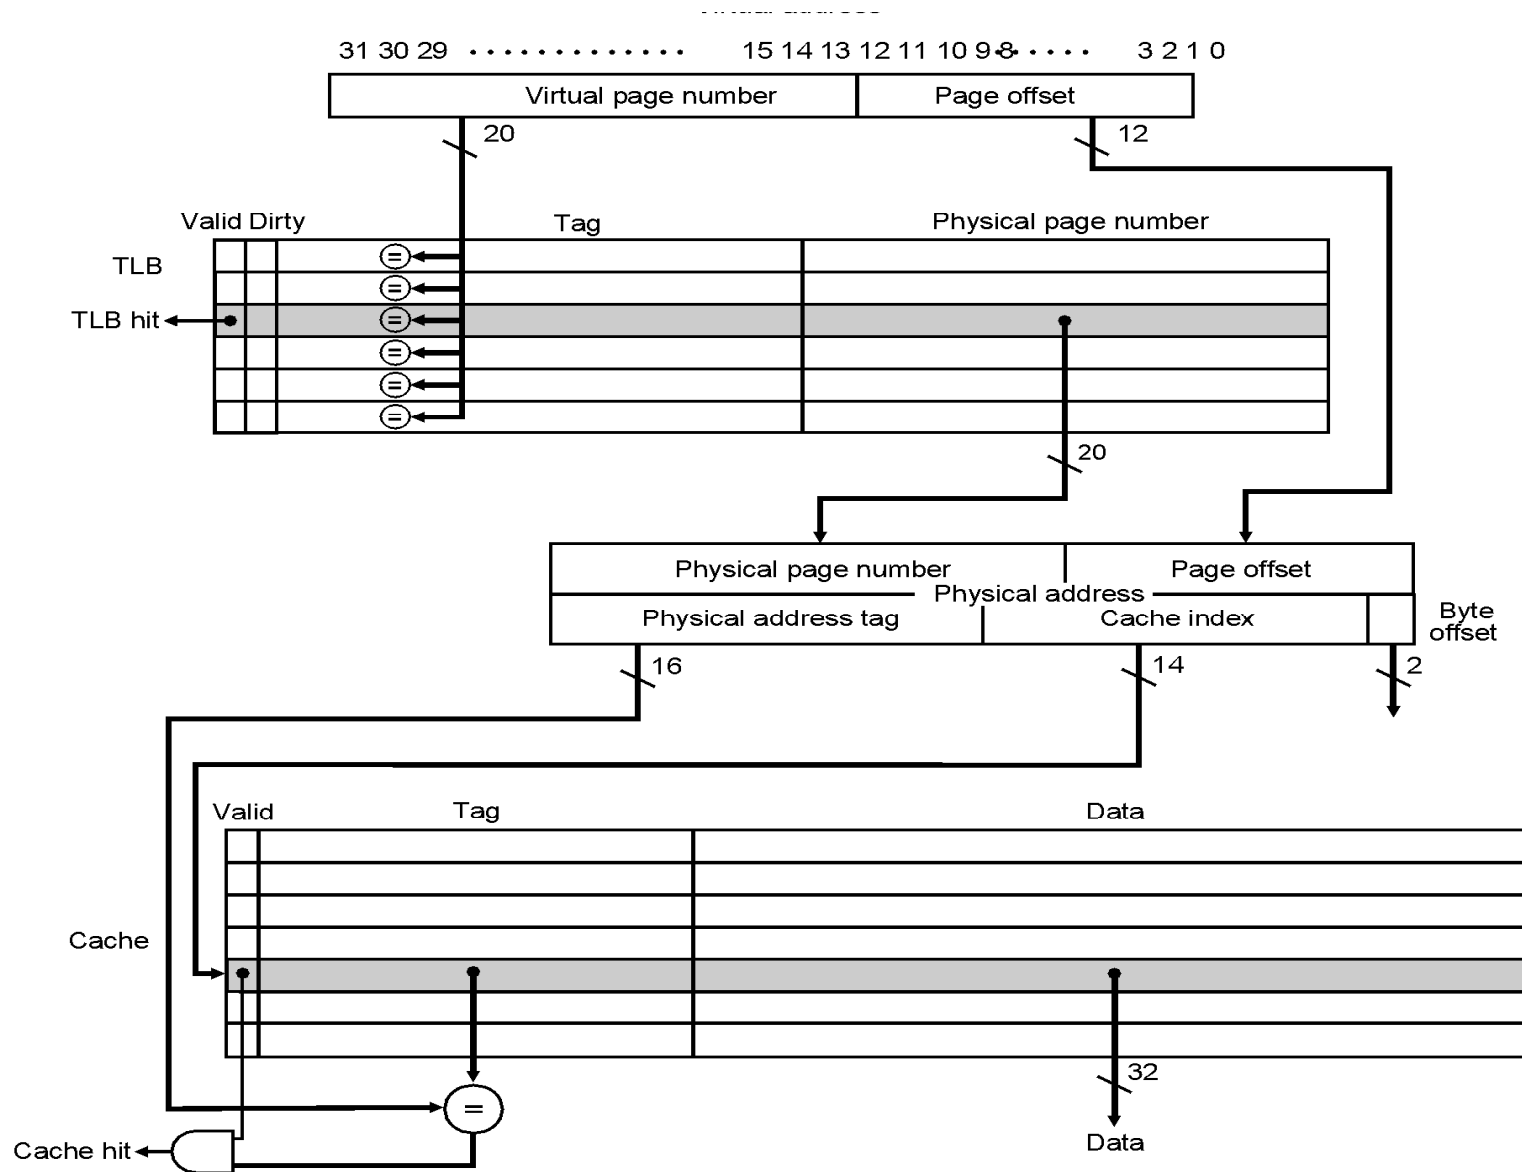
Address-Translation Scheme



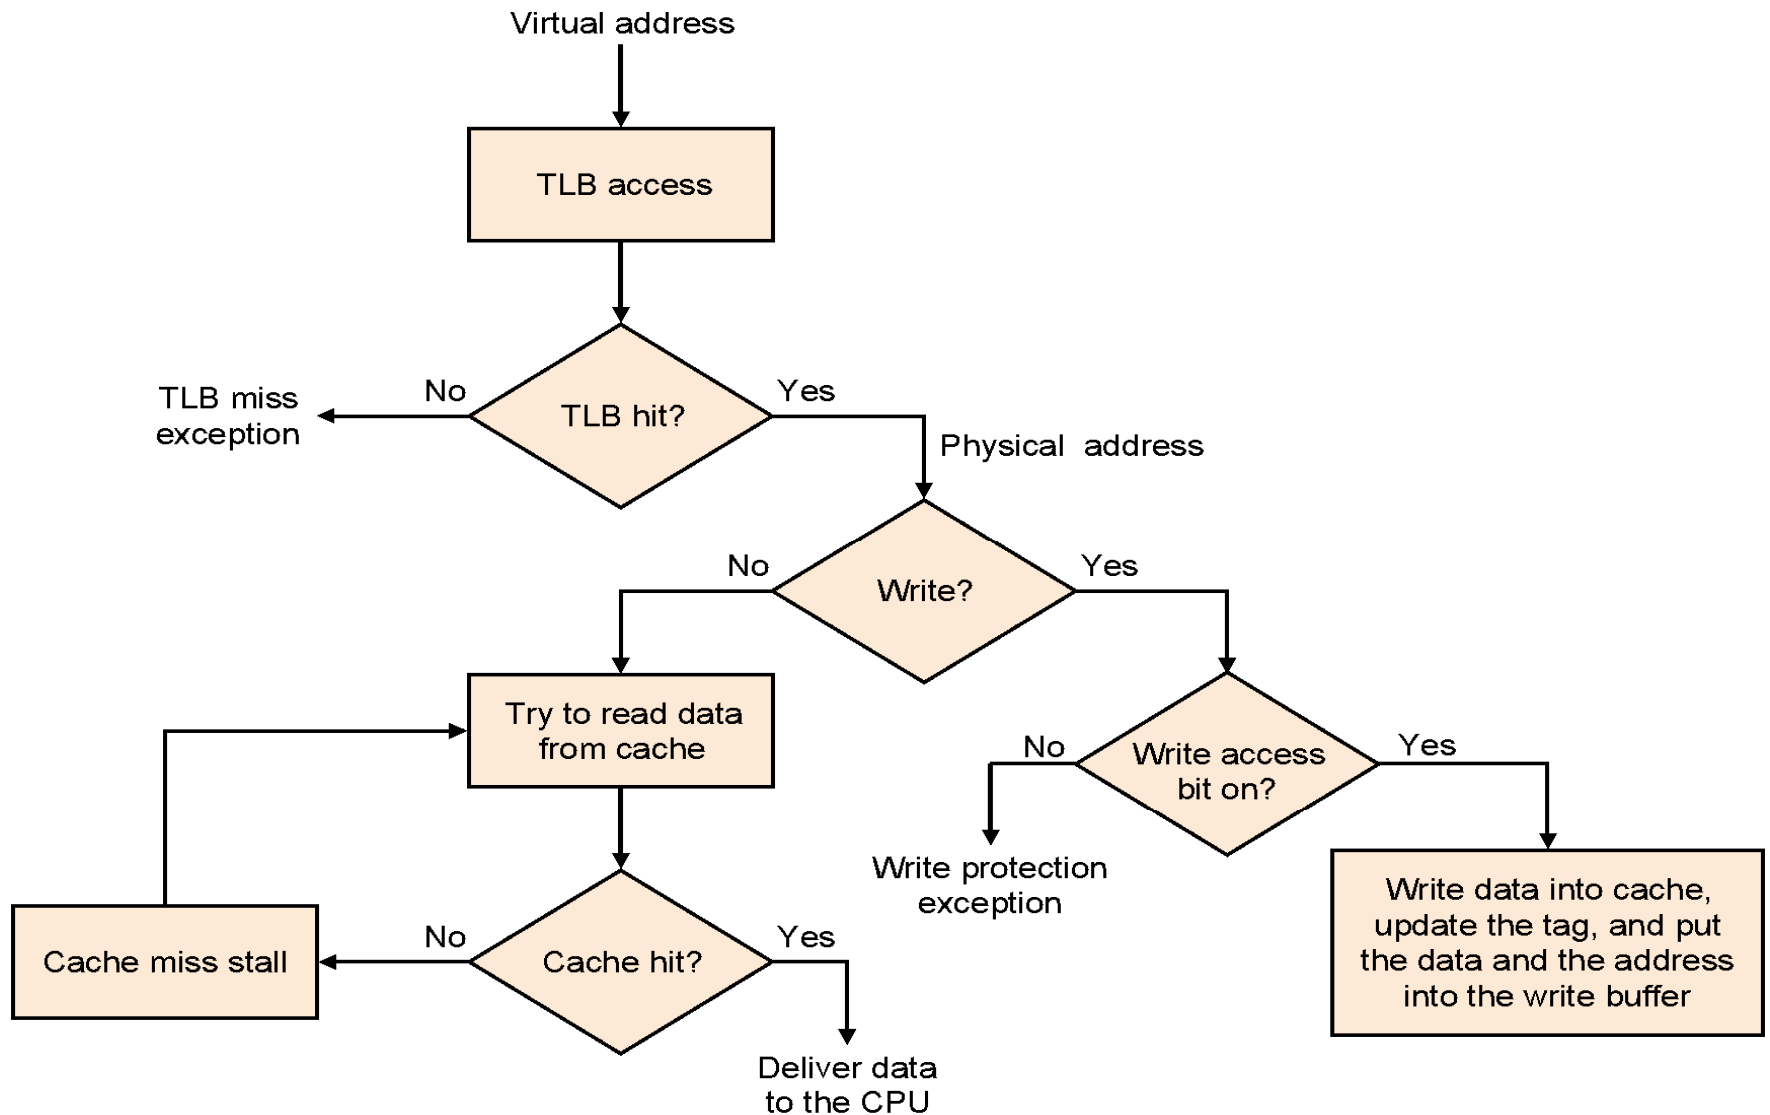
Paging in 64 bit Linux

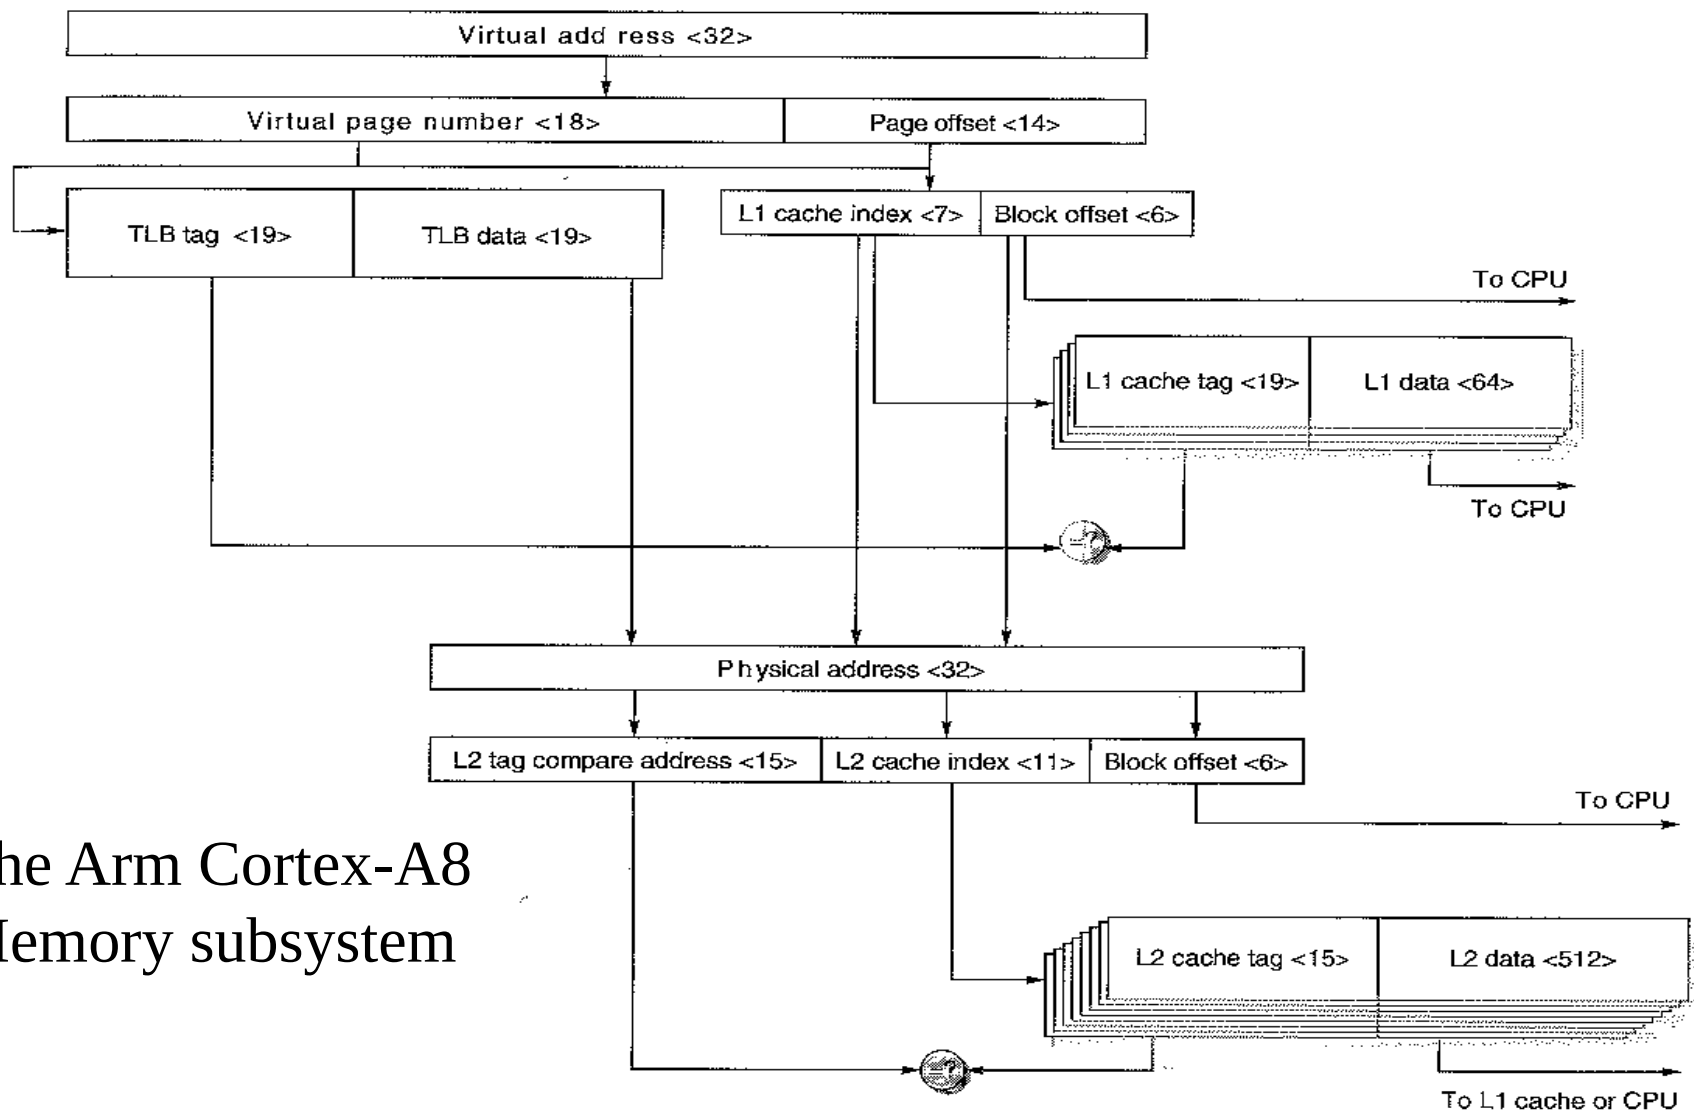
Platform	Page Size	Address Bits Used	Paging Levels	Address Splitting
Alpha	8 KB	43	3	10+10+10+13
IA64	4 KB	39	3	9+9+9+12
PPC64	4 KB	41	3	10+10+9+12
sh64	4 KB	41	3	10+10+9+12
X86_64	4 KB	48	4	9+9+9+9+12

TLB look up and cache access



Cache, sequence of events for





The Arm Cortex-A8 Memory subsystem

Figure 2.16 The virtual address, physical address, indexes, tags, and data blocks for the ARM Cortex-A8 data caches and data TLB. Since the instruction and data hierarchies are symmetric, we show only one. The TLB (instruction or data) is fully associative with 32 entries. The L1 cache is four-way set associative with 64-byte blocks and 32 KB capacity. The L2 cache is eight-way set associative with 64-byte blocks and 1 MB capacity. This figure doesn't show the valid bits and protection bits for the caches and TLB, nor the use of the way prediction bits that would dictate the predicted bank of the L1 cache.

Summary

- **SRAM** is fast but expensive and not very dense:
 - Good choice for providing the user FAST access time.
- **DRAM** is slow but cheap and dense:
 - Good choice for presenting the user with a BIG memory system
- **Flash Memory : Read** is fast but write is slow
- **Two Different Types of Locality:**
 - **Temporal Locality** (Locality in Time): If an item is referenced, it will tend to be referenced again soon.
 - **Spatial Locality** (Locality in Space): If an item is referenced, items whose addresses are close by tend to be referenced soon.
- **Cache**
 - By taking advantage of the principle of locality:
 - Present the user with as much memory as is available in the cheapest technology.
 - Provide access at the speed offered by the fastest technology.
- **Direct Mapped , Set associative, full associative**
- **Virtual Memory**