

# CS354: Database

## Recap: SQL Queries

---

SELECT      [DISTINCT] <attribute list>  
FROM        <table list>  
[WHERE      <condition on the tables>]  
[GROUP BY <grouping attributes>]  
[HAVING     <group condition>]  
[ORDER BY <attribute list>    ASC | DESC]  
[LIMIT      <number of tuples>]

# SQL Query: Temporal Relation

---

- Result of a **SELECT** clause that exists temporally, which assists you in formulating a query
- Syntax:  
**SELECT** <attributes>  
**FROM** R1, R2, (SELECT ... ) <alias>, ..., RN  
**WHERE** <condition>;
- Must always use an alias to denote the result relation of the **SELECT** command

# SQL Example: Temporal Relation

---

Find fname, lname of male employees with salary > 50K

```
SELECT *  
FROM   (SELECT fname, lname, salary  
        FROM   employee  
        WHERE  sex = 'M') r1  
WHERE  r1.salary > 50000
```

# SQL Query: Temporal Relation Notes

---

- You can use multiple temporal relations
- You cannot use a temporal relation to create another temporal relation

Example of incorrect usage:

```
SELECT ...  
FROM    ...,  
        (SELECT ...) r1,  
        (SELECT ... FROM r1 ...) r2,  
WHERE ...;
```

# SQL Query: WITH

---

- SQL-99 standard introduced **WITH** clause to help refine the result of a query (another way to achieve temporal relation)
  - Some vendors do not support **WITH** (e.g., MySQL)
- Syntax:  
**WITH** <alias> AS (SELECT ...)[  
    <alias2> AS (SELECT ...)]  
**SELECT** <query>;
- Can be used to perform “refinement” on a query
  - Subsequent queries in the **WITH** clause can use the results of the previous query

# SQL Example: WITH

---

Find all information on dependents of John Smith

```
WITH r1 as (SELECT *  
             FROM employee  
             WHERE fname = 'John'  
                   AND lname = 'Smith')  
SELECT *  
FROM   dependent  
WHERE  essn IN (SELECT ssn from r1);
```

# SQL Query: JOIN Operations

---

- SQL-99 standard added several join operations:
  - **INNER JOIN** (normal join)
  - **LEFT JOIN** (left outer join)
  - **RIGHT JOIN** (right outer join)
  - **FULL JOIN** (outer join)
- Each operation results in a relation
- Operation can only appear in:
  - **FROM** clause of **SELECT** command
  - **WHERE** clause of **SELECT** command with an operator that uses a sub-query



# SQL Query: [INNER] JOIN

---

- Compute the (inner) join between tables `r1` and `r2` with a given join condition
- Syntax:  
`r1 JOIN r2 ON <join-condition>;`  
or  
`r1 INNER JOIN r2 ON <join-condition>;`
- **JOIN** operator makes the SQL query look a lot like RA query
- Can join more than 2 relations

# SQL Example: INNER JOIN

---

Find fname, lname of employees in the 'Research' department

RA Query:

$\pi_{\text{fname, lname}}(\sigma_{\text{dname}='Research'}(\text{EMPLOYEE} \bowtie_{\text{dno}=\text{dnumber}} \text{DEPARTMENT}))$

SQL Query:

```
SELECT fname, lname
FROM   (employee JOIN department
        ON dno = dnumber)
WHERE  dname = 'Research';
```

# SQL Query: OUTER JOIN

---

- Compute the outer join between tables `r1` and `r2` with a given join condition - see RA slides for details on difference between left, right, and full outer joins
- Syntax:  
`r1 LEFT | RIGHT | FULL [OUTER] JOIN r2 on <join condition>;`
- Results in NULL values for the attributes where non-matching tuples occur

# SQL Query: NATURAL JOIN

---

- Compute the natural join on attributes with the same names from two or more tables with the common attribute appearing only once in the result
- Syntax:  
**r1 NATURAL JOIN r2;**
- Example:  
**SELECT \***  
**FROM works\_on NATURAL JOIN dependent;**

# SQL Query: CROSS JOIN

---

- Cross join is the same as a Cartesian Product
- Syntax:  
**r1 CROSS JOIN r2;**
- Example:  
**SELECT** ssn, fname, lname, dno, dnumber, dname  
**FROM** employee **CROSS JOIN** dependent;

# SQL Outline

---

- Data definition
  - Database Creation
  - Table Creation
- Query (SELECT)
- Data update (INSERT, DELETE, UPDATE)
- View definition



# SQL Modifications/Updates

---

- A modification command does not return a result but it changes the database
- There are 3 kinds of modifications
  - **INSERT** tuple(s)
  - **DELETE** tuple(s)
  - **UPDATE** the value(s) of existing tuples

# SQL Modification: INSERT

---

- Add one more more tuples to an existing relation
- Two forms of **INSERT**:
  - Literal values (constant or known values)
  - Result from a **SELECT** command



# SQL Modification: INSERT (2)

---

Inserting a tuple using literal/constant values

Syntax:

**INSERT INTO <table name>[(<attr names>)]  
VALUES (<list of values>);**

- Complete tuple: omitting [(<attr names>)] means you must specify all attribute values in the exact order defined in relation
- Partial tuple: specify a subset of the attribute values in the same order as the list of attributes [(<attr names>)]

## SQL Modification: INSERT (3)

---

Inserting a tuple using **SELECT** command

Syntax:

**INSERT INTO** <table name>[(<attr names>)] (<**SELECT** subquery>)

- Multiple tuples may be added dependent on the **SELECT** subquery relation

# SQL Example: INSERT

---

- Complete tuple:

```
INSERT INTO employee VALUES ('Joyce', 'C', 'Ho',  
'111223333', '1985-02-05', '400 Dowman Drive,  
Atlanta, GA', 'F', '150000', '987654321', 5);
```

- Partial tuple:

```
INSERT INTO employee(fname, lname, ssn) VALUES  
('Joyce', 'Ho', '111223333');
```

# SQL Example: INSERT w/ SELECT

---

Suppose we want a new table that has the name, number of employees, and total salaries for each department. We first create the table then load it with the information from the database.

```
CREATE TABLE dept_info  
( dept_name  VARCHAR(10),  
  no_of_emps INT,  
  tot_salary  INT);
```

```
INSERT INTO dept_info  
  (SELECT      dname, count(*), sum(salary)  
   FROM        department, employee  
   WHERE       dnumber = dno  
   GROUP BY dname);
```

# MySQL: Bulk Import

---

- All respectable RDBMS provide utilities to import data from text files
  - Syntax for uploading data will vary based on vendor
- MySQL allows the **LOAD DATA INFILE**  
(<http://dev.mysql.com/doc/refman/5.7/en/load-data.html>)
  - For a pipe-delimited file (| separates each column):  
**LOAD DATA LOCAL INFILE <filename>**  
**{REPLACE | IGNORE} INTO TABLE <table name>**  
**FIELDS TERMINATED BY '|';**

# SQL Modification: DELETE

---

- Remove tuples from a relation
- Syntax:  
**DELETE FROM** <relation>  
**WHERE**           <condition>;
- Be careful! All tuples that satisfy the condition clause are deleted
- Tuples are deleted from only one table at a time unless **CASCADE** is specified on a referential integrity constraint
- What happens if we don't specify a **WHERE** clause?

# SQL Example: DELETE

---

Delete all employees with the last name Brown

```
DELETE FROM employee  
WHERE lname = 'Brown';
```

## SQL Example: DELETE (2)

---

Delete all employees from the 'Research' department who have more than 2 dependents

```
DELETE FROM employee
WHERE dno IN (SELECT dnumber
               FROM department
               WHERE dname = 'Research')
AND ssno IN (SELECT essno
              FROM dependent
              GROUP BY essno
              HAVING COUNT(name) > 2);
```



# SQL Modification: UPDATE

---

- Modify/change certain attributes in certain tuples of a relation
- Syntax:  
**UPDATE** <relation>  
**SET** <list of attribute assignments>  
**WHERE** <condition>;
- **UPDATE** command modifies tuples in the same relation

# SQL Example: UPDATE

---

Change the location and controlling department number of project 10 to 'Bellaire' and 5, respectively.

```
UPDATE project  
SET      plocation = 'Bellaire', dnum = 5  
WHERE   pnumber = 10;
```

## SQL Example: UPDATE (2)

---

Give all employees in the 'Research' department a 10% raise

```
UPDATE employee
SET      salary = salary * 1.1
WHERE dno IN (SELECT dnumber
                FROM department
                WHERE dname = 'Research');
```

- Reference to `salary` attribute on the right of `=` refers to the salary value before modification
- Reference to `salary` attribute on the left of `=` refers to salary value after modification

# SQL: VIEW


---

- A view is a virtual table, a relation that is defined in terms of the contents of other tables and views
- A view does *not* exist in the physical form
- In contrast, a relation whose value is really in the database is called a base table
- Syntax:  
**CREATE VIEW <name> AS <query>;**

# SQL: View & Logical Data Independence

---

- Recall Logical Data Independence (class on Database Concepts)
  - Ability to present the stored information in a different way to different users
- View can be adapted to the need of the user
- If conceptual schema changes, only the **SELECT** query needed to construct view needs to change



# SQL Example: VIEW

---

- Suppose an administrator maintains a list of activities of all employees which contains the following information:  
fname, lname, project\_name, hours\_worked
- Regular SELECT query:  

```
SELECT fname, lname, pname, hours  
FROM   employee, works_on, project  
WHERE  ssn = essn AND pno = pnumber;
```
- Create VIEW for the admin:  

```
CREATE VIEW emp_activity  
AS (SELECT fname, lname, pname, hours  
    FROM   employee, works_on, project  
    WHERE  ssn = essn AND pno = pnumber);
```

# SQL: VIEW Advantages

---

- View can be used in queries like an ordinary relation
  - When a view is used in a **SELECT** query, the virtual relation is computed first
- Simplify complex queries by hiding them from the end-user and applications
- Limit data access to specific users (expose only non-sensitive data) and provides extra security for read/write access
- Enables backward compatibility - changes to database won't affect changes to other applications

# SQL: VIEW Disadvantages

---

- Querying data from database view can be slow (since view is computed each time)
- Tables dependency - updates to the underlying tables will force changes to the view itself to make it work properly
- Most data manipulation statements (**INSERT**, **DELETE**, **UPDATE**) are not possible on the view



# SQL Data Update & View: Recap

---

- Query
  - Temporal Relation / WITH
  - JOIN
- SQL Modification
  - INSERT
  - DELETE
  - UPDATE
- SQL Views



# More SQL Practice

---

Find the fname, lname of employees with more than 2 dependents and work on all projects controlled by department #1

# More SQL Practice

---

Find the fname, lname of employees with more than 2 dependents and work on all projects controlled by department #1

```
SELECT fname, lname
FROM   employee
WHERE  <employee has more than 2 dependents>
      AND <works on all projects controlled by dept #1>;
```

First formulate in words  
conquer each subquery separately

## More SQL Practice: Subquery #1

---

Find employee that has more than 2 dependents

```
SELECT essn  
FROM dependent  
GROUP BY essn  
HAVING COUNT(name) > 2;
```

## More SQL Practice: Subquery #2

---

Find employees that works on all projects controlled by department #1 - set difference technique

```
SELECT ssn
FROM   employee e
WHERE  NOT EXISTS
      < set of projects controlled by department #1 >
      - <set of projects worked on by e.ssn>;
```

## More SQL Practice: Subquery #2

---

Find employees that works on all projects controlled by department #1 - set difference technique

```
SELECT ssn
FROM   employee e
WHERE NOT EXISTS (SELECT pnumber
                  FROM   project
                  WHERE pnumber IN
                        ( SELECT pnumber
                          FROM   project
                          WHERE dnum = 1)
                  AND pnumber NOT IN
                        ( SELECT pno
                          FROM   works_on
                          WHERE essn = e.ssn));
```

## More SQL Practice: Putting it Together

Find the fname, lname of employees with more than 2 dependents and work on all projects controlled by department #1

[illegible]

## More SQL Practice (2)

---

Find the department name, and the number of employees in that department that earns more than 40K for departments with at least 2 employees



## More SQL Practice (2)

---

Find the department name, and the number of employees in that department that earns more than 40K for departments with at least 2 employees

```
SELECT    dname, COUNT(ssn)
FROM      department, employee
WHERE     dnumber = dno
        AND salary > 40000
GROUP BY  dname
HAVING    COUNT(ssn) >= 2;
```

What is wrong with this solution?

## More SQL Practice (2): Solution

---

Find the department name, and the number of employees in that department that earns more than 40K for departments with at least 2 employees

```
SELECT    dname, COUNT(ssn)
FROM      department, employee
WHERE     dnumber = dno
        AND salary > 40000
        AND dno IN ( SELECT    dno
                      FROM      employee
                      GROUP BY dno
                      HAVING     COUNT(ssn) >= 2 )
GROUP BY dname;
```

## More SQL Practice (3)

---

Find fname, name of employees who work on 2 or more projects together with John Smith

## More SQL Practice (4)

---

Find departments who have 2 or more employees working on all projects controlled by 'Research' department

