

# CS354: Database



**Dr. Raju Halder**

# CS354: Database

- Instructors: Dr. Raju Halder/Dr. Mayank Agarwal
  - halder@iitp.ac.in
  - Meeting time: Wednesday
- Teaching Assistants
  - Ashish Kumar Ranjan(ashish.mtcs17@iitp.ac.in)
  - TBA


# References

- Elmasri Ramez and Navathe Shamkant, **Fundamentals of Database System**. Pearson Education, Seventh edition.
- Silberschatz, Korth, and Sudarshan, **Database System Concepts**. McGraw Hill Education, Sixth edition
- Hector Garcia-Molina, Jeffrey D. Ullman, Jennifer D. Widom, **Database System Implementation**. Pearson Education.



# References

- Lex de Haan and Toon Koppelaars, **Applied Mathematics for Database Professionals**. Apress, first edition (June 19, 2007)



- Introduction
  - What is a database?
  - Characteristics of a database
  - Database users
- History of databases



# What is Database


- **Managing** large amounts of **data** is an integral part of most nowadays business and governmental activities
  - Collecting taxes
  - Bank account management
  - Bookkeeping
  - Airline reservations
  - Human resource management
  - ...





# What is Database

- **Databases** are needed to manage that **vast amount of data**
- A database (**DB**) is a collection of **related data**
  - Represents some aspects of the **real world**
    - **Universe of discourse**
  - Data is logically **coherent**
  - Is provided for an intended group of **users** and **applications**



# What is Database

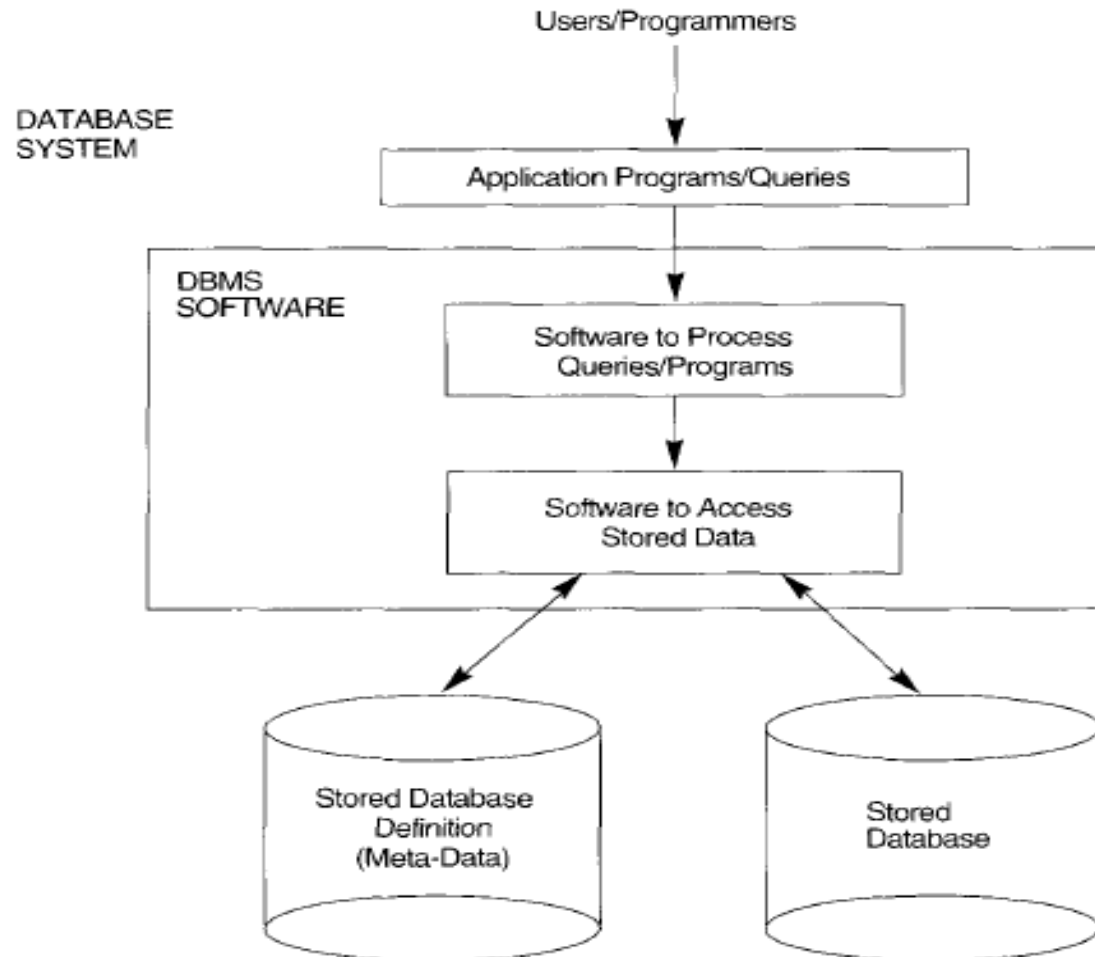
- A database management system (**DBMS**) is a collection of programs to maintain a database, that is, for

- Definition of data and structure
- Physical construction
- Manipulation
- Sharing/Protecting
- Persistence/Recovery






# A Simplified Database System Environment



# File System

- A file system is not a database!
- File management systems are **physical** interfaces



# File System

- **Advantages**

- Fast and easy access

- **Disadvantages**

- Uncontrolled redundancy
  - Inconsistent data
  - Limited data sharing and access rights
  - Poor enforcement of standards
  - Excessive data and access paths maintenance
- 

# Database Vs. File System

- Databases are **logical** interfaces
  - Retrieval of data using **data semantics**
  - Controlled redundancy
  - Data consistency & integrity constraints
  - Effective and secure data sharing
  - Backup and recovery
- However...
  - More complex
  - More expensive data access



# Database Vs. File System

- **DBMS** replaced previously dominant file-based systems in **banking** due to special requirements
  - **Simultaneous** and quick access is necessary
  - Failures and loss of data **cannot** be tolerated
  - Data always has to remain in a **consistent** state
  - Frequent queries and modifications




# Database Characteristics

- Databases **control redundancy**
  - Same data used by different applications or tasks is **stored only once**
  - Access via a **single interface** provided by DBMS
  - Redundancy only purposefully used to speed up data access (e.g. materialized views)
- Problems of **uncontrolled redundancy**
  - Difficulties in consistently updating data
  - Waste of storage space



# Database Characteristics

- Databases are **well-structured**, e.g. ER model:
  - Simple banking system



- Relational Databases provide
  - **Catalog** (data dictionary) contains all **meta data**
  - Defines the **structure** of the data in the database

# Database Characteristics

- Databases support **declarative querying**
  - Just specify what you want, not how and from where to get it
  - Queries are separated and abstracted from the actual physical organization and storage of data
- Get the first name of all customers with last name “Smith”
  - File system: Trouble with physical organization of data
    - Load file “c:\datasets\customerData.csv”
    - Build a regular expression and iterate over lines:  
If 2<sup>nd</sup> word in line equals “Smith,” then return 3<sup>rd</sup> word
    - Stop when end-of-file marker is reached
  - Database system: simply query
    - `SELECT first_name FROM data WHERE last_name='Smith'`



# Database Characteristics

- Databases aim at **efficient** manipulation of data
  - Physical tuning allows for good data allocation
  - Indexes speed up search and access
  - Query plans are optimized to improve performance
- Example: Simple Index

Index File

(checking accounts)

number
4543032
7809849
8942214

Data File

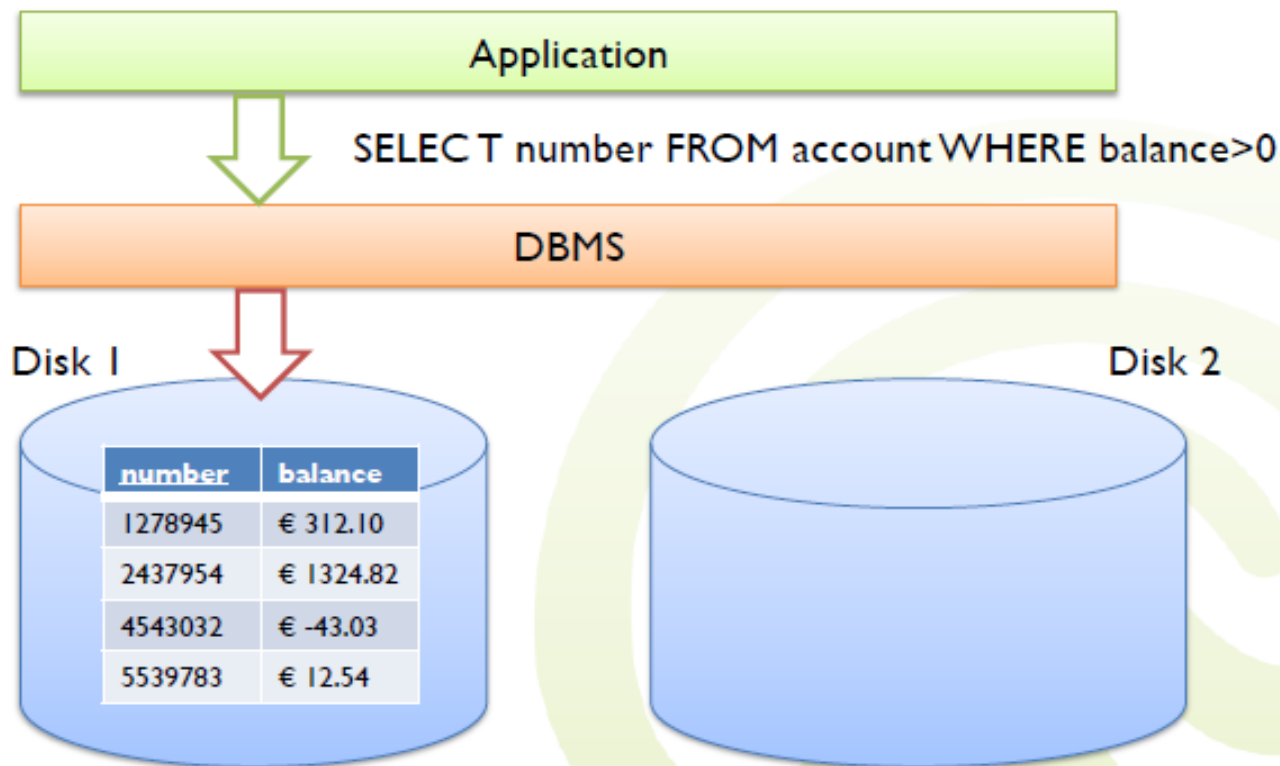
number	type	balance
1278945	saving	€ 312.10
2437954	saving	€ 1324.82
4543032	checking	€ -43.03
5539783	saving	€ 12.54
7809849	checking	€ 7643.89
8942214	checking	€ -345.17
9134354	saving	€ 2.22
9543252	saving	€ 524.89

# Database Characteristics

- **Isolation** between applications and data
  - Database employs **data abstraction** by providing **data models**
  - Applications work only on the **conceptual representation** of data
    - Data is strictly **typed** (Integer, Timestamp, Varchar, ...)
    - Details on where data is actually stored and how it is accessed are **hidden** by the DBMS
    - Applications can access and manipulate data by invoking **abstract operations** (e.g. SQL select statements)
  - DBMS-controlled parts of the file system are **protected** against external manipulations (tablespaces)

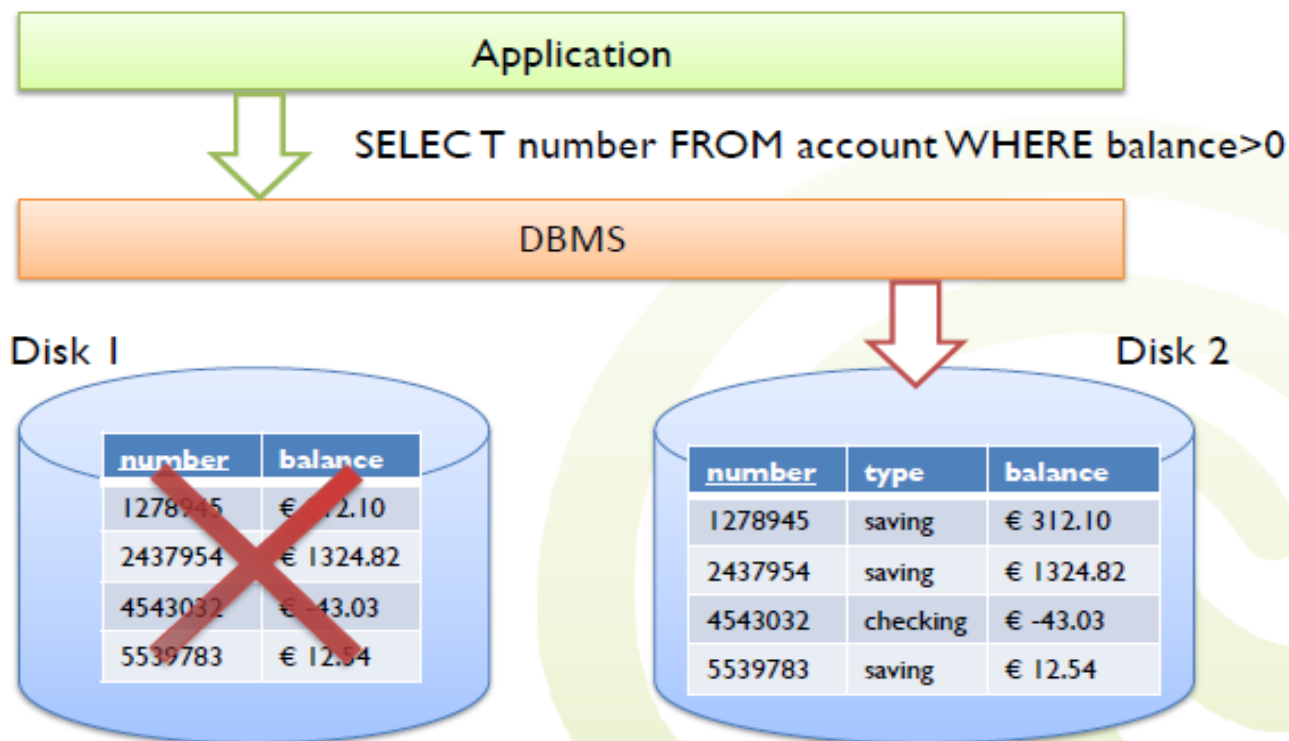
# Database Characteristics

- **Example:** Schema can be changed and tablespace moved without being noticed by app



# Database Characteristics

- **Example:** Schema can be changed and tablespace moved without being noticed by app





# Database Characteristics

- Supports multiple **views** of the data



- Views provide a different perspective of the DB
  - A user's conceptual understanding or task-based excerpt of the data (e.g. aggregations)
  - Security considerations and access control (e.g. projections)
- For applications, a view does not differ from a table
- Views may contain **subsets** of a DB and/or contain **virtual data**
  - Virtual data is **derived** from the DB (mostly by simple SQL statements, e.g. joins over several tables)
  - Can either be computed at query time or **materialized** upfront

# Database Characteristics

- Example views: **Projection**
  - Saving account clerk vs. checking account clerk

Original Table

<u>number</u>	type	balance
1278945	saving	€ 312.10
2437954	saving	€ 1324.82
4543032	checking	€ -43.03
5539783	saving	€ 12.54
7809849	checking	€ 7643.89
8942214	checking	€ -345.17
9134354	saving	€ 2.22
9543252	saving	€ 524.89

Saving View

<u>number</u>	balance
1278945	€ 312.10
2437954	€ 1324.82
5539783	€ 12.54
9134354	€ 2.22
9543252	€ 524.89

Checking View

<u>number</u>	balance
4543032	€ -43.03
7809849	€ 7643.89
8942214	€ -345.17

# Database Characteristics

- **Sharing** of data and support for **atomic multi-user transactions**
  - Multiple users and applications may access the DB at the same time
  - **Concurrency control** is necessary for maintaining consistency
  - **Transactions** need to be **atomic** and **isolated** from each other



# Database Characteristics

- Example: Atomic transactions

- **Program:**

Transfer  $x$  Euros from Account 1 to Account 2

1. Debit amount  $x$  from Account 1
2. Credit amount  $x$  to Account 2



# Database Characteristics

- Example: Atomic transactions

- **Program:**

- Transfer  $x$  Euros from Account 1 to Account 2

1. Debit amount  $x$  from Account 1
2. Credit amount  $x$  to Account 2

- **But what happens if the system fails after performing the first step?**






# Database Characteristics

- **Example:** Multi-user transactions
  - Program: Withdrawal of amount  $x$  from Account I
    1. Read old balance from DB
    2. New balance  $:=$  old balance  $- x$
    3. Write new balance back to the DB
  - Problem: Dirty Read
    - Account I has €500
    - User 1 wants to withdraw €20
    - User 2 wants to withdraw €80 **at the same time**
  - Without multi-user transactions, the account's balance is either €480 or €420, but not €400 (which would have been correct)



# Database Characteristics

- **Persistence** of data and disaster **recovery**
  - Data needs to be persistent and accessible at all times
  - **Quick** recovery from system crashes **without data loss**
  - Recovery from natural disasters (fire, earthquake, ...)



- Introduction
  - What is a database?
  - Characteristics of a database
  - Database users
- History of databases



# Database Users

- Usually **several groups of persons** are involved in the daily usage of a large DBMS (many job opportunities for smart DB people...)
- Persons directly involved on DB level
  - **Database administrators**
    - Responsible for tuning and maintaining the DBMS
    - Management of storage space, security, hardware, software, etc.
  - **Database designers**
    - Identifies the data that needs to be stored and chooses appropriate data structures and representations
    - Integrates the needs of all users into the design



# Database Users

- **Application developers**

- Identify the requirements of the end-users
- Develop the software that is used by (naïve) end-users to interact with the DB
- Cooperate closely with DB designers



- Persons working behind the scenes

- **DBMS designers and implementers**

- Implement the DBMS software

- **Tool developers**

- Develop generic tools that extend the DBMS' functionalities

- **Operators and maintenance personnel**

- Responsible for actually running and maintaining the DBMS hardware



# Database Users

- **End users**
  - All people who use the DB to do their job
- End users split into
  - **Naïve end users**
    - Make up most DB users
    - Usually **repeat** similar tasks over and over
    - Are supported by predesigned interfaces for their tasks
    - Examples: bank tellers, reservation clerks, ...





# Database Users

## – Sophisticated end users

- Require **complex** non-standard operations and views from the DB
- Are familiar with the facilities of the DBMS
- Can solve their problems themselves, but require complex tools
- Examples: engineers, scientists, business analysts, ...



## – Casual end users


- Use DB only from time to time, but need to perform different tasks
- Are familiar with query languages
- Examples: People in middle or senior management






# History of DBs

- Databases have an exceptional history of development
  - Many synergies between **academic**, **governmental** and **industrial** research
  - Much to be learned from it
  - Most popular concepts used today have been invented decades ago



# History of DBs

- The beginnings
  - 1880: U.S. Bureau of Census instructs **Herman Hollerith** to develop a machine for storing census data
  - Result: **Punch card** tabulator machine
    - The evaluation of 1880's census took 8 years
    - 1890's has been finished after only one year
  - Leads to the foundation of **IBM**
    - International Business Machines
  - Data processing machines soon established in accounting





# History of DBs

- During WWI, many **data collection, sorting** and **reporting** tasks in industry and governmental organizations was performed using **punch cards**
- **1935:** U.S. Social Security Act required continuous report on all 26 million governmental employees
  - “World’s biggest bookkeeping job”
  - Mechanical punch card systems not powerful enough
    - IBM develops the electric **UNIVAC I** punch card processor in 1951





# History of DBs

- In **1959**, U.S. dominated the (still highly active) punch card machine market
  - Within the U.S., the Pentagon alone used more than 200 data processing computers, costing \$70 million per year
- In **1964**, the term “data base” appeared for the first time in military computing using **time sharing systems**
  - Data could be shared among users
  - Data model very close to punch cards
  - **Master files** bound to a specific application



# History of DBs

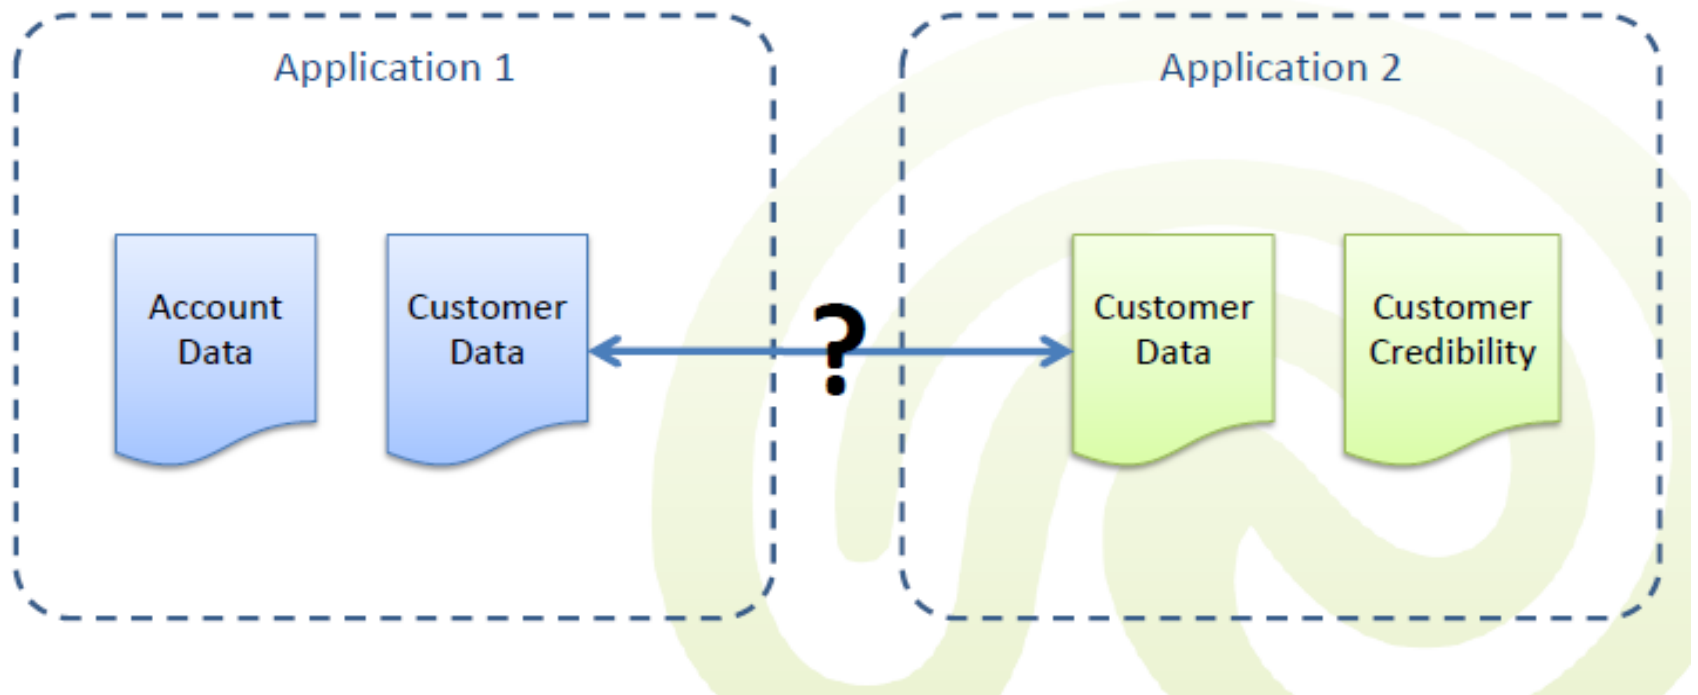
- **Master files**

- Used to maintain continuity between program runs
- Each application had its own master files
- Similar data needed by multiple applications had to be duplicated
  - Consistency problems when updating data
- Highly-dependent on the hardware and (low-level) programming language used
  - Inspired by punch cards and optimized for magnetic tapes
  - Usually, no **relationships** between different records have been stored, just plain data

# History of DBs

- **Master files**

- Highly hardware-oriented approach
- Data stored in independent flat files





# History of DBs

- To turn stored data into a proper **database**, the following goals had to be achieved (McGee, 1981):
  - **Data consolidation**
    - Data must be stored in a central place, accessible to all applications
    - Knowledge about relationships between records must be represented
  - **Data independence**
    - Data must be independent of the specific quirks of the particular low level programming language used
    - Provide high-level interfaces to physical data storage
  - **Data protection**
    - Data must be protected against loss and abuse

# History of DBs

- **Data consolidation** motivated the development of data models
  - Hierarchical data model
  - Network data model
  - **Relational data model**
  - Object-oriented data model
  - Semantic data model
- **Data independence** inspired the development of query models and high-level languages
  - **Relational Algebra, SQL**
- **Data protection** led to development of transactions, backup schemes, and security protocols

# History of DBs

- **Hierarchical data model**



- First appearance in **IBM's IMS** database system, designed for the Apollo Program in **1966**

- Still, as of 2006, 95% of all Fortune 1000 companies use IBM IMS in their data backbone...

- Benefits from **advances** in **hardware** design

- Random access main memory and tape media available

- Data may be organized in a **tree structure**

- Initially, tree had maximum depth of two



# History of DBs

- **Hierarchical data model**

- Each type of record has one or multiple own files/tables
- Hierarchical **one-to-many** relationships
- Vaguely resembles a file system organization

Customer

customerID	Name
1000	Mickey Mouse
1001	Scrooge McDuck
1002	Donald Duck
1003	Goofy

Accounts – Mickey Mouse

number	balance
1000	€ 73.68

Accounts – Scrooge McDuck

number	balance
6000	\$ 934,3243,435,322
6001	€ 4,543,123,987
6002	¥ 12,432,355,112



# History of DBs


- **Hierarchical data model**

- **Advantages**

- 1:n relationships can be expressed
    - Can easily be stored on tape media

- **Disadvantages**

- No n:m relationships
    - No data independence



# History of DBs

- **Network data model**

- In the mid-1960th, direct access storage devices (DASD) gained momentum
  - Primarily hard disks
  - More complex storage schemes possible
- Hierarchical model failed, e.g. for bill-of-material-processing (BOMP)
  - Many-to-many relationships needed
  - Development of the IBM DBOMP system (1960)
- Result: Network model
  - Two types of files: Master files, chain files
  - Chain file entries could chain master file entry to one another



# History of DBs

- **Network data model**

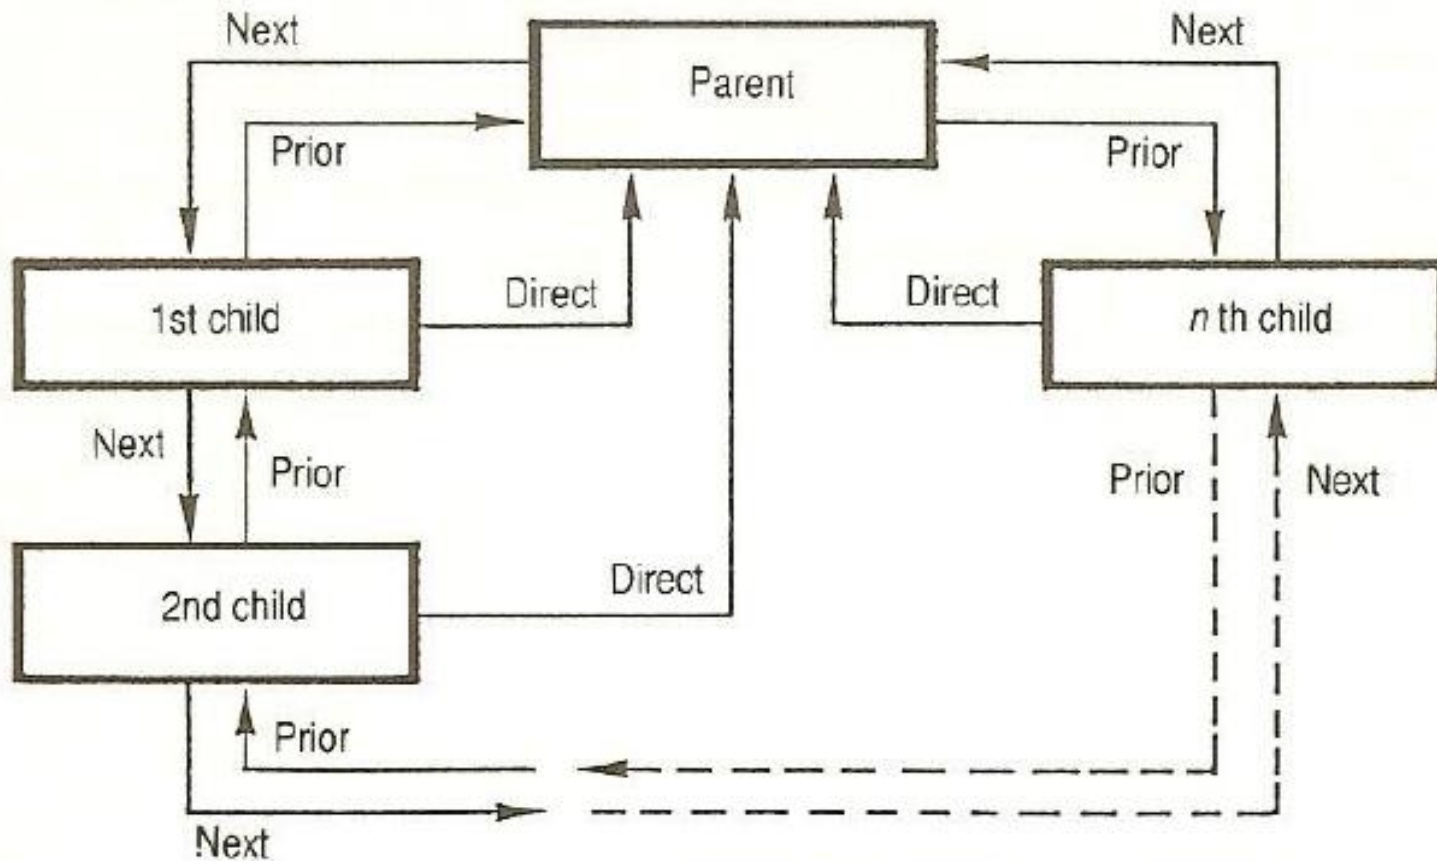
- The model was standardized by Charles W. Bachman for the **CODASYL** Consortium in 1969
  - CODASYL = Conference of Data Systems Languages
  - Thus, also called the CODASYL model
- Allowed for more natural modeling of **associations**
- Advantage
  - **Many-to-many-relationships**
- Disadvantages
  - No declarative queries
  - Queries must state the data access path





# History of DBs

- **An example of the network data model:**



# History of DBs

- What's wrong with all that?
  - Strong degree of **hardware dependence**
  - **No proper abstraction** of data
  - No decoupling of data and its application
  - Each database needed to be **“hand-crafted”** for its application
  - To change something in the data-schema, “a sharp-looking guy in a white shirt and black rims had to do the programming by hand”
    - No formal/structural/mathematical foundation  
→ **no high-level data languages**



Figure 1.1. A data communications specialist looks over computer program printed on Model 15 Tektronix for early data communications system. Datacom data communications process (background) monitors and stores messages to and from company.

# History of DBs




- **The relational data model**
- Published by **Edgar F. “Ted” Codd** in 1970, after several years of work
  - *A Relational Model of Data for Large Shared Data Banks*, Communications of the ACM, 1970
  - Employee of IBM Research
    - IBM **ignored** his idea for a long time as not being “practical” while pushing its hierarchical IMS database system
    - Other researchers in the field also **rejected** his theories
    - Finally, he received the Turing Award in 1981

# History of DBs

- **Idea underlying the relational model:**
  - Database is seen as a collection of **predicates** over a finite set of **predicate variables**
    - Example:
      - is\_supervisor\_of( $x, y$ )
      - is\_supervisor\_of('W.-T. Balke', 'B. Köhncke') (TRUE)
      - is\_supervisor\_of('C. Lofi', 'B. Köhncke') (FALSE)
    - The set of all true assignments is called a **relation**
    - Relations are stored in **tables**
  - Contents of the DB = a collection of **relations**
  - Queries are also **predicates**
    - Queries and data are very similar
    - Allows for **declarative querying**

# History of DBs

- It's really like a collection of index cards
  - More details during the next weeks...





# History of DBs

- Beginning 1977, **Lawrence J. Ellison** picked up the idea and created **Oracle DB** (currently in version 11g)
  - And became insanely rich – long time in the Top 10 of the richest people
  - In 2007 Oracle ranked third on the list of largest software companies in the world, after Microsoft and IBM
  - Oracle's expected net income in 2009: **\$5.59 billion**



# History of DBs

- Oracle also sells a suite of **business applications**
  - Oracle eBusiness Suite
  - Includes software to perform financial- and manufacturing-related operations, customer relationship management, enterprise resource planning, and human resource management
- Basically gained from high-value acquisitions beginning in 2003
  - JD Edwards, PeopleSoft, Siebel Systems, BEA, ...





# History of DBs

- During the 1970s, IBM had also decided to develop a relational database system
  - **System R** with the first implementation of the **SQL** declarative query language (SEQUEL)
  - At first, mostly a research prototype, later became the base for **IBM DB2**
- Today, the relational model is the **de-facto standard** of most modern databases



# History of DBs

<b>Company</b>	<b>Revenue Estimates for Database Products, 2006 (in million USD)</b>	<b>Market Share (in %)</b>
Oracle	7,168.0	47.1
IBM	3,204.1	21.1
Microsoft	2,654.4	17.4
Teradata	494.2	3.2
Sybase	486.7	3.2
Other Vendors	1,206.3	7.9
Total	15,213.7	100.0

**Source: Gartner Research, 2007**

**Gartner**

# History of DBs

Year	Event
1880	Hollerith census machine
1951	Univac I electrical data machine
1959	First CODASYL Conference
1960	Flight reservation system SABRE
1966	IMS hierarchical database
1969	Network model
1971	CODASYL Recommendation for DDL and 3-Layer-Architecture
1975	System R introduces SEQUEL query language
1976	System R introduces transaction concepts
1976	Peter Chen proposes entity relationship modeling
1980	Oracle, Informix and others start selling DBMS with SQL support

# History of DBs

Year	Event
1983	Work on ACID transactions published by Theo Haerder and Andreas Reuter
1986	SQL standardized as SQL-1 ANSI/SQL
1987	SQL internationally standardized as ISO 9075
1989	SQL 2 standard supports referential integrity
1991	SQL 2 supports domains and key definitions
1993	Object-oriented data model
1995	Preliminary SQL 3 supporting sub-tables, recursion, procedures, and triggers
1996	First object-oriented databases
1999	First part of the SQL 3 standard finalized
2003	SQL 2003 finalized with support for object-relational extensions
...	<b>To be continued ...</b>

# History of DBs

- **Beyond the relational model...**
  - Data models based on formal logic
    - Deductive databases and expert systems
  - Object-oriented data models
    - Main Idea: Object-oriented design (garage metaphor)
    - Very easy integration in OO programming languages
    - Today, mostly integrated into the relational model
  - Semi-structured data models
    - Most important: XML
    - Allows a large degree of structural freedom
  - For details, take the master's courses on these topics ...

# Summary

- Databases
  - are **logical interfaces**
  - **control redundancy**
  - are **well-structured**
  - support **declarative querying**
  - aim at **efficient manipulation** of data
  - support **multiple views** of the data
  - support **atomic multi-user transactions**
  - support **persistence** and **recovery** of data
- (Several groups of database users)
- (History of databases)

